

Ford Motor Data Engineer Interview Guide – Experienced 3+

Technical Round 1

1. Introduction

The interview began with a standard introduction where I briefly talked about my professional journey, current role, key projects, and technologies I work with.

2. Pipeline Design and Explanation

I was asked to explain a data pipeline from my project. I described an end-to-end GCP-based pipeline handling data ingestion from Oracle into Google Cloud Storage (GCS), processing with Dataproc (PySpark jobs), and storing results in BigQuery for reporting.

3. GCP Tools

I highlighted key tools:

- Dataproc for running Spark jobs
- BigQuery for analytical querying
- Pub/Sub for messaging
- Firestore for semi-structured data
- GCS for storage

4. Day-to-Day Activities

I shared a typical day involving data pipeline monitoring, developing new transformations, resolving production issues, and collaborating with data analysts to improve query performance.

5. Clustering in Real-Time

I explained clustering with a real-time example of customer segmentation in retail data using k-means for grouping similar purchasing behaviors.

6. Production Issues and Approach

I described how I resolved a production issue involving delayed Pub/Sub messages by debugging logs, identifying bottlenecks, and optimizing Spark job configurations.

7. Agile: Adding New Requirements

Explained how new requirements are added via change requests reviewed during backlog grooming sessions.

8. Sprint Duration

Sprint duration in my project is two weeks.

9. Source and Target Systems

Source: Oracle database

Target: BigQuery

10. Joins in SQL

Explained INNER, LEFT, RIGHT, and FULL joins with examples.

11. Types of Views

Covered Simple, Materialized, and Indexed views.

12. Native vs. External Tables

Native tables store data within the database, while external tables reference data stored externally (e.g., in GCS).

Technical Round 2

13. Access Granting for Views

I mentioned that access is managed by a DBA or Security Team based on roles and policies.

14. Source Issue Resolution

Explained interaction with the source team, using tools like Jira to raise issues and Slack for real-time communication.

15. Script Implementation

Described deployment through CI/CD pipelines using Git and Jenkins.

16. P1 Issues

Shared experience handling a P1 issue involving data loss by restoring from a backup and reprocessing missed data.

17. P2 Ticket SLA

P2 tickets typically have a resolution time of 24-48 hours.

18. Removing Duplicates with Partitioning

Using ROW_NUMBER() over (PARTITION BY...) to remove duplicates.

19. Inserting Columns Between Others

Mentioned it requires creating a new table with desired schema and migrating data (as direct insertion isn't supported).

20. Additional Calls

Apart from standard daily stand-ups, we have weekly sync-ups with stakeholders.

21. Ticket Beyond Story Point Duration

Explained how we request an extension with proper justification during sprint reviews.

22. User Interaction

Yes, I interact directly with business users for understanding requirements and gathering feedback.

23. Notice Period and Relocation

Shared my notice period and willingness to relocate.

Glassdoor Ford Review –

<https://www.glassdoor.co.in/Reviews/Ford-Motor-Company-Reviews-E263.htm>

Ford Careers –

<https://www.careers.ford.com/en/home.html>

Subscribe to my YouTube Channel for Free Data Engineering Content –

<https://www.youtube.com/@shubhamwadekar27>

Connect with me here –

<https://bento.me/shubhamwadekar>

Checkout more Interview Preparation Material on –

https://topmate.io/shubham_wadekar

© Shubham Wadekar