

# EMR SCENARIO BASED Q&A

**BY - SHUBHAM WADEKAR**

Copyright © Shubham Wadekar. All rights reserved.

This material is for personal use only. No part may be copied, shared, resold, or published without written permission. Unauthorized distribution is strictly prohibited and may result in action.

For permissions or licensing, contact: [shubham.p.wadekar@gmail.com]

Disclaimer: This content is for educational purposes. Accuracy is attempted but not guaranteed. No job outcome is promised.

## 1. How would you decide the number and type of master, core, and task nodes in an EMR cluster for large Spark ETL workloads?

I start from the workload pattern (throughput, peak hours, data size, shuffle intensity) and map that to node roles and instance families.

1. Pick the right instance families
  - Primary node (formerly master): this runs cluster services and the Spark driver for cluster mode. I choose a balanced instance (m5/m6i) with enough memory and network, not the smallest one. If I expect very large query plans or many concurrent jobs, I size it bigger for driver memory and UI responsiveness.
  - Core nodes: these store HDFS blocks (if I use HDFS) and run executors. For shuffle-heavy ETL, memory and network matter most, so I prefer r5/r6i for memory or m5/m6i if balanced is fine. If I need very fast local disk (large spills), i3/i4i can help.
  - Task nodes: these do compute only (no HDFS). I almost always use Spot here to cut cost. Same family as core nodes for homogeneity.
2. Choose storage pattern before sizing
  - If I rely on EMRFS (S3) for storage and shuffle to disk, I can keep few or zero HDFS core nodes and scale with task nodes.
  - If I need HDFS (e.g., very high shuffle throughput, external services expecting HDFS), I add enough core nodes to hold the working set with headroom and leave task nodes for burst compute.
3. High availability and control plane
  - For production with HDFS, I use 3 primary nodes (multi-master) for HA. For EMRFS-only, 1 primary is usually fine, but I still add instance-fleet diversification across AZs where possible.
  - Enable EMR managed scaling so the cluster grows during peaks and shrinks off-peak.
4. How I estimate node counts
  - Start from needed executors: target total cores for your daily SLA. For example, if I want about 1,000 vCPUs at peak and I choose r5.4xlarge (16 vCPU), I need roughly 70 worker instances at peak (allowing a few vCPUs for the OS/YARN).
  - Keep a core:executor layout like this: core nodes sized to hold HDFS (if used) and minimum steady executors, task nodes to add burst capacity. A common split is 20–30% cores, 70–80% tasks.
  - Spot for tasks with 2–3 instance types per fleet and 30–50% max Spot loss tolerance. On-Demand for core nodes to protect HDFS.
5. Per-node executor sizing (simple, repeatable method)
  - Leave ~1 vCPU per node for the OS/YARN.
  - Use 4–5 cores per executor (good balance for shuffle and GC).
  - Leave 6–10% memory for overhead.

Example on r5.4xlarge (16 vCPU, 128 GB): leave 1 vCPU → 15 usable. Use 3 executors × 5 cores each. Memory per executor  $\approx (128 \text{ GB} \times 0.9) / 3 \approx 38 \text{ GB}$ . I would round to something like 36G executor memory with 4G overhead.

  - Keep the same layout on task and core nodes for predictability.
6. Network and disks
  - Prefer instances with higher network bandwidth for shuffle-heavy jobs.
  - Use larger, fewer files in S3 (128–512 MB) to avoid tiny-file overhead. If using EBS, provision enough throughput; for heavy spills, i3/i4i with NVMe is excellent.

7. Autoscaling and guardrails
  - Enable EMR managed scaling with metrics like YARN pending memory and vCores.
  - Cap max cluster size to protect cost.
  - Scale up early before the daily peak and scale down gradually to avoid losing cached executors too aggressively.
8. Validate and iterate
  - Run a peak-load test, watch stage times, shuffle read/write, and executor GC in the Spark UI. Adjust instance family or counts if you see sustained spill or network saturation.
  - If drivers run out of memory on the primary, upsize the primary or move drivers to client mode from a separate orchestrator host for very large plans.

In short, I select memory-optimized workers for shuffle, use On-Demand core nodes plus Spot task nodes, size executors consistently per node, enable managed scaling, and iterate with UI metrics until the SLA is comfortably met.

## 2. Your Spark jobs are slowing during peak load. How would you analyze drivers and executors across nodes to find bottlenecks?

I follow a simple checklist: verify the driver, then the cluster resources, then the stages and shuffles, and finally storage and skew.

1. Check the driver first
  - Spark UI: look for long “Job submitted → running” gaps, many retries, or the driver spending time in GC.
  - Driver memory: if the driver has large collect operations or big query plans, increase driver memory or move expensive collect operations to aggregates.
  - Concurrency: too many jobs on the same driver can serialize scheduling. Limit concurrent jobs or use a job server pattern with queues.
2. Verify cluster resource pressure
  - YARN/EMR metrics: pending vCores and memory mean you need more executors or bigger nodes. If pending stays high during peak, scale out.
  - Executor counts: did autoscaling add task nodes? If not, adjust thresholds or upper bounds.
  - Executor lost/replaced events: Spot interruptions or unstable nodes will slow stages; diversify Spot and add On-Demand buffer.
3. Inspect stages and tasks in Spark UI
  - Look at the slowest stages. If shuffle read time dominates, the bottleneck is network or disk spill.
  - Task timeline: if tasks start slowly in waves, you might have too few partitions or uneven partition sizes. Increase parallelism or repartition to a higher count.
  - GC and memory: long GC in tasks indicates oversized partitions or too much data per task. Reduce executor cores per executor or increase memory; also tune the number of partitions.
4. Shuffle and spill diagnostics
  - Check “Shuffle Read/Write Size” and “Bytes Spilled”. High spill means not enough memory or too large partitions.
  - Raise shuffle partitions for wide transformations or enable adaptive query execution so Spark adjusts at runtime.
  - Ensure local disks are fast; on EBS, throughput might be the choke point. Instances with NVMe (i3/i4i) help.
5. Skew and hot keys
  - If a few tasks run much longer than the rest, it’s data skew. Confirm by looking at per-task input sizes.
  - Mitigations: pre-aggregate, add salting to skewed keys, use skew join handling or AQE skew hints, or change join order to broadcast the small side.
6. Join strategy
  - If one table is small enough, make sure it broadcasts; if it is just above the limit, increase the auto-broadcast threshold.
  - If both are large, sort-merge join is typical; ensure both sides are partitioned consistently to reduce shuffle.
7. File layout and source pressure
  - Many tiny files in S3 will slow listing and planning. Compact to 128–512 MB files.
  - Prune data early: push down filters and columns so tasks read less from S3.
  - Check schema mismatches that cause excessive casting or parsing overhead.

8. Configuration quick wins
  - Enable adaptive query execution to handle skew and adjust shuffle partitions.
  - Set reasonable executor cores (4–5) and memory with overhead. Too many cores per executor can cause GC storms.
  - Tune shuffle partitions to reflect input size; a simple starting point is  $2\text{--}3\times$  the total executor cores, then let AQE refine it.
9. Observe, change one thing at a time, and retest
  - Capture a baseline: stage times, shuffle sizes, GC time, spilled bytes.
  - Apply one change, rerun the same workload slice, and compare. Keep changes that reduce the bottleneck.

In short, I use the Spark UI and EMR metrics to see whether the driver, resources, shuffle, or skew is the culprit. Then I fix the specific bottleneck: scale executors, improve file layout, enable AQE, adjust partitions, upgrade disk/network, or change join strategies.

### 3. How would you design EMR with multiple masters to reduce downtime in case of master node failure?

I would create the cluster with three primary (master) nodes so the control plane is highly available. In EMR this gives me high availability for both HDFS NameNode and YARN ResourceManager. If one primary node goes down, the others can take over and the cluster keeps running.

My simple design steps are:

1. Create the cluster with three primary nodes, not one. This turns on HDFS and YARN high availability automatically, so failover is handled for me.
2. Keep core nodes on On-Demand so the HDFS data layer is stable. Add task nodes on Spot for cheap extra compute, because losing a task node does not lose HDFS data.
3. Use the same instance family for core and task nodes to keep executors consistent. Pick memory-optimized or balanced instances depending on how heavy the shuffles are.
4. Put the driver of long jobs on the cluster only if I'm using cluster mode and I sized primary nodes well. For very large drivers, I run in client mode from an external orchestrator so a primary node failover does not interrupt the driver.
5. Store checkpoints and critical intermediates in S3, not only in HDFS. This way, if a node fails and a job retries, it can continue safely.
6. Enable managed scaling and set safe minimum core capacity so the cluster can grow during peaks and still keep quorum for HDFS and YARN.
7. Add monitoring and alarms on primary node health, HDFS, and YARN metrics. If a primary fails, I get alerted, but the workload should continue.

In short, using three primary nodes gives me automatic failover. Keeping state in S3, using stable core nodes, and good monitoring makes the cluster ride through a master failure with minimal impact.

#### 4. When would you choose long-running EMR clusters vs transient clusters triggered by Step Functions?

I choose based on workload pattern, cost, and isolation.

When I prefer long-running clusters:

1. There is a steady, all-day workload or many jobs spread across the day. A warm cluster avoids repeated startup time.
2. I want interactive work (notebooks, ad-hoc queries) or low latency between jobs.
3. I benefit from caching and local HDFS (for example, repeated joins on the same data during the day).
4. I need multi-tenant scheduling with YARN queues and fair sharing across teams.
5. Tooling is attached to the cluster (Ganglia, persistent UIs) and I want continuous visibility.

Trade-offs: can cost more if idle; needs capacity and patch management; version upgrades are harder because the cluster is always on.

When I prefer transient (per-job) clusters orchestrated by Step Functions or workflows:

1. Batch jobs that run on a schedule (hourly/daily) and can tolerate a few minutes of cluster startup.
2. Strong need for cost control and isolation. Each job gets its own cluster, then auto-terminates, so there is no idle cost and no noisy neighbors.
3. Easy version pinning and upgrades. I can spin a cluster with a specific EMR release and libraries per pipeline.
4. Better blast-radius control. One bad job does not affect others because they run on separate clusters.
5. Compliance needs where each pipeline must be isolated with its own IAM roles, security settings, and tags.

Trade-offs: startup overhead per run, less caching benefit, more clusters to view in logs, and you need orchestration to pass parameters and collect outputs.

A simple rule I follow:

- If jobs are spiky, independent, and cost control/isolation is important, use transient clusters with Step Functions or managed workflows.
- If jobs run most of the day, are interactive, or benefit from caching/HDFS, use a right-sized long-running cluster with managed scaling.

## 5. A Spark job runs poorly even with many task nodes. How would you use YARN and node types to improve parallelism?

I start by checking if the cluster can actually run enough executors in parallel and whether each node is sliced correctly for Spark.

First, I size executors per node in a simple, repeatable way. I leave about one vCPU for the OS on each node, then divide the remaining cores into executors with around four to five cores each. I do the same with memory: keep a small headroom for the OS and Spark overhead, then split the rest equally across those executors. This usually gives me multiple executors per node, which increases parallel tasks without causing garbage collection issues.

Second, I align YARN resources with the instance capacity. I set YARN's node memory and vCores to match the real resources on the node so YARN can place all the executors I planned. If YARN shows high pending containers, it means executors are waiting and I need either bigger nodes, more nodes, or smaller executors so more can fit.

Third, I make node roles work for me. Core nodes are stable (On-Demand) and hold HDFS if I use it. Task nodes are compute only, so I use more of them on Spot to raise parallelism cheaply. I keep the instance family the same across core and task nodes so executors behave consistently. If I must mix families, I try to keep similar CPU and memory ratios.

Fourth, I remove hidden bottlenecks that block parallelism. I increase the number of shuffle partitions so I have at least as many tasks as total executor cores. I turn on adaptive query execution so Spark can split skewed partitions at runtime. I also fix tiny-file problems by compacting inputs to 128–512 MB so tasks are balanced and start quickly.

Fifth, I let the cluster scale correctly. I enable EMR managed scaling on YARN signals like pending memory and pending vCores so the cluster adds task nodes when the queue backs up. I also make sure speculative execution is on to replace slow tasks that hold up a stage.

Finally, I watch the Spark UI and YARN metrics. If I see long queues or many pending containers, I shrink executor size a bit so more can run at once. If I see heavy spilling, I choose instances with faster local disks (like i3 or i4i) or add larger EBS volumes with enough throughput.

In short, I right-size executors per node, make sure YARN can actually place them, use many task nodes of the right type, raise the number of partitions, and let autoscaling react to real pressure.



## **6. Spark ETL creates large shuffle data. Why does using EMRFS instead of HDFS reduce performance, and how would you fix it?**

Using S3 through EMRFS is slower for heavy shuffle because S3 is an object store, not a filesystem. There is no data locality, operations like rename are expensive, and every read or write goes over the network with higher latency. HDFS, on the other hand, keeps data on the cluster's local disks, supports fast sequential IO, and lets executors read from nearby disks, which is much faster during shuffle.

To fix this, I keep all intermediates on the cluster and use S3 only for final outputs. I enable HDFS and let Spark spill shuffle data to local disks or HDFS, not to S3. On EMR, this is straightforward: choose instances with fast local NVMe disks (for example, i3 or i4i) or attach large, high-throughput EBS volumes, and make sure Spark's shuffle and spill directories point to those local paths. This keeps the most expensive part of the job on fast storage.

For the final write to S3, I reduce S3 overhead by using the S3 optimized/committers so large jobs don't suffer from many slow renames. I also write in Parquet and coalesce to fewer, larger files to cut down on small-object churn. If I must stage data between steps, I stage on HDFS and only copy to S3 at boundaries.

Additionally, I improve network and parallelism. I pick instance families with better network bandwidth for shuffle-heavy jobs, increase shuffle partitions so work is spread evenly, and turn on adaptive query execution to handle skew automatically.

In short, I avoid using S3 for shuffle or temporary data. I keep intermediates on HDFS or fast local disks and only write the final results to S3, using optimized committers and properly sized files to keep performance high.

## 7. For long-term data storage, would you recommend HDFS or EMRFS, and why?

For long-term storage, I would recommend EMRFS (S3) instead of HDFS. The main reason is durability and scalability.

HDFS stores data on the local disks of the EMR cluster. This means if the cluster is terminated, or if several nodes fail, the data could be lost unless I copy it out. HDFS is good for temporary, high-speed shuffle and intermediate storage, but it is not reliable for long-term data storage. It also requires me to keep the cluster running, which adds cost even when jobs aren't running.

EMRFS (which is basically Spark/Hadoop using S3) solves these issues. S3 is highly durable (eleven 9s), scalable, and cost-effective for long-term data. I can store petabytes of data without managing disks, and I don't have to keep EMR clusters alive to keep my data safe. Any new EMR cluster or Athena query can immediately reuse the data from S3.

So my general approach is:

- Use HDFS or local disks for temporary shuffle or intermediate data that needs very fast access during job execution.
- Use EMRFS (S3) for all raw data, curated datasets, and long-term storage.

In short: for long-term storage, I would always recommend EMRFS (S3) because it's durable, cost-effective, and accessible across clusters, while HDFS is only good for temporary, short-term storage within a single cluster.

## 8. A team sees inconsistent query results with EMRFS + S3. How would you explain strong vs eventual consistency?

The difference comes from how S3 handles metadata updates.

Strong consistency means when you write or delete a file, any new read immediately sees the latest state. Eventual consistency means there can be a short delay before changes are visible, so some reads may return stale results.

Before 2020, S3 had eventual consistency for overwrite and delete operations. That meant if one job overwrote or deleted files, another job could still see the old data for a short time. This often caused inconsistent query results with EMRFS, because Spark or Hive might list a prefix and pick up old files or miss new ones.

Now, S3 provides strong read-after-write consistency for both new objects and overwrite/delete across all regions. This has reduced the issue a lot, but teams can still see “inconsistent” results if:

- They have multiple writers writing to the same S3 prefix at the same time (no coordination).
- They use non-optimized committers, which rely on renames (renames in S3 are actually copy+delete and can be temporarily inconsistent).
- They read from partially written data before a job finishes writing.

My explanation to the team would be: EMRFS + S3 is strongly consistent for object writes, but query inconsistency can still happen due to how writes and renames are implemented. To avoid it, I would use the S3 optimized committer, write data to unique output paths for each job run, and only point queries at a final “success” marker folder once the write is complete.

In simple terms: S3 now guarantees strong consistency for file operations, but to avoid inconsistent queries, we need to design the pipeline so jobs don't read incomplete or overlapping writes.

## 9. How would you combine HDFS (fast access) with EMRFS (durable storage) in a pipeline?

I use a simple pattern: keep all permanent data in S3 via EMRFS, and use HDFS or local disks only for temporary, high-speed work during Spark jobs.

Here's the flow I follow:

1. Land and catalog in S3: Raw and curated datasets live in S3 (Parquet/ORC), registered in the Glue Data Catalog so they're reusable by EMR, Athena, and Redshift.
2. Read from S3, stage on HDFS: Each Spark job reads input from S3, then performs heavy transformations. During shuffles and spills, Spark writes temporary data to HDFS or local NVMe (fast). This keeps the most expensive I/O off S3.
3. Write final results back to S3: When the job completes, I coalesce to right-sized files (128–512 MB), use the S3-optimized committer to avoid slow renames, and write the final outputs back to S3.
4. Use checkpointing wisely: If the pipeline has long lineage or needs fault tolerance, I checkpoint to HDFS for speed during the run. If I need durability across runs, I checkpoint to S3 at logical boundaries.
5. Cleanup: Because HDFS is ephemeral, I do not depend on it between jobs. Each run recreates what it needs locally and writes durable results to S3.

In short, S3 is the system of record and sharing layer; HDFS is the job's scratchpad for speed. This gives fast execution and durable outputs without keeping clusters alive just to retain data.

## 10. HDFS usage is driving up costs. How would you balance HDFS (ephemeral) vs EMRFS (durable) for cost and reliability?

I reduce HDFS to the minimum needed for performance, and move everything else to S3.

What I change:

1. Store all long-term data in S3: Raw, cleaned, and curated tables live in S3 (Parquet/ORC). This lets me shut down clusters with zero data loss and reuse data across new clusters.
2. Keep HDFS small: Size core nodes for just enough HDFS to handle shuffle/intermediate needs. Push most compute to task nodes (often Spot) so I get parallelism without paying for large persistent HDFS capacity.
3. Use fast local disks for shuffle: Prefer instances with NVMe (i3/i4i) or attach EBS with enough throughput. Point Spark's temp and shuffle dirs to local storage so I don't need big HDFS space.
4. Write once, compact, and control files: On S3, write with the S3-optimized committer and compact to larger files to cut S3 request costs and speed reads. Avoid tiny files.
5. Autoscale instead of overprovisioning: Enable EMR managed scaling based on YARN pending metrics. Scale up during peaks and scale down after, so I'm not paying for idle HDFS.
6. Clear the boundaries: Treat HDFS as ephemeral. If a job needs durable checkpoints or handoff between steps, write that to S3. Don't keep intermediate data on HDFS longer than the job lifetime.
7. Guardrails and monitoring: Track HDFS utilization and spill metrics. If HDFS usage stays high, either add fast local storage or optimize Spark partitions to reduce spill. Alert when HDFS passes a safe threshold so jobs don't fail and I don't overspend.

In short, I keep HDFS lean for speed within a run, and I use S3 for everything durable. This lowers steady-state cost, keeps reliability high, and still gives good performance for heavy transforms.

## 11. A Spark job fails due to executor memory errors. How would you tune driver and executor configs without over-provisioning?

I first confirm where the memory error happens: on the driver or on the executors. I check the Spark UI and logs for “OutOfMemoryError” and see which stage and which component failed. Then I fix the root cause step by step, keeping resources reasonable instead of just throwing big machines at it.

How I fix executor memory errors in a simple way:

1. Right-size cores per executor  
I avoid very large executors. I keep about 4–5 cores per executor so each task gets a fair share of memory and GC is stable. Many small-to-medium executors are safer than a few huge ones.
2. Set memory and overhead cleanly  
I leave 6–10% for overhead. For example, if I give an executor 24G, I add about 3G overhead. On EMR, I make sure executor memory + overhead fits within the YARN container size.
3. Reduce data per task  
If tasks handle too much data, they blow the heap. I increase shuffle partitions so each task gets a smaller slice. I also compact input files to 128–512 MB so partitions are even.
4. Cut unnecessary data early  
I read only the columns I need (Parquet/ORC pushdown) and filter early. I avoid SELECT \*. Less data in equals less memory pressure.
5. Tame joins and skew  
I broadcast only truly small tables. If a join key is skewed (few keys have huge data), I add salting or enable adaptive query execution so Spark splits skewed partitions. This removes a handful of giant tasks that cause OOM.
6. Manage caching and lineage  
I cache only what I will reuse several times, and I unpersist as soon as I am done. If lineage is very long, I add a checkpoint at a logical point so Spark does not keep huge DAG histories.
7. Safer serialization and GC  
I use Kryo serialization. If I still see long GC pauses, I reduce cores per executor slightly (for example from 5 to 4) so each Java heap serves fewer active tasks.
8. Keep shuffle on fast local storage  
I point Spark’s shuffle dirs to NVMe or good EBS so spilling is cheap when it happens. Spilling to slow disks makes OOM more likely because tasks run longer.

How I fix driver memory errors:

1. Avoid collect and show on big data  
I replace collect with aggregates or samples. If I must inspect data, I write a small sample to S3 instead of pulling everything to the driver.
2. Increase driver memory modestly  
I give the driver a reasonable bump (for example 4G → 8–12G) and keep the plan size under control by breaking a huge job into stages.

Simple starting template I use:

- 4–5 cores per executor
- executor memory sized so it fits well with 6–10% overhead
- shuffle partitions around 2–3 times total executor cores, then let adaptive query execution refine it
- only needed columns and rows, avoid big collect, cache sparingly

This usually solves OOM without overspending.

© Shubham Wadekar

## 12. A client insists on MapReduce for large log processing. How would you justify Spark vs MapReduce on EMR?

I would explain it in three plain points: speed, simplicity, and cost on EMR.

### 1. Speed

Spark keeps data in memory between steps and runs a whole DAG at once. MapReduce writes to disk between every map and reduce phase. For multi-step ETL on logs, Spark avoids many disk writes and is typically much faster. Even when data does not fit in memory, Spark spills efficiently and still beats MapReduce for most pipelines.

### 2. Simplicity and features

With Spark SQL and DataFrames, I can express complex joins, windows, and aggregations in a few lines. I get built-in optimizations, adaptive query execution, and easy JSON/Parquet handling. MapReduce needs a lot of boilerplate for every step and is harder to maintain. Spark also gives me one engine for batch, interactive analysis, and even streaming, so the team learns one tool.

### 3. Cost and operations on EMR

On EMR, Spark benefits from the EMR runtime optimizations, managed scaling, Spot capacity for cheap task nodes, and S3 integration with Parquet. Jobs finish faster, clusters run for fewer hours, and I can use transient clusters per pipeline for isolation. MapReduce jobs tend to run longer and require more disk I/O, which translates to more hours and more storage throughput.

When MapReduce can still be okay:

- Very simple, single-step batch transforms where performance is not critical
- Legacy code that is rarely touched and already stable

How I propose a safe path:

- Rebuild one representative MR job in Spark using DataFrames and Parquet
- Measure runtime, cost, and code size
- Show the result and roll out gradually

In short, Spark on EMR gives faster pipelines, simpler code, and lower total cost for typical large log processing. MapReduce still works, but Spark is the better long-term choice for most data engineering workloads.



### 13. Spark jobs are slow even on memory-optimized nodes. What would you check (executor parallelism, node types, cluster sizing)?

I start by confirming whether the cluster can run enough tasks in parallel and whether each task is doing a fair amount of work. Then I check storage and shuffle speed.

1. Executor parallelism

I avoid very large executors. I target about 4–5 cores per executor and leave some memory overhead. If executors have too many cores, a few tasks hog memory and GC gets heavy. If they are too small, I waste scheduler overhead. I also set shuffle partitions so total tasks  $\geq 2\text{--}3 \times$  total executor cores; this keeps all cores busy.

2. File and partition layout

If inputs are many tiny files, tasks start slowly and stages underutilize the cluster. I compact inputs to 128–512 MB files. I read only needed columns and date partitions to reduce scanned bytes.

3. Adaptive query execution and skew

I enable adaptive query execution so Spark adjusts shuffle partitions and handles skew. If I see a few slow tasks, I check for skewed keys and add salting or pre-aggregation.

4. Node types and disks

Even on memory-optimized nodes, slow local disk or network can bottleneck shuffles. For heavy shuffle, I prefer instances with fast NVMe (i3/i4i) or high-throughput EBS and point Spark's temp/shuffle to those disks. I also verify network bandwidth is adequate.

5. Storage path

I keep intermediates on HDFS or local disks, not S3. Writing shuffle or temp data to S3 via EMRFS slows everything.

6. Cluster sizing and autoscaling

I check YARN pending vCores/memory. If there is backlog, I let managed scaling add task nodes. If pending is zero but CPUs are idle, I reduce cores per executor and raise shuffle partitions to improve concurrency.

7. Joins and broadcast

I confirm small tables actually broadcast. If both sides are large, I sort-merge with good partitioning and ensure both sides are filtered early.

8. Driver pressure

If the driver is GC'ing or running out of memory, I avoid collect/show on big data, bump driver memory modestly, or split the job.

In short, I right-size executors, increase useful parallelism, fix file sizes, enable AQE, use fast local storage for shuffle, and scale nodes only after the software-side bottlenecks are removed.

**14. You need to load processed data into Redshift. Would you use Spark or Hadoop, and how would you integrate it?**

I would use Spark for the transform step and Redshift COPY for the actual load, because COPY is the fastest and most reliable way to bulk-load Redshift.

1. Write curated data to S3  
Spark writes the processed dataset to S3 in Parquet (or compressed CSV if required), coalesced into 100–512 MB files for efficient COPY.
2. Load pattern  
I run a COPY into a staging table in Redshift using an IAM role attached to Redshift. After COPY, I merge from staging into the final table to avoid duplicates. If the target is append-only by date, I can delete that date's slice first and then insert.
3. Orchestration  
I trigger Spark on EMR (or EMR Serverless), then call Redshift COPY using the Redshift Data API or a Lambda step. For idempotency, each run writes to a unique S3 path and places a success marker. The loader reads only that path.
4. Performance settings  
I use Parquet for fewer bytes and faster ingest, gzip for CSV if needed, and set dist/sort keys for the final table. I keep reasonable file counts (not thousands of tiny files). Auto analyze/vacuum helps, but I still analyze changed tables after large loads.
5. When I might use JDBC from Spark  
Only for very small trickle updates. For bulk loads, JDBC is slower than COPY and risks long transactions.

So the integration is: Spark → S3 curated data → Redshift COPY to staging → merge into final → analyze. This gives speed, reliability, and simpler recovery.

### 15. For real-time fraud detection, would you choose Spark Streaming or Hadoop MapReduce? Why?

I would choose Spark Structured Streaming because it is built for low-latency, continuous processing, while MapReduce is batch and too slow for real-time decisions.

How I would design it in simple terms:

1. Sources  
Ingest events from Kinesis or Kafka into Spark Structured Streaming.
2. Processing  
Do windowed aggregations, joins with reference data, and feature engineering in-memory. Use watermarks to handle late events and keep state bounded. If a small reference table is needed, broadcast it.
3. Scoring  
Load a trained model and score events in the stream. For simple rules, apply them inline; for ML models, use a UDF or a model-serving endpoint.
4. Outputs  
Write alerts to a fast store or queue (for example DynamoDB, Elasticsearch, or a notifications topic) and also write a durable trail to S3 for audits and retraining.
5. Reliability  
Enable exactly-once sinks where supported, checkpoint to S3, and use unique output paths so replays do not duplicate results.

MapReduce would buffer data, run in batches, and add many minutes of latency, which is unacceptable for fraud detection. Spark Streaming gives sub-minute latency, stateful processing, and simpler code for windows and joins, so it is the right fit.

**16. In a Spark + Hive workflow, how would you allow Spark steps to continue if a Hive step fails?**

I would design the workflow so Hive is “optional” and Spark can self-heal or fall back. My approach is simple:

1. Make Hive steps non-blocking at the orchestrator  
In EMR step configs, set the Hive step’s failure action to CONTINUE, or if I use Step Functions, add a Retry with a small backoff and then a Catch that routes the flow forward. This way, a Hive failure won’t stop later Spark steps.
2. Add existence checks and fallbacks in Spark  
Before Spark uses Hive objects, it checks the catalog. If a table or partition is missing (because the Hive step failed), Spark creates it on the fly with CREATE TABLE IF NOT EXISTS or ALTER TABLE ADD PARTITION. This lets Spark proceed even when Hive DDL failed.
3. Prefer a shared catalog and use Spark SQL instead of Hive when possible  
If I use the Glue Data Catalog as the metastore, both Hive and Spark see the same metadata. Many Hive DDL tasks (create external table, add partitions) can be done by Spark SQL directly, so I reduce dependency on a separate Hive step.
4. Use idempotent, unique output paths and success markers  
Every step writes to a run-specific S3 path and drops a small \_SUCCESS marker. Downstream Spark steps check for the marker first; if the Hive step failed but the data is already available from a previous run, Spark just continues.
5. Keep business logic in Spark; keep Hive to light DDL only  
If Hive is only doing DDL and Spark owns the transformations, Spark can recover by creating missing objects. If Hive does transforms, I mirror the logic in Spark as a fallback.

In short, I make Hive non-blocking at orchestration level, let Spark create missing metadata on demand, and rely on markers so downstream steps know when to continue safely.

**17. A nightly EMR job often fails at step 5 due to S3 errors. How would you design retries or skipping?**

I would add smart retries with backoff, make the step idempotent, and allow safe skipping when yesterday's slice is already complete.

1. Retries with exponential backoff at orchestration  
If I use Step Functions, I wrap step 5 with a Retry policy (for example, maxAttempts 3–5, exponentialBackoff). On EMR steps, I can re-submit the same step from the orchestrator when it fails. This absorbs transient S3 issues like throttling or slow-downs.
2. Idempotency via unique paths and markers  
Step 5 writes to a date-stamped S3 prefix (for example, dt=YYYY-MM-DD) and creates a `_SUCCESS` marker after finishing. On retry, the code first checks for `_SUCCESS`; if present, it skips the write. This prevents duplicates and makes retries safe.
3. Use S3-optimized committers and right-sized files  
I enable the S3 optimized committer so the job does not rely on slow rename patterns, and I coalesce outputs to 128–512 MB files. This reduces the chance of partial writes and too many small PUTs that trigger throttling.
4. Safe “skip and continue” logic  
If step 5 is non-critical for downstream steps, I add a Catch path in Step Functions that logs the failure, sends an alert, and continues. If it is critical, I continue only when yesterday's partition already exists with a `_SUCCESS` marker. Otherwise, the workflow stops.
5. Application-level resilience  
In Spark, I tune S3 client retries and timeouts, avoid listing huge prefixes repeatedly, and ensure the job reads from an immutable input path (no overlapping writers). If the S3 error was permissions-related, the run will fail deterministically and retries won't help, so I surface the exact error and stop.
6. Pre-flight checks  
Before step 5 starts, I verify IAM/KMS access to the target bucket, and test a tiny write to the exact prefix. This catches configuration issues early and avoids burning the full run time just to fail at the end.

With these changes, transient S3 issues are retried automatically, duplicates are avoided through markers and unique paths, and the pipeline can safely skip or proceed based on whether the required partition is already complete.

## 18. How would you restructure EMR workflows to run independent jobs in parallel instead of sequentially?

I start by identifying which steps are truly independent. I draw a simple dependency graph: only keep arrows where one step really needs the output of another. Everything else can run at the same time.

Then I change the orchestrator to run branches in parallel:

- If I use Step Functions, I put independent steps inside a Parallel state. Each branch runs its own EMR step or transient EMR cluster.
- If I use Airflow, I set task dependencies so parallel tasks have the same upstream and different downstream, and let the scheduler run them simultaneously.
- If I use EMR managed workflows, I split the single long pipeline into multiple smaller workflows kicked off at the same time.

For resources, I avoid one cluster becoming a bottleneck:

- Option 1: Multiple transient clusters, one per branch, so they don't compete for YARN resources.
- Option 2: One large cluster with YARN queues. I give each branch a queue with guaranteed capacity so a heavy job doesn't starve others.

For data paths, each branch writes to its own date-stamped S3 prefix and drops a `_SUCCESS` marker. A final "merge" step depends only on those markers, not on time, so it runs as soon as all branches finish.

For safety, I add retries and timeouts per branch, and I publish per-branch metrics (row counts, durations) so I can see which branch is slow. This setup keeps fast jobs fast, slow jobs don't delay others, and the final dependency only waits for what it truly needs.

## 19. A Spark step fails silently. How would you debug using YARN logs, Spark UI, or S3 logs?

I follow a quick, repeatable path:

1. Get the application ID  
From the EMR step console or CloudWatch logs, I note the Spark application ID (starts with application\_...). This is the key to all logs.
2. Check YARN first  
Using the EMR console or CLI, I open YARN application/attempt logs. I download:
  - Driver stdout/stderr to see exceptions and configuration
  - Executor logs for the last failed stage  
If the driver died early, the reason is usually in driver stderr. If only some tasks died, I look at executor logs for stack traces.
3. Open the Spark History Server  
I load the app in the History Server to see stages, failed tasks, and error messages. I look for:
  - Which stage failed
  - The exact task error and the input file/partition
  - Signs of data skew, OOM, or KMS/S3 access errors
4. S3 interaction clues  
If the job reads/writes S3, I check:
  - The step's CloudWatch logs for AccessDenied, throttling, or KMS errors
  - S3 server-side access logs (if enabled) to confirm GET/PUT failures and which keys were involved
  - Whether the output prefix has partial files but no \_SUCCESS marker
5. Narrow the blast radius  
I re-run only the failing stage inputs: filter by date/partition or LIMIT a small sample to reproduce the error quickly. If one file is corrupt, I isolate it and confirm with a simple select.
6. Common silent-failure causes and fixes
  - Driver OOM → avoid collect/show, modestly raise driver memory, or split the job
  - Executor OOM → increase partitions, reduce cores per executor, or fix skew
  - Permissions/KMS → validate IAM/KMS on the exact bucket/key, test a tiny write first
  - Malformed data → use try\_cast, validate schemas, or quarantine the bad files
  - Spot losses → diversify instance types and add On-Demand buffer

This method always gives me either a stack trace from logs or a failing input key I can act on, turning a silent failure into a clear root cause.

## 20. Would you use Step Functions or Airflow for centralized orchestration of EMR workflows, and why?

I choose based on who operates the platform and what level of control is needed.

When I pick Step Functions:

- I want a fully managed service with no scheduler servers to maintain.
- Most tasks are AWS-native (EMR, Athena, Glue, Lambda, Redshift). Service integrations are simple, with built-in retries, backoff, and error handling.
- I need tight IAM integration, easy per-step permissions, and event-driven starts with EventBridge.
- I prefer visual state machines and JSON/YAML definitions that infra teams can version with IaC.

Trade-offs: complex branching can get verbose; advanced scheduling and cross-environment promotion patterns are simpler in Airflow if the team already uses it.

When I pick Airflow:

- The team already has Airflow skills and wants Python DAGs with rich custom logic.
- We orchestrate both AWS and non-AWS systems (Databricks, Snowflake, on-prem) and need many providers.
- We need advanced scheduling, SLAs, sensors, and backfills across many DAGs.
- We want a central platform with plugins, macros, Jinja-templated SQL, and rich lineage/observability add-ons.

Trade-offs: someone must operate it (or pay for a managed Airflow). Permissions and retries are flexible but require more coding than Step Functions' built-ins.

Simple rule I use:

- Mostly AWS services, want zero ops, straightforward retries and parallelism → Step Functions.
- Heterogeneous ecosystem, heavy Python logic, complex backfills and cross-system orchestration → Airflow.

Both can coexist: Step Functions can orchestrate specific AWS pipelines, while Airflow triggers those state machines and handles enterprise-wide scheduling and dependencies.



## 21. How would you support both Spark 2.x and Spark 3.x workloads when migrating jobs from on-prem to EMR?

I would run Spark 2 and Spark 3 side-by-side, each on its own EMR release, and migrate in small batches with clear guardrails so nothing breaks.

My simple plan:

1. Separate runtimes  
Keep two clusters or two transient workflows: one pinned to a Spark 2.x EMR release for legacy jobs, another to Spark 3.x for new and migrated jobs. This avoids library and Scala version conflicts.
2. One shared catalog, isolated data paths  
Use the Glue Data Catalog so both clusters see the same table metadata, but write outputs to versioned S3 prefixes (for example, s3://.../curated\_v2 and s3://.../curated\_v3). This lets me compare results and roll back safely.
3. Compatibility checklist before moving a job
  - Libraries and Scala/Python versions used by the job
  - Behavior changes in SQL (ANSI mode, null handling, case sensitivity)
  - Config differences (broadcast threshold, shuffle partitions, dynamic partition pruning, adaptive query execution)
  - UDFs and serialization (Kryo vs Java), and any timezone/timestamp settings
  - Connectors and committers for S3 writes
4. Migrate by canary  
Pick one representative job, run it on Spark 3 against the same inputs, and compare row counts, aggregates, and a sample diff. If it matches and performance improves, move the rest of that job family.
5. Guardrails in code  
Parameterize runtime-specific settings so the same code can run on both versions by switching a config file or job parameter. Avoid hardcoding spark.sql.\* flags.
6. Orchestration  
Use Step Functions or Airflow to route each job to the right cluster version. Keep per-version success markers in S3 so downstream steps only consume from the intended version.
7. Decommission plan  
Once a domain is fully validated on Spark 3, freeze its Spark 2 job, switch consumers to the v3 S3 prefix, and remove the old path after a retention window.

In short, I keep both versions alive during migration, control blast radius with separate clusters and S3 paths, validate with canaries, and move workloads in waves until Spark 2 can be retired.

## 22. A Spark ML job keeps running out of memory. How would you choose instance types and tune executor configs?

I first cut memory pressure in the algorithm and data flow, then size the cluster to match, instead of only throwing bigger machines at it.

What I change in the job:

1. Limit data held per task  
Increase partitions so each task handles less data. For wide stages, enable adaptive query execution so Spark splits heavy partitions automatically.
2. Use the right feature representation  
Keep vectors sparse if most features are zeros. Avoid exploding one-hot encodings. Do feature selection or dimensionality reduction before training.
3. Cache only what is reused  
Persist training data as MEMORY\_AND\_DISK, not memory-only. Unpersist intermediates immediately after use. Add checkpoints for very deep pipelines to shorten lineage.
4. Tame model hyperparameters  
For tree models, lower maxDepth and numTrees; for GBT, reduce maxBins and step size. For ALS, use moderate rank and regularization. Use TrainValidationSplit instead of huge cross-validation grids.
5. Avoid large collect operations  
Do not collect feature matrices or models to the driver unless they are truly small. Save samples to S3 when you need inspection.

How I size instances and executors:

1. Pick memory-optimized workers  
Choose r5/r6i (or similar) for high RAM per vCPU. If the job spills a lot, prefer i3/i4i with fast NVMe; otherwise attach EBS with enough throughput.
2. Keep executors medium-sized  
Use about 2–4 cores per executor so each heap serves fewer concurrent tasks. Give generous memory and overhead to each executor. This reduces GC pressure for ML workloads.
3. Set memory and overhead cleanly  
Allocate executor memory plus a larger overhead than default, especially for Python ML and heavy shuffles. For example, if executor memory is 16G, set overhead around 2–4G.
4. Driver sizing  
Increase driver memory modestly and avoid collecting large data. If the driver still hits OOM, break the pipeline into stages and checkpoint between them.
5. Parallelism and file layout  
Aim for enough partitions to keep all executor cores busy (often 2–3× total cores). Compact inputs to 128–512 MB files so partitions are even.
6. Spill paths on fast disks  
Point spark.local.dir to NVMe or high-throughput EBS so spills are cheap when they happen.

A simple starting template that works well for ML:

- Instance family: r5/r6i for memory, or i3/i4i if spill heavy
- 2–4 cores per executor
- Executor memory sized high with 10–20% overhead
- Adaptive query execution on, and partitions  $\geq 2 \times$  total executor cores
- Persist training data as MEMORY\_AND\_DISK and unpersist aggressively

This combination reduces per-task footprint, stabilizes GC, and turns executor OOMs into short, recoverable spills rather than job failures.

### **23. A data scientist needs pandas/scikit-learn on EMR nodes. How would you install these reliably across clusters?**

I make sure the libraries are installed on every node that can run executors, with pinned versions, and that Spark uses the same Python on driver and executors.

Simple, reliable approach:

1. Put a bootstrap script in S3  
The script installs a fixed Python environment on every node at launch. I usually install Miniconda, create an environment with exact versions, and write a small marker file so retries are idempotent.
  - Create requirements.txt with exact versions (for example: pandas==2.2.2, scikit-learn==1.5.1, numpy==1.26.4, scipy==1.13.1).
  - In the script: install OS build tools and BLAS/LAPACK, install Miniconda, create env, pip install from requirements.txt.
  - Save wheels in S3 if the cluster has no internet or to speed up installs.
2. Point Spark to this Python  
Set environment variables so both driver and executors use the conda Python:
  - spark-env: export PYSPARK\_PYTHON=/opt/conda/envs/ds/bin/python
  - spark-env: export PYSPARK\_DRIVER\_PYTHON=/opt/conda/envs/ds/bin/python
3. Cover autoscaled nodes  
Bootstrap actions run for nodes that join later too, so new task nodes get the same environment automatically.
4. Pin and test once  
Keep the script and requirements.txt in S3 under versioned paths. Test on a small cluster, then reuse for every environment.
5. Faster alternative for large fleets  
If bootstrap time matters, bake an EMR custom image with the conda env preinstalled. Then every new node comes with pandas and scikit-learn already present.

This gives reproducible installs, correct versions, and the same Python across driver and executors.

## 24. How would you configure EMR clusters to automatically push Spark/Hadoop logs to CloudWatch?

I push logs to CloudWatch with an agent and a small amount of configuration, so I can search errors in one place.

Step by step:

### 1. IAM permissions

The EMR EC2 instance profile needs CloudWatch Logs permissions: `CreateLogGroup`, `CreateLogStream`, `PutLogEvents`, `PutRetentionPolicy`.

### 2. Agent config in S3

Create a CloudWatch agent config file in S3 that tails common EMR logs, for example:

- `/var/log/spark/, /var/log/hadoop/, /var/log/hadoop-yarn/*`
- `/mnt/var/log/hadoop/steps/, /var/log/livy/, /var/log/hive/*`  
Use a log group naming pattern like `/emr/cluster-${cluster_id}/${component}` and set a retention period.

### 3. Bootstrap action to install and start the agent

In a bootstrap script, install the CloudWatch agent on every node, download the config from S3, substitute cluster id and region, start the agent, and enable it on reboot.

### 4. Spark and YARN settings for clearer logs

- Set Spark log level and rolling: `spark.executor.logs.rolling.enabled=true`, with a reasonable size and max files.
- Ensure Spark History Server writes event logs to S3 (`spark.eventLog.enabled=true`, `spark.history.fs.logDirectory=s3://.../spark-events`) so you have both CloudWatch stream logs and the history UI.
- Keep EMR's logUri to S3 for long-term retention; CloudWatch is for live viewing and alerts.

### 5. Alarms and dashboards

Create CloudWatch metric filters for "ERROR" on key log groups and attach alarms to notify on-call. Optionally build a dashboard with recent error counts per cluster.

With this setup, every node ships Spark driver and executor logs, Hadoop and YARN logs, and step logs to CloudWatch automatically, making failures easy to search and alert on.

## 25. How would you use Step Functions + Lambda to run transient EMR clusters nightly instead of long-running ones?

I would build a small state machine that creates a fresh EMR cluster each night, runs the steps, and then terminates the cluster to avoid idle cost.

Simple flow:

1. Schedule  
Use EventBridge to trigger the Step Functions state machine every night at a fixed time.
2. Create cluster  
First state calls a Lambda that uses the EMR API to create a cluster with the exact EMR release, instance types, bootstrap scripts, security, and tags. I keep core nodes On-Demand and task nodes on Spot for cost. The cluster is created with auto-terminate disabled at this point.
3. Submit steps  
Next state submits one or more EMR steps (Spark or Hive) using the same cluster id. Each step writes outputs to a unique, date-stamped S3 path and drops a success marker to make retries idempotent.
4. Wait and monitor  
A loop state polls step status using the EMR API or the Data API until steps finish. If a step fails, the state machine retries with exponential backoff. If retries are exhausted, it catches the error, sends an alert via SNS, and proceeds to cleanup.
5. Terminate cluster  
Final state terminates the cluster regardless of success or failure (finally block). This guarantees no idle cost.
6. Parameters and versioning  
All inputs like processing date, S3 paths, EMR release, and Spark configs are parameters to the state machine, so I can promote the same definition across dev, test, and prod.
7. Observability and guardrails  
CloudWatch logs capture step output, and I set cost tags on the cluster. I also cap the maximum cluster size, so even a bad config cannot explode cost.

This gives me zero idle time, clean isolation per run, and simple retries with a predictable nightly bill.

## 26. Spark jobs write results to S3, but Athena queries show inconsistent data. How would you fix EMRFS + S3 consistency?

I would make the write path atomic and make reads target only completed data, so Athena never sees partial or old files.

What I change:

1. Write to a temporary path, then commit  
Each Spark job writes to a unique temp prefix and uses the S3 optimized committer to avoid slow rename patterns. After success, I promote by writing a small success marker (`_SUCCESS`) and optionally updating a pointer location.
2. Use run-specific, date-stamped paths  
Never overwrite the same prefix in place. For example, write to `s3://.../curated/dt=YYYY-MM-DD/runId=...` and then point consumers to the latest run for that date. This avoids readers seeing old and new files together.
3. Readers only look at completed data  
Downstream Athena tables or views point to a stable prefix that is updated only after the writer finishes. Readers check for the `_SUCCESS` marker and ignore paths without it.
4. Avoid concurrent writers to the same prefix  
If multiple jobs write the same partition, coordinate through a lock or write to separate run prefixes and merge later. This removes overlap that causes inconsistent reads.
5. Compact files and avoid tiny objects  
Coalesce output to 128–512 MB files. Fewer, larger files reduce S3 metadata operations and stale listings.
6. Catalog updates and partition management  
If using partitioned tables, add the new partition after the write completes, or use partition projection so Athena does not list half-written locations.
7. Permissions and encryption  
Ensure IAM and KMS permissions are correct so failed partial writes don't leave unreadable objects that confuse queries.

In simple words: write to a temp, unique path, commit with a success marker, point Athena only at finished data, and avoid overlapping writers. This removes inconsistent reads and makes results predictable.

## 27. How would you integrate EMR with Glue Data Catalog for schema management across accounts?

I use Lake Formation cross-account sharing so every EMR cluster, in any account, reads the same schemas with fine-grained permissions.

Simple plan:

1. Central catalog account  
Keep all Glue databases and tables in one “producer” account. Point them at S3 paths that are also centrally owned.
2. Share through Lake Formation  
In the producer account, grant Lake Formation permissions on selected databases/tables to the consumer accounts. This creates resource links in the consumer accounts, so they “see” the same tables without copying schemas.
3. Give EMR the right identity  
In each consumer account, the EMR cluster’s IAM role must be added in Lake Formation permissions (database/table level; column/row filters if needed). Least privilege: SELECT on the curated tables is usually enough for read jobs.
4. Point EMR to Glue as the metastore  
Launch EMR with Glue Data Catalog enabled (this is default on recent EMR). Now Spark SQL, Hive on EMR, and Presto on EMR read the shared schemas automatically.
5. Control access at multiple layers  
S3 bucket policies allow only the producer and consumer roles. Lake Formation enforces table/column permissions. Athena/EMR workgroups use separate result buckets to avoid leaks. Encrypt data with KMS and grant the same roles to the key.
6. Keep schemas in code  
Store table DDL (or Glue Crawlers’ configs) in IaC so changes are reviewed and promoted. When a schema evolves, update once in the producer account; consumers get it immediately via the resource links.

This setup avoids schema drift, keeps permissions consistent, and removes the need to duplicate crawlers or DDL in every account.

## 28. How would you efficiently push Spark-aggregated results from EMR into Redshift?

I let Spark do the heavy transform, then use Redshift COPY from S3 for fast, cheap bulk load.

Step by step:

1. Write curated output to S3  
From Spark, write results in Parquet (preferred) to a unique, date-stamped S3 path. Coalesce to 100–512 MB files so there aren't thousands of tiny files. Drop a `_SUCCESS` marker.
2. Load via COPY into a staging table  
Use the Redshift Data API (or Lambda) to run COPY from that S3 path into a staging table. Attach an IAM role to Redshift with read access to the S3 bucket and KMS key. Parquet makes COPY very fast and avoids CSV parsing issues.
3. Merge into target tables  
Run an upsert pattern: either MERGE (if enabled) or delete the target slice for that date/business key, then insert from staging. Keep transactions short and commit once.
4. Optimize for throughput  
Match file count to cluster slices (a few files per slice). Use automatic compression/encoding or ANALYZE after large loads. Choose good dist/sort keys (for example, sort by date, dist by a high-cardinality key used in joins).
5. Idempotency and retries  
Each run writes to a unique S3 prefix. The loader checks `_SUCCESS` before COPY. If a retry happens, it loads the same prefix again without duplicates because the merge pattern is deterministic.
6. When to consider the Spark-Redshift connector  
Only for small trickle writes or interactive notebooks. For daily or hourly bulk, S3 + COPY is simpler, faster, and more reliable.

In short: Spark → Parquet in S3 (right-sized files) → Redshift COPY to staging → merge to final → analyze. It's fast, cheap, and easy to operate at scale.



## 29. How would you enrich Spark streaming data on EMR with lookups from DynamoDB?

I would pick a lookup pattern based on the size and freshness needs of the reference data, and keep the streaming job backpressure-safe.

1. If the reference table is small to medium and changes infrequently  
Load it once per job (or on a schedule) and broadcast it to executors.
  - At startup or every N minutes, read the latest snapshot of the reference data (either via DynamoDB export to S3 or a batch read) into a DataFrame.
  - Convert to a key→values map and broadcast it.
  - In Structured Streaming, use mapPartitions to enrich events by key from the broadcast map.
  - To refresh, swap in a new broadcast variable when a checkpointed timer hits; old broadcasts get unpersisted.
2. If the table is medium/large or must be near-real-time fresh  
Do batched, rate-limited lookups to DynamoDB in the stream.
  - In foreachBatch (preferred) or mapPartitions, collect distinct keys per micro-batch, call BatchGetItem in chunks (e.g., 100 keys per request), and join the results back to the micro-batch DataFrame.
  - Add a small in-executor cache (TTL map) to absorb hot keys and reduce RCU usage.
  - Apply rate limiting and retries with exponential backoff to avoid throttling. Monitor RCUs and set on-demand or provisioned capacity with auto scaling.
3. If the reference is very large or joins are heavy  
Maintain a continuously updated snapshot in S3 and do a stream-static join.
  - Keep DynamoDB → S3 snapshots current using DynamoDB Streams + Lambda (or Glue/DMS) writing Parquet by key.
  - In each micro-batch (foreachBatch), read only the needed partition(s) from S3 and join with the streaming batch. Partition the snapshot by a natural key prefix to prune reads.
4. Operational guardrails
  - Idempotency: write enriched outputs to date/run-scoped S3 paths with a success marker so replays don't duplicate.
  - Checkpointing: enable checkpointLocation in S3; if a micro-batch fails mid-lookup, Spark reprocesses safely.
  - Backpressure: prefer foreachBatch so you can control external I/O; cap concurrent DynamoDB calls.
  - Schema drift: validate reference schema and default missing fields to avoid null pointer issues.
  - Security: grant the EMR role least-privilege to DynamoDB and KMS (if tables are encrypted).

A simple, reliable default is: export DynamoDB to S3 regularly, broadcast when small, and use foreachBatch + BatchGetItem with caching and throttling when you need fresher data.

### 30. How would you structure EMR outputs in S3 so analysts can run fast Athena queries?

I would organize S3 as a lake with clear zones, partitioning, and file sizes that minimize scanned data.

1. Zones and naming
  - raw/ (as-is), curated/ (cleaned, typed), and analytics/ (pre-aggregated).
  - Use predictable, lowercase, snake\_case table and column names.
  - Keep one dataset per top-level prefix, not many small scattered folders.
2. Columnar formats and compression
  - Write Parquet (or ORC) with snappy compression for all curated and analytics tables.
  - Select only needed columns upstream to keep Parquet files lean.
3. Partitioning that matches query filters
  - Use Hive-style keys: dataset/col1=val1/col2=val2/... (for example, dt=YYYY-MM-DD/region=US).
  - Partition first by date/time (most common filter), then by a low-cardinality dimension like region or source.
  - Avoid over-partitioning; too many tiny partitions slow planning. If partitions explode, use partition projection in the table definition.
4. Right-sized files and compaction
  - Target 128–512 MB per Parquet file. Too many small files hurt Athena performance.
  - At the end of each job, coalesce/repartition before writing. Schedule periodic compaction for late or small partitions.
5. Stable, atomic writes for consistency
  - Write to a temporary run path, then promote by creating a \_SUCCESS marker or updating a view/pointer.
  - Never overwrite a partition in place while readers are active; use new run prefixes and swap.
6. Glue Data Catalog hygiene
  - Define schemas as code (DDL/IaC) or run narrow crawlers only on curated/analytics, not raw.
  - Prefer explicit DDL with partition projection for large tables to avoid expensive MSCK/repair scans.
7. Access patterns and helpers
  - Create “presentation” tables or views with only analyst-needed columns and pre-aggregations to cut scan cost.
  - Keep a lightweight audit table (row counts, min/max dates) so analysts sanity-check before heavy queries.

## 8. Security and isolation

- Separate S3 buckets or prefixes for results, with SSE-KMS and per-team IAM/Lake Formation permissions
- Use dedicated Athena workgroups with query limits and their own result buckets.

In short, Parquet + sensible partitioning, large files, atomic writes, and clean catalog definitions make Athena queries fast and cheap, while zones and views keep things easy for analysts.

## 31. How would you design auto-scaling for EMR so that jobs don't fail during peaks but clusters aren't idle all day?

I keep a small, steady base and let the cluster grow only when there is real queue pressure. I also make scale-in gentle so running jobs are not hurt.

Simple plan:

1. Set a safe minimum and a hard maximum  
Keep a few core nodes On-Demand as the steady base (for HDFS/quorum and small daytime jobs). Allow a larger max by adding task nodes when needed. This prevents idle cost yet gives room for peaks.
2. Use EMR Managed Scaling (or EMR Auto Scaling) on real signals  
Scale out when YARN pending vCores or pending memory stay high for a few minutes. Scale in when overall YARN utilization stays low for a while. This reacts to actual demand, not just time.
3. Separate roles and pricing  
Core nodes on On-Demand for stability. Task nodes mostly on Spot for cheap burst capacity. Diversify Spot across 2–3 instance types and sizes with capacity-optimized allocation so losses are rare.
4. Protect running work when scaling in  
Enable graceful decommission so EMR drains tasks off nodes before removing them. Set longer scale-in cooldowns than scale-out, so the cluster shrinks slowly and doesn't kill performance mid-stage.
5. Right-size executors so added nodes actually help  
Use about 4–5 cores per executor with enough memory overhead. Increase shuffle partitions so total tasks  $\geq$  total executor cores; otherwise extra nodes sit idle.
6. Pre-warm before known peaks  
If you know the daily spike time, add a scheduled "nudge" 10–15 minutes earlier. This avoids a backlog while autoscaling catches up.
7. Keep HDFS small and use fast local disks for shuffle  
Rely on S3 for durable data. Use i3/i4i or adequate EBS for fast spill. Smaller HDFS means you can scale cores down without risking data loss.
8. Observability and guardrails  
Alert on sustained pending vCores, failed Spot replacements, and long stage times. Cap the maximum nodes to prevent runaway cost. Tag clusters for cost tracking.

In short, keep a small On-Demand core, burst with diversified Spot task nodes based on YARN pressure, drain slowly on the way down, and ensure executors/partitions are tuned so extra nodes translate into real parallelism.

### 32. Spark ML jobs need high memory for short bursts. How would you pick instance types and pricing (Spot, On-Demand, Reserved)?

I match the workload shape: big memory for a short window, minimal cost outside that window, and resilience to interruptions.

Choosing instances:

1. Memory first  
Pick memory-optimized families (r5/r6i) when the model mostly needs heap. If the job spills a lot during feature engineering, pick i3/i4i for fast NVMe. For Python-heavy ML, give extra executor overhead.
2. Executor template  
Use small-to-medium executors (2–4 cores each) with generous memory and 10–20% overhead. More, smaller executors lower GC pressure compared to a few huge ones.
3. Storage for spills  
Point spark.local.dir to NVMe (i3/i4i) or high-throughput EBS. Spills are then short and do not crash tasks.

Picking pricing:

1. Baseline on On-Demand or Reserved/Savings  
Keep primaries and a minimal set of cores On-Demand. If this base runs every day, cover it with Reserved Instances or Savings Plans to reduce cost.
2. Burst on Spot  
Add most of the task nodes on Spot for the ML training window. Use instance fleets with 2–3 alternative types and capacity-optimized allocation. Set a reasonable Spot cap and allow EMR to replace lost nodes. Checkpoint models and use retry-friendly stages so Spot interruptions don't lose progress.
3. Transient clusters for short runs  
For nightly or ad-hoc training, spin up a transient cluster, run the job, then terminate. You pay only for the burst. Bake libraries via bootstrap or a custom image so startup is quick.
4. Guardrails  
Set max cluster size and timeouts so experiments don't overshoot. Use Step Functions or Airflow retries around training steps. Store checkpoints and intermediate datasets in S3 so retries resume fast.

A simple default that works well: r6i or i4i workers, 2–4 cores per executor with high memory and overhead, primaries On-Demand (covered by Savings Plans if steady), most task capacity on Spot via diversified fleets, transient clusters per training run, and checkpoints to S3 for safe retries. This gives high memory when you need it and low cost when you don't.

### 33. How would you reduce EMR costs when running separate clusters for dev, test, and prod?

I keep prod reliable but squeeze costs in dev/test with a few simple rules.

1. Use transient clusters for dev/test so they start for a job and auto-terminate after steps finish. Keep prod long-running only if needed.
2. Turn on managed scaling. Set a small On-Demand core base and let cheap Spot task nodes handle bursts (use 2–3 instance types per fleet for better Spot capacity).
3. Right-size executors and files so extra nodes actually help. Many small executors (2–5 cores each) and 128–512 MB files keep parallelism high and runtime short.
4. Prefer memory- or storage-optimized instances only when the workload needs it. For most ETL, balanced families work; switch to i3/i4i only for heavy shuffle.
5. Share non-prod when possible. One multi-tenant dev cluster with YARN queues is cheaper than many tiny clusters; isolate teams by queues, not clusters.
6. Use Savings Plans or Reserved Instances to cover the steady prod base. Keep dev/test mostly on Spot.
7. Bake a custom image (AMI) or use bootstrap caching so clusters start fast; less paid warm-up time.
8. Store all durable data in S3 (EMRFS). Keep HDFS small so you can scale down safely and never pay to keep clusters up “just for data.”
9. Enforce tags and budgets. Require owner, env, and ttl tags; add Cost Explorer and Budgets alerts per env so surprises are caught early.
10. Schedule non-prod quiet hours. EventBridge can shut down nightly or on weekends if clusters are still up.

In short, make dev/test transient and Spot-heavy, keep prod’s steady base covered by discounts, right-size everything, and use tagging plus budgets to keep spend in check.

### 34. Data scientists forget to terminate ad-hoc clusters. How would you auto-detect and shut down idle clusters?

I combine prevention, detection, and automatic shutdown.

1. Prevent idle costs at creation
  - Default to transient clusters with auto-terminate after the last step.
  - For interactive work, set an idle termination policy (idle timeout) so EMR auto-terminates when no YARN apps run for N minutes.
  - Require ttl and owner tags via a launch template or policy; clusters without a ttl are rejected.
2. Detect idleness automatically
  - Create a small Lambda that runs every 10–15 minutes (EventBridge). It lists clusters and checks signals like no running steps, low CPU, zero YARN pending/running apps for longer than a threshold, and past-due ttl.
  - Optionally read a “do\_not\_kill\_until” tag so training jobs can extend their window.
3. Shut down safely
  - If idle threshold is exceeded, the Lambda sends a notification (Slack/SNS). If still idle after a grace period, it calls `TerminateCluster`.
  - Enable graceful decommission so nodes drain tasks before termination.
4. Make idleness visible
  - Push key metrics (running apps, pending vCores, CPU) to a dashboard per cluster.
  - Set CloudWatch alarms: “no running apps for 60–120 minutes” triggers an alert or auto-termination.
5. Reduce the chance of false positives
  - Exclude clusters with active steps or tags like “scheduled\_maintenance.”
  - For studios/notebooks, kill only after idle timeout and last user disconnect.

With idle timeouts, a janitor Lambda, and strict tagging, ad-hoc clusters shut themselves down when unused, and you stop paying for forgotten capacity.

### 35. ETL jobs hit out-of-memory errors on compute-optimized instances. How would you reconfigure cluster sizing?

I would treat this as a mismatch between workload and hardware plus overly large executors. I would switch to memory-friendly nodes and reshape executors so each task handles less data.

What I change first

- Move workers from compute-optimized (like c5/c6i) to memory-optimized (r5/r6i) or storage-optimized with fast NVMe (i3/i4i) if shuffle spill is heavy. These give larger RAM per vCPU and faster local disks.
- Reduce cores per executor to 2–4 so each heap serves fewer concurrent tasks. This lowers GC pressure and avoids big spikes.
- Increase executor memory and set overhead to at least 10–20 percent. Make sure executor memory + overhead fits in YARN container size.
- Increase partitions so each task processes a smaller slice. A simple start is total shuffle partitions around 2–3 times total executor cores, then let adaptive query execution refine it.

Layout on the node

- Leave about one core for the OS/YARN. On an r6i.4xlarge (16 vCPU, 128 GB), use three or four executors with 3–4 cores each. Give each executor generous memory and overhead, keeping some gap for the OS.
- Point `spark.local.dir` to fast local disks (NVMe on i3/i4i or high-throughput EBS). Keep shuffle and temp on local, not S3.

Data and query tweaks that lower memory

- Compact tiny input files into 128–512 MB. Read only needed columns and filter early.
- Broadcast only truly small tables. If there is key skew, add salting or rely on adaptive skew handling to split the few giant partitions.
- Cache only data reused multiple times and unpersist as soon as done. Add checkpoints to cut very long lineage.

Cluster shape and cost

- Keep core nodes On-Demand for stability, add task nodes on Spot for burst. Memory-optimized Spot often costs less than pushing c5 nodes to the limit.
- Enable managed scaling so more task nodes join when YARN shows pending memory or vCores.

A simple, reliable template

- Instance family r6i or i4i, 2–4 cores per executor, generous executor memory with 10–20 percent overhead, adaptive query execution on, partitions at least 2× total executor cores, shuffle on local NVMe. This usually turns OOM failures into fast, controlled spills and finishes the job.

### **36. Ad-hoc ETL jobs take too long to set up clusters. When would you recommend EMR Serverless instead?**

I recommend EMR Serverless when jobs are short-lived, sporadic, and do not need custom OS-level setup or HDFS. It removes cluster startup and management, so you pay only for compute while the job runs.

#### **Good fit**

- On-demand Spark or Hive jobs that run a few minutes to a couple of hours, with unpredictable schedules.
- Teams that want zero cluster ops, fast cold starts, and automatic scale-out/in per job.
- Workloads that read/write S3 and do not require HDFS, SSH access, or complex daemons on the nodes.
- Multi-tenant ad-hoc analytics where you want per-application quotas and spend limits.

#### **How I set it up**

- Create an EMR Serverless application for Spark. Define pre-initialized capacity if you want near-instant starts for peak hours.
- Package dependencies in job artifacts or use runtime configuration to pull libraries at start. Keep them light to avoid cold-start drag.
- Submit jobs via API, console, or orchestrators like Step Functions or Airflow. Pass a unique S3 output path per run and write a success marker.
- Set maximum capacity and concurrency to cap spend. Use per-job tags for cost allocation. Monitor with the built-in metrics and logs.

#### **When I would not use it**

- You need custom native binaries, long-running services, or SSH access to tune the OS.
- Heavy reliance on HDFS or engines beyond Spark/Hive.
- Very tight network constraints that require complex VPC setups not supported by serverless defaults.

#### **Simple rule**

- If the pain is cluster wait time and idle cost for ad-hoc Spark or nightly short ETL, EMR Serverless is a strong choice. If you need deep host customization, HDFS, or many non-Spark components, stick with EMR on EC2 but speed it up with custom images and autoscaling.



**37. An ML team wants Spark on S3 without cluster tuning. How would EMR Serverless help, and what are its limits?**

EMR Serverless lets the team run Spark jobs directly on S3 data without managing clusters. There is no need to choose instance types, set executor sizes, or worry about scaling. I just create an EMR Serverless Spark application, submit jobs with the script and dependencies, and EMR Serverless automatically provisions and scales compute up and down. This removes cold starts from full cluster boot (especially if I enable pre-initialized capacity during peak hours), and I only pay for the compute used while the job runs. It integrates well with S3, Glue Data Catalog, and IAM, so the team can focus on code and data, not infrastructure. It also supports per-application limits, concurrency controls, and tagging for cost governance, which helps with multi-tenant ML teams.

Its limits are mainly around deep customization and certain performance patterns. I do not get SSH access to nodes, so I cannot do low-level OS tuning, custom daemons, or complex system packages that require root changes. HDFS is not the storage model; it is designed to read and write to S3, so extremely shuffle-heavy workloads may be slower than using EC2 clusters with fast local NVMe, unless the code and partitioning are carefully optimized. Startup latency is lower than spinning a cluster, but still exists if the application is cold and needs to pull many dependencies; keeping dependencies lean or using pre-initialized capacity helps. Some niche networking or VPC topologies can be more constrained than EMR on EC2. If I need engines or components outside Spark/Hive, or very specific hardware (like large NVMe fleets), Serverless may not fit.

In short, EMR Serverless removes cluster tuning and scaling, is great for Spark on S3 with minimal ops, but it is not ideal when I need deep host customization, HDFS-centric workloads, or the absolute best shuffle performance from local NVMe.

**38. For quarterly Hive reporting, would you choose EMR Serverless or a provisioned EMR cluster? Why?**

I would choose EMR Serverless for quarterly reporting because the workload is infrequent and predictable. With Serverless, I do not keep a cluster running for months between runs. I submit the Hive job on schedule, it scales up to finish, and then scales down to zero, so I pay only for the hours I use each quarter. It is simpler to operate, easier to secure with IAM and Glue Catalog, and I avoid cluster patching or right-sizing work. If the reports run for a few hours and mostly scan S3 data, Serverless fits perfectly.

I would choose a provisioned EMR cluster only if the quarterly report needs heavy, iterative tuning with custom system packages, relies on HDFS for speed, or requires multiple interactive sessions over several days where having a warm cluster with NVMe disks gives a big performance advantage. For most quarterly batch Hive jobs that read Parquet from S3 and write results back to S3, EMR Serverless is the simpler and more cost-effective option.

### **39. Spark jobs finish in 10 minutes but clusters run longer. How would EMR Serverless reduce costs?**

With EMR on EC2, you pay from the moment the cluster starts until it terminates, including idle time before and after jobs. If your Spark job runs 10 minutes but the cluster takes 8–12 minutes to bootstrap and then sits around waiting for other steps or manual termination, you are paying for a lot of non-use time.

EMR Serverless removes most of that waste. You create a serverless Spark application once and just submit jobs. Compute spins up automatically, runs the 10-minute job, and scales back to zero when the job ends. You are billed only for the compute used while the job is executing (plus brief provisioning), not for an always-on cluster. If you enable pre-initialized capacity only during known busy windows, you can keep start latency low without paying outside that window. Because there is no HDFS to keep alive, there is no incentive to keep the environment running; all data is on S3. In short, EMR Serverless matches cost to runtime, so those extra idle minutes essentially disappear from the bill.

### **40. When is EMR Serverless a better choice than Glue for Spark ETL, and when is Glue better?**

Choose EMR Serverless when:

- You want Spark with minimal infrastructure but still need full Spark control, custom Spark configs, and flexibility with runtimes and dependencies.
- You have bursty or ad-hoc jobs and care about quick job submission without managing clusters.
- You need features common in EMR's Spark ecosystem, like certain libraries, fine-grained Spark tuning, or using the Glue Data Catalog without adopting Glue Job abstractions.
- You prefer to orchestrate with Step Functions, Airflow, or your own tooling and just need a serverless Spark execution engine.

Choose AWS Glue (Jobs) when:

- You want a managed ETL service with built-in job authoring (Glue Studio), workflows, bookmarks for incremental reads, and tight integration with crawlers and the Data Catalog.
- Your team prefers minimal Spark tuning and likes out-of-the-box connectors, transforms, and job bookmarking with less code.
- You need simple scheduling, retries, and lineage-friendly patterns without building your own orchestration.
- You have many small to medium ETL jobs where Glue's per-DPU pricing and autoscaling are easy to manage.

A simple rule I use: if you need "Spark the way Spark behaves on EMR," with custom configs and the EMR runtime, go with EMR Serverless. If you want "managed ETL with wizards, crawlers, bookmarks, and minimal Spark tuning," go with Glue. Both read and write S3 and can use the Glue Data Catalog; the choice is about how much Spark control you want versus how much ETL convenience you want.