

# **AZURE FUNCTION SCENARIO BASED Q&A**

**BY - SHUBHAM WADEKAR**

Copyright © Shubham Wadekar. All rights reserved.

This material is for personal use only. No part may be copied, shared, resold, or published without written permission. Unauthorized distribution is strictly prohibited and may result in action.

For permissions or licensing, contact: [shubham.p.wadekar@gmail.com]

Disclaimer: This content is for educational purposes. Accuracy is attempted but not guaranteed. No job outcome is promised.

## 1. How would you handle cold starts in Azure Functions?

**Scenario: Users report slow responses for the first request after inactivity in an HTTP-triggered function.**

When users report slow responses for the first request, it usually means the function app is facing cold start issues. This happens mostly in the Consumption Plan, where Azure stops the function app after some idle time to save resources. When a new request comes in, Azure needs to start the app again, which causes a delay.

To solve this, I would use one of the following approaches:

1. **Switch to the Premium Plan:** In the Premium Plan, Azure keeps one or more instances always warm, even during idle time. This removes the cold start issue completely. It is the most effective way if performance is critical.
2. **Use a timer-triggered warm-up function:** I can create a simple timer-triggered function that runs every few minutes and calls the HTTP endpoint. This keeps the function app active and reduces cold start. However, this is more like a workaround and may not be reliable under all conditions.
3. **Optimize startup code:** I make sure there is no heavy logic or large dependency loading in the startup path of the function. This helps reduce the impact of cold start even if it happens.

So, in production scenarios where user experience is important, I prefer the Premium Plan to make sure there is no delay for the first user request.

## 2. How would you secure sensitive configuration values in Azure Functions?

**Scenario: You need to protect API keys and DB credentials used in your function.**

To protect sensitive values like API keys, database passwords, or storage account keys, I avoid hardcoding them in the function code. Instead, I use the following secure methods:

1. **Use Application Settings in Azure Portal:** I store the secrets as environment variables under Configuration > Application Settings in the Function App. Azure automatically injects these values into the runtime environment. In my code, I read them using environment variables like this in C#:

```
string connectionString = Environment.GetEnvironmentVariable("MyDbConnectionString");
```

2. **Use Azure Key Vault:** For stronger security and centralized secret management, I store the secrets in Azure Key Vault. I link the Function App to the Key Vault using a managed identity. This way, the function can securely fetch secrets without needing credentials in the code. Here's a simple flow:

- Enable managed identity for the Function App.
- Give the identity permission to access secrets in Key Vault.
- Use Key Vault references in application settings like:

```
@Microsoft.KeyVault(SecretUri=https://mykeyvault.vault.azure.net/secrets/MyDbPassword/)
```

3. **Restrict access to settings:** I ensure only trusted users or systems can access the function configuration using role-based access control (RBAC) and least privilege principles.

By following these steps, I keep all sensitive information outside the code and make the function secure and compliant with best practices.

### 3. How would you handle a function that fails due to a dependency timeout?

**Scenario: The function times out when waiting for a slow external API.**

When an Azure Function fails due to a timeout while calling an external API, it usually means the API is taking longer than expected to respond. Here's how I would handle this:

#### 1. Increase the timeout setting:

In the **Consumption Plan**, the maximum timeout is 5 minutes by default and can be extended to 10 minutes using the host.json file:

```
{  
  "functionTimeout": "00:10:00"  
}
```

In the **Premium Plan**, I can set a longer timeout or even unlimited by changing the same setting:

```
{  
  "functionTimeout": "-1"  
}
```

#### 2. Use retry policies:

I would add retry logic either in the code or through bindings. For example, if the function calls an HTTP API, I use a library like Polly in C# to add retries with exponential backoff:

```
var retryPolicy = Policy  
    .Handle<HttpRequestException>()  
    .WaitAndRetryAsync(3, retryAttempt => TimeSpan.FromSeconds(Math.Pow(2,  
retryAttempt)));
```

#### 3. Move the call to a Durable Function:

If the external call can take a very long time and I don't want the main function to timeout, I use Durable Functions. The orchestrator can wait for the result and automatically manage state and retries without worrying about timeout limits.

#### 4. Async call and queue-based processing:

If the API does not need to respond immediately, I offload the call to a queue. The function pushes a message to a queue and another function picks it up and handles the API call separately. This way, the main function is quick and doesn't timeout.

#### 5. Monitor and log slow dependencies:

I log API response times and set up alerts using Application Insights to detect repeated slowness and take further actions like alerting the vendor or changing service-level agreements.

These steps help ensure the function remains reliable even if the external service is slow or temporarily unavailable.

#### 4. How would you implement dependency injection in Azure Functions?

**Scenario:** You want to inject shared services like logging and DB access across multiple functions.

##### 1. Create shared services

I define reusable service classes for logging or database access.

```
# services/logger.py
```

```
class MyLogger:
```

```
    def log_info(self, message):
```

```
        print(f"[INFO] {message}")
```

```
# services/database.py
```

```
class MyDatabaseService:
```

```
    def get_data(self):
```

```
        return "Sample data from DB"
```

##### 2. Initialize services globally

I create a global file like startup.py to initialize and share instances.

```
# startup.py
```

```
from services.logger import MyLogger
```

```
from services.database import MyDatabaseService
```

```
my_logger = MyLogger()
```

```
my_database_service = MyDatabaseService()
```

##### 3. Use the services in your function by importing them

```
# MyFunction/__init__.py
```

```
import datetime
```

```
import azure.functions as func
```

```
from startup import my_logger, my_database_service
```

```
def main(mytimer: func.TimerRequest) -> None:
```

```
    current_time = datetime.datetime.utcnow()
```

```
    my_logger.log_info(f"Function ran at {current_time}")
```

```
    data = my_database_service.get_data()
```

```
    my_logger.log_info(f"Retrieved from DB: {data}")
```

**Summary:**

- Python Azure Functions don't support constructor-based injection.
- I simulate dependency injection by creating services and importing shared instances globally.
- This keeps the code organized and allows service reuse across multiple functions.

For more advanced scenarios, I can create a ServiceContainer or use libraries like dependency-injector, but for most practical cases in Azure Functions, this simple approach works well.

© Shubham Wadekar

## 5. How would you handle concurrency in Azure Functions?

**Scenario: Your function modifies Azure Table Storage but suffers from data inconsistencies due to simultaneous executions.**

When multiple Azure Function instances try to update the same row in Azure Table Storage at the same time, it can cause data conflicts like overwrites. In Python, I handle this using a few strategies:

### 1. Use ETag for optimistic concurrency

Azure Table Storage supports ETags, which act like version numbers. When I fetch an entity, it comes with an ETag. While updating, I include that ETag to ensure no other instance has modified it.

If someone else has updated the row, the ETag will have changed, and my update will fail. I can then retry safely.

Here's a simplified example using the Azure Data Tables SDK:

```
from azure.data.tables import TableServiceClient, UpdateMode
```

```
from azure.core.exceptions import ResourceModifiedError
```

```
connection_string = "<your_connection_string>"
```

```
table_name = "MyTable"
```

```
service = TableServiceClient.from_connection_string(conn_str=connection_string)
```

```
table_client = service.get_table_client(table_name=table_name)
```

```
try:
```

```
    # Read the existing entity
```

```
    entity = table_client.get_entity(partition_key="Users", row_key="123")
```

```
    original_etag = entity.metadata['etag']
```

```
    # Modify the entity
```

```
    entity['score'] = 90
```

```
    # Update with ETag to ensure concurrency control
```

```
    table_client.update_entity(mode=UpdateMode.REPLACE, entity=entity, etag=original_etag,
                              match_condition="IfMatch")
```

except ResourceModifiedError:

```
# Handle concurrency conflict
```

```
print("ETag mismatch: Entity was modified by another instance. Retrying...")
```

```
# Retry logic goes here
```

## 2. Serialize updates using queues

Another way I solve this is by sending update messages to a queue instead of letting all function instances write to Table Storage directly.

- I let multiple clients put update requests into Azure Queue Storage.
- I create a queue-triggered Azure Function that reads and processes each message one at a time.
- This serializes the writes and avoids concurrent access altogether.

This pattern is especially helpful when multiple updates are expected on the same partition or row.

## 3. Use a proper partitioning strategy

If concurrency is high, I also review my table design. I try to spread updates across different partition keys, so that rows are less likely to be written at the same time.

For example, instead of having all users in a single partition, I might partition by region or user type.

## 4. Limit parallelism in host.json

If I want to avoid too many function instances working in parallel, I can configure the host.json to limit concurrency.

```
{
  "extensions": {
    "queues": {
      "batchSize": 1,
      "newBatchThreshold": 1
    }
  }
}
```

This ensures only one message is processed at a time, which is safer for high-conflict operations like writing to the same table row.

By combining ETag version control, queue-based serialization, partitioning, and parallelism tuning, I can safely handle concurrency in Azure Table Storage using Python Azure Functions and avoid any data inconsistencies.

## 6. How would you monitor and troubleshoot an Azure Function in production?

**Scenario: Your production function fails intermittently without clear errors.**

When a production function fails randomly and doesn't show clear error messages, I follow a step-by-step process to monitor and troubleshoot the issue:

### 1. **Enable and check Application Insights:**

I always connect the Function App to Application Insights. It captures logs, exceptions, custom events, and performance data. In this case, I check the **Failures** tab to see detailed stack traces, request details, and telemetry around the failed execution.

### 2. **Use structured logging:**

In my function code, I use ILogger to log key information like function start, inputs, outputs, and catch block messages. This helps trace the flow even if the error is random:

```
_logger.LogInformation("Starting function with input: " + inputValue);
```

```
_logger.LogError("Error occurred: " + ex.Message);
```

### 3. **Use custom telemetry:**

I add custom logs or metrics using Application Insights SDK. For example, if there's a logic branch that might fail silently, I log it:

```
telemetryClient.TrackEvent("Null value encountered during processing");
```

### 4. **Check retry and timeout settings:**

Sometimes a function fails due to timeouts or retries. I review retry policies in host.json or binding configuration to make sure I'm not hitting retry loops or function time limits.

### 5. **Use Live Metrics:**

In Application Insights, I use the Live Metrics Stream to watch function executions and failures in real time. It helps spot patterns quickly.

### 6. **Check environment dependencies:**

If the function calls external APIs or depends on other services like storage or databases, I verify their availability and error responses. I also check network and permission issues like firewalls or expired credentials.

### 7. **Capture and analyze exception types:**

I look at exception types and stack traces to identify what kind of error is occurring. If the error is unhandled, I update the function to add better try-catch blocks.

By using a combination of structured logs, Application Insights, retry configuration, and custom telemetry, I can find the root cause of the failure and improve the function's stability in production.



## 7. How would you optimize an Azure Function processing millions of messages from a queue?

**Scenario: The queue function is slow and delayed at scale.**

When an Azure Function is processing millions of messages and becomes slow, I take several steps to improve its performance and throughput:

### 1. **Increase batch size and concurrency settings:**

Azure Queue and Service Bus triggers allow tuning through `host.json`. I increase the batch size and control how many messages are processed in parallel:

```
{
  "extensions": {
    "queues": {
      "batchSize": 32,
      "newBatchThreshold": 16,
      "maxDequeueCount": 5
    }
  }
}
```

This helps the function fetch and process more messages in fewer executions.

### 2. **Enable parallel execution inside the function:**

If the logic allows, I process messages in parallel using asynchronous calls (like `Task.WhenAll`) inside the function. This is useful when making multiple I/O-bound operations like API calls or database writes.

### 3. **Scale the function using the Premium Plan:**

The Premium Plan supports more powerful instances, unlimited execution time, and faster scaling. This helps handle large volumes of messages more efficiently compared to the Consumption Plan.

### 4. **Use separate queues for different priorities or workloads:**

If the queue contains mixed types of messages, I separate them into different queues. Each one can have its own function and scaling behavior, which prevents bottlenecks.

### 5. **Ensure function is idempotent:**

At large scale, retries may happen due to transient failures. I design the function to handle duplicate messages gracefully, especially when writing to a database.

### 6. **Monitor and adjust poison message handling:**

If some messages repeatedly fail, they may delay the rest. I configure the maximum dequeue count and use a dead-letter queue to remove problematic messages from the pipeline.

### 7. **Review storage and dependencies:**

I check if the function is slow due to dependency issues like database latency or blob access. Optimizing these external systems also improves overall throughput.

By tuning host settings, improving parallelism, and upgrading the hosting plan, I can optimize the queue-triggered function to handle millions of messages efficiently.

## 8. How would you implement durable workflows using Azure Durable Functions?

**Scenario:** You need to execute multiple steps where each depends on the success of the previous one.

In Azure Durable Functions using Python, I handle such step-by-step workflows by using the orchestrator pattern. This pattern allows me to define a workflow where each step (called an activity function) runs in sequence, and the next one only starts if the previous one finishes successfully.

Here's how I implement it in Python:

### 1. Define the activity functions

These functions do the actual work and return the result. For example:

#### activity1.py

```
import azure.functions as func
```

```
def main(name: str) -> str:
```

```
    return f"Step 1 completed for {name}"
```

#### activity2.py

```
def main(data: str) -> str:
```

```
    return f"Step 2 received: {data}"
```

### 2. Define the orchestrator function

This function controls the workflow. It calls activity functions one after another and passes the output of one step into the next.

#### orchestrator.py

```
import azure.durable_functions as df
```

```
def orchestrator_function(context: df.DurableOrchestrationContext):
```

```
    name = context.get_input()
```

```
    step1_result = yield context.call_activity('activity1', name)
```

```
    step2_result = yield context.call_activity('activity2', step1_result)
```

```
    return step2_result
```

```
main = df.Orchestrator.create(orchestrator_function)
```

### 3. Define the HTTP starter function

This function lets me trigger the durable workflow using an HTTP request.

#### **starter.py**

```
import azure.functions as func
```

```
import azure.durable_functions as df
```

```
async def main(req: func.HttpRequest, starter: str) -> func.HttpResponse:
```

```
    client = df.DurableOrchestrationClient(starter)
```

```
    name = req.params.get('name')
```

```
    if not name:
```

```
        return func.HttpResponse("Please pass a name", status_code=400)
```

```
    instance_id = await client.start_new('orchestrator_function', None, name)
```

```
    return client.create_check_status_response(req, instance_id)
```

#### **How it works:**

- When I send an HTTP request to the starter function with a name (like "John"), it starts the orchestrator.
- The orchestrator runs activity1, waits for it to finish, then runs activity2 using its output.
- If any activity fails, the orchestrator pauses and retries based on built-in settings.
- Azure Durable Functions automatically handles state, checkpoints, and retries behind the scenes.

#### **Why this works well for this scenario:**

- Each step only starts if the previous one succeeds.
- If something fails, it can be retried or compensated.
- The function is durable, so it can survive restarts or long waits.

This pattern is ideal for orchestrating ETL processes, approvals, data enrichment, or any scenario where order and success dependency between steps is important.

## 9. How would you deploy an Azure Function using CI/CD?

**Scenario: You want automated deployment from GitHub or Azure DevOps to your Function App.**

To automate deployment of an Azure Function using CI/CD, I follow a structured process using either GitHub Actions or Azure DevOps, depending on the team's setup.

### Using GitHub Actions:

1. I connect my GitHub repository to Azure using a publish profile or a service principal.
2. I create a workflow file in `.github/workflows/deploy.yml`. This file defines steps to build and deploy the function.
3. In the workflow, I include steps to:
  - Check out the code
  - Set up the language environment (like .NET, Node.js, or Python)
  - Build the function project
  - Use `azure/functions-action` to deploy the code to Azure

Here's an example for a Node.js Azure Function:

name: Deploy Function App

on:

push:

branches:

- main

jobs:

build-and-deploy:

runs-on: ubuntu-latest

steps:

- uses: actions/checkout@v2

- uses: azure/functions-action@v1

with:

app-name: 'my-function-app'

package: '.'

publish-profile: \${ secrets.AZURE\_FUNCTIONAPP\_PUBLISH\_PROFILE }

### Using Azure DevOps:

1. I create a pipeline in Azure DevOps with two stages: Build and Release.
2. In the Build stage:
  - I use tasks like dotnet build, npm install, or pip install depending on the language.
  - Then, I publish the build artifacts (zipped output) to be used in the release stage.
3. In the Release stage:
  - I use the Azure Function App Deploy task.
  - I select the target Function App and provide the artifact path.
  - I also store secrets like connection strings or credentials in the Azure DevOps Library or Key Vault.

With this setup, every time I push to the main branch, the pipeline builds the function and deploys it automatically to Azure, ensuring consistent and quick deployments without manual steps.

## 10. How would you optimize cost for high-volume Azure Functions?

**Scenario: The function scales heavily, causing a spike in consumption costs.**

When a function is handling a very high volume of requests or messages, and the cost increases suddenly in the Consumption Plan, I take the following steps to reduce cost without affecting performance:

1. **Switch to the Premium Plan for predictable pricing:**

The Premium Plan gives better scaling control and fixed instance pricing. If the workload is consistent or high for long periods, Premium can be cheaper than paying per execution.

2. **Batch the processing:**

Instead of processing one message at a time, I increase the batch size (especially for queue or event triggers). This reduces the number of executions:

```
{  
  "extensions": {  
    "queues": {  
      "batchSize": 32,  
      "newBatchThreshold": 16  
    }  
  }  
}
```

3. **Reduce execution duration:**

I profile the function and remove unnecessary steps or dependencies that slow it down. The faster the function completes, the less I pay, since billing is based on execution time and memory used.

4. **Avoid unnecessary triggers:**

I make sure the function is not triggered too frequently without real work to do. For example, I reduce polling frequency for timer triggers or add conditions to check if processing is needed.

5. **Use durable functions only when needed:**

Durable Functions can be more expensive due to storage and orchestration overhead. I use them only for workflows that truly need state management.

6. **Move non-critical tasks to background:**

I push tasks that are not urgent (like sending a notification or logging analytics) to queues and process them during off-peak times using low-priority function apps.

7. **Monitor and analyze cost trends:**

I use Azure Cost Management and Application Insights to monitor execution counts, average duration, and scale-out behavior. This helps me spot unnecessary executions and optimize further.

By batching messages, reducing execution time, and selecting the right hosting plan, I can keep Azure Function costs under control even during high-volume operations.

## 11. Can Azure Functions run offline?

Yes, Azure Functions can run offline for local development and testing. This is very helpful when I want to develop and debug the function on my local machine before deploying it to Azure.

To run Azure Functions offline, I follow these steps:

1. **Install Azure Functions Core Tools:**

I install the Azure Functions Core Tools on my local machine. This provides a runtime environment that simulates Azure Functions locally.

2. **Set up a local project:**

I create a function project using Visual Studio Code or the command line. The project includes files like `host.json`, `local.settings.json`, and function-specific code.

3. **Use `local.settings.json` for config:**

This file stores app settings like connection strings, API keys, and environment variables needed during local execution. These settings are not uploaded to Azure when deploying, so they are safe for development use.

Example:

```
{
  "IsEncrypted": false,
  "Values": {
    "AzureWebJobsStorage": "UseDevelopmentStorage=true",
    "FUNCTIONS_WORKER_RUNTIME": "dotnet"
  }
}
```

4. **Run the function locally:**

I use the command `func start` in the terminal. This starts the function runtime, and I can test triggers like HTTP, Timer, or Queue locally using tools like Postman or curl.

5. **Use local emulators if needed:**

For services like Azure Storage, Cosmos DB, or Azure Service Bus, I can use local emulators during testing to avoid connecting to live cloud resources.

So, while Azure Functions are meant to run in the cloud, I can fully develop, test, and debug them offline using the local runtime. This speeds up development and allows me to catch errors before deploying to Azure.

## 12. How do you handle exceptions in Azure Functions?

In Python-based Azure Functions, I handle exceptions using standard try-except blocks. This allows me to gracefully catch errors, log them properly, return meaningful responses, and ensure failed executions are visible through monitoring tools like Application Insights.

Here are the main ways I manage exceptions:

### 1. Using try-except blocks inside the function

For example, in an HTTP-triggered function:

```
import azure.functions as func
```

```
import logging
```

```
def main(req: func.HttpRequest) -> func.HttpResponse:
```

```
    try:
```

```
        name = req.params.get('name')
```

```
        if not name:
```

```
            raise ValueError("Missing required parameter: name")
```

```
        # Simulate processing
```

```
        greeting = f"Hello, {name}"
```

```
        return func.HttpResponse(greeting, status_code=200)
```

```
    except ValueError as ve:
```

```
        logging.error(f"Validation error: {str(ve)}")
```

```
        return func.HttpResponse(str(ve), status_code=400)
```

```
    except Exception as e:
```

```
        logging.exception("Unhandled exception occurred")
```

```
        return func.HttpResponse("Internal Server Error", status_code=500)
```

- This function checks for a required input and handles both validation and unexpected errors.
- It returns proper HTTP status codes and logs the details.



## 2. Logging exceptions for monitoring

I always use `logging.exception()` inside the `except` block for unexpected errors. This logs the full stack trace to Application Insights if monitoring is enabled:

```
import logging

try:

    # risky operation

    result = 10 / 0

except Exception as e:

    logging.exception("Error while processing request")
```

## 3. Custom error handling in queue or blob triggered functions

In non-HTTP functions, like queue or blob triggers, I also use `try-except` to ensure the function does not crash silently:

```
def main(msg: func.QueueMessage):

    try:

        data = msg.get_body().decode('utf-8')

        # process data

    except Exception as e:

        logging.exception("Failed to process message")

        # optionally move the message to a poison queue or dead-letter it
```

## 4. Poison Queue handling

If a queue-triggered function fails several times (by default, 5), the message can be moved to a poison queue. This lets me investigate problematic messages separately. I use this with manual logging and alerts.

## 5. Retry policies (if applicable)

In `host.json`, I can also configure automatic retries for transient errors (for example, in queue-triggered functions):

```
"retry": {

    "strategy": "fixedDelay",

    "maxRetryCount": 3,

    "delayInterval": "00:00:10"

}
```

**Summary:**

- I wrap risky code in try-except blocks.
- I log all exceptions using logging.exception for full traceability.
- I return proper status codes in HTTP responses.
- I use poison queues and retry settings for automated error recovery.

This structured approach makes my Azure Functions more reliable, easier to debug, and production-safe

© Shubham Wadekar

### 13. How do you manage state in Azure Functions?

Azure Functions by default are stateless. This means they don't remember anything between executions. But if I need to manage or preserve state, I use one of the following options:

1. **Azure Durable Functions:**

This is the best choice for managing state across multiple steps or long-running workflows. Durable Functions automatically store the state between executions using Azure Storage. I can use orchestrator functions to manage the workflow and keep track of each step.

2. **External storage:**

If I want to store state myself, I use services like Azure Table Storage, Cosmos DB, or Blob Storage. For example, I save user session data or job progress in a database and read it back in the next function call.

3. **State passing via input/output bindings:**

Sometimes I can pass state in and out using bindings. For example, I can bind a blob or table entity to a function input, update it, and output the changed value to keep state updated externally.

So, although Azure Functions are stateless by design, I can manage state by using Durable Functions or storing it in external systems like databases or storage services.

### 14. How do you configure deployment slots in an Azure Function App?

Deployment slots allow me to create different environments within a single Function App — like staging, testing, and production. Each slot has its own hostname and settings, which makes it safe to test new code before making it live.

Here's how I configure and use deployment slots:

1. **Enable slots:**

Deployment slots are available in Premium and App Service plans, not in the Consumption Plan. So I first make sure the Function App is using the correct plan.

2. **Create a new slot:**

In the Azure Portal, I go to my Function App, click on Deployment slots, and then click Add Slot. I name it (e.g., "staging") and choose whether to clone the settings from the production slot.

3. **Deploy code to the slot:**

I deploy my code (using CI/CD or manually) to the staging slot instead of production. This allows me to test the new version without affecting users.

4. **Use slot-specific settings:**

I set configuration values that should be different between slots, like API keys or connection strings, using slot settings so they stay separate from production.

5. **Swap slots when ready:**

Once testing is complete, I swap the staging slot with production. This switch is fast and avoids downtime. If something goes wrong, I can swap back to roll back the change.

Using deployment slots helps me deploy safer, test better, and avoid breaking production with untested changes.

## 15. How do you configure Function App settings?

To configure Function App settings, I use the Azure Portal, Azure CLI, or Infrastructure as Code tools like ARM templates or Bicep. These settings include environment variables, connection strings, timeouts, and runtime behavior.

Here's how I configure them:

### 1. Using Azure Portal:

- I go to my Function App in the Azure Portal.
- Under the Configuration section, I add Application settings which act like environment variables.
- These include keys like AzureWebJobsStorage, FUNCTIONS\_WORKER\_RUNTIME, database connection strings, or custom values like MyApiKey.
- If I want a setting to be different per deployment slot, I check the Deployment slot setting option.

### 2. Using local.settings.json for local development:

When testing locally, I use local.settings.json:

```
{
  "IsEncrypted": false,
  "Values": {
    "AzureWebJobsStorage": "UseDevelopmentStorage=true",
    "MySetting": "value123"
  }
}
```

These settings don't get deployed to Azure and are only used offline.

### 3. Using Azure CLI:

I can also set settings from the command line:

```
az functionapp config appsettings set --name my-func-app --resource-group my-rg --settings MySetting=value123
```

### 4. Using ARM or Bicep:

For automated deployments, I define settings inside a template and deploy them using CI/CD pipelines.

This approach ensures my configuration is flexible, secure, and consistent across environments.

## 16. How do you implement versioning in Azure Functions?

Azure Functions doesn't have built-in versioning like APIs in API Management, but I can still implement versioning in different ways depending on how the functions are triggered. In Python, I usually follow one of these approaches:

### 1. Version in the Function Name or Folder

The simplest way is to create separate functions for each version. I add the version number to the function name and folder:

#### Folder structure:

MyFunctionApp/

├─ V1/

| └─ \_\_init\_\_.py → function V1

├─ V2/

| └─ \_\_init\_\_.py → function V2

#### Example: function.json in V2:

```
{
  "scriptFile": "__init__.py",
  "entryPoint": "main",
  "bindings": [
    {
      "authLevel": "function",
      "type": "httpTrigger",
      "direction": "in",
      "name": "req",
      "methods": ["get"],
      "route": "v2/myfunction"
    },
    {
      "type": "http",
      "direction": "out",
      "name": "$return"
    }
  ]
}
```

**In `__init__.py`:**

python

return "This is version 2 of the function"

- The route is `/api/v2/myfunction`, which clearly indicates the version.
- I can keep V1 and V2 running side-by-side safely.

## **2. Version in the Route or URL Parameter (for HTTP-triggered functions)**

If I don't want to duplicate the whole function, I can accept a version as a query parameter or part of the route:

def main(req):

version = req.route\_params.get("version")

if version == "v1":

return func.HttpResponse("Processing with version 1")

elif version == "v2":

return func.HttpResponse("Processing with version 2")

else:

return func.HttpResponse("Unsupported version", status\_code=400)

**Route in `function.json`:**

"route": "myfunction/{version}"

This way, a single function handles multiple versions based on input.

## **3. Versioning with Function Apps**

In some cases, I deploy separate Function Apps for different versions like:

- my-function-app-v1
- my-function-app-v2

This helps isolate deployments, configs, and environments completely. It's useful in production if major logic changes are required.

## **4. Version control via source code (Git)**

I also use Git branches or tags for versioning the function logic at the code level. This helps me manage releases even if only one version is deployed live at a time.

**Summary:**

In Python Azure Functions, I implement versioning by:

- Creating separate folders or function names like v1, v2
- Using versioned HTTP routes (/api/v1/...)
- Deploying different Function Apps for each major version
- Managing versioned code using Git or pipelines

This helps me support older clients, safely test new logic, and roll back if needed.

© Shubham Wadekar

## 17. How do you monitor and log Azure Functions?

In Python-based Azure Functions, I use logging, Application Insights, and Azure Monitor to track how my functions behave, troubleshoot issues, and analyze performance. Here's how I monitor and log Azure Functions effectively:

### 1. Use built-in logging in Python Functions

I use the standard logging module inside my functions. The logs automatically go to Application Insights (if it's connected) and can be viewed in the Azure portal.

#### Example:

```
import logging

def main(req):

    logging.info("Function started")

    try:

        name = req.params.get('name')

        if not name:

            raise ValueError("Missing name")

        logging.debug(f"Received name: {name}")

        return f"Hello, {name}"

    except Exception as e:

        logging.exception("Error occurred")

        return "Something went wrong"
```

- logging.info() logs general messages
- logging.debug() logs detailed technical info
- logging.exception() logs error with full stack trace



## 2. Enable Application Insights

When I create the Function App in Azure, I make sure Application Insights is enabled.

This allows me to:

- See real-time logs in Live Metrics Stream
- Track function invocations, duration, failures
- Use Kusto Query Language (KQL) in Log Analytics to run custom queries

### Example KQL query:

traces

| where operation\_Name == "MyHttpFunction"

| sort by timestamp desc

## 3. Use the Azure Portal for Monitoring

Under the Function App in the Azure Portal:

- The "Monitor" tab shows recent executions with status and duration.
- I can click on any execution to see logs, inputs, and outputs.
- If Application Insights is linked, I get deeper analytics and performance tracking.

## 4. Set up alerts

In production, I configure alerts in Azure Monitor. For example:

- Alert if the function fails more than 5 times in 5 minutes
- Alert if average execution time exceeds a threshold
- Alert if no runs happen (indicating trigger issues)

## 5. Log custom metrics or telemetry (optional)

If needed, I can send custom telemetry using opencensus-ext-azure:

```
from opencensus.ext.azure.log_exporter import AzureLogHandler
```

```
logger = logging.getLogger(__name__)
```

```
logger.addHandler(AzureLogHandler(connection_string='InstrumentationKey=<key>'))
```

```
logger.warning("Custom telemetry warning")
```

**Summary:**

To monitor and log Azure Functions in Python, I:

- Use logging module for info, debug, and errors
- Enable Application Insights for centralized telemetry
- View execution logs in the portal's Monitor tab
- Use KQL to analyze logs and failures
- Configure alerts for proactive monitoring

This helps me detect failures, debug issues quickly, and ensure the function is running smoothly in production.

© Shubham Wadekar

## 18. What are the best practices for error handling in Azure Functions?

In Azure Functions, proper error handling is important to make sure failures are caught, logged, and handled properly. Below are some best practices that I follow, using Python:

### 1. Use Try-Except Blocks

I wrap critical logic inside a try-except block to catch runtime errors. This helps in preventing the function from crashing unexpectedly.

```
import logging
```

```
import azure.functions as func
```

```
def main(req: func.HttpRequest) -> func.HttpResponse:
```

```
    try:
```

```
        name = req.params.get('name')
```

```
        if not name:
```

```
            raise ValueError("Name parameter is missing.")
```

```
        return func.HttpResponse(f"Hello, {name}!")
```

```
    except ValueError as ve:
```

```
        logging.error(f"Validation error: {str(ve)}")
```

```
        return func.HttpResponse(f"Bad Request: {str(ve)}", status_code=400)
```

```
    except Exception as e:
```

```
        logging.error(f"Unexpected error: {str(e)}")
```

```
        return func.HttpResponse("Internal Server Error", status_code=500)
```

### 2. Return Proper HTTP Status Codes

I return 400 for bad input, 500 for internal errors, etc., so that clients can understand what went wrong.

### 3. Log Errors Using Azure Monitor Logs

I use the built-in logging module to log all exceptions. These logs can be viewed in Application Insights for troubleshooting.

### 4. Avoid Swallowing Exceptions

Instead of catching an error and doing nothing, I always log it or raise it again after logging.

```
except Exception as e:
```

```
    logging.error(f"Unhandled exception: {str(e)}")
```

```
    raise
```

### 5. Use Custom Exceptions When Needed

For complex applications, I define custom exceptions to handle specific error scenarios cleanly.

```
class InvalidInputError(Exception):
```

```
    pass
```

## 6. **Retry Policies for Transient Failures**

If the function depends on external services like storage or APIs, I handle transient errors using retry logic or Durable Functions retry features.

import time

```
def retry_operation():
```

```
    for i in range(3):
```

```
        try:
```

```
            # simulate external call
```

```
            result = make_external_call()
```

```
            return result
```

```
        except ConnectionError:
```

```
            logging.warning("Retrying after failure...")
```

```
            time.sleep(2)
```

```
    raise Exception("Failed after 3 retries")
```

## 7. **Dead-lettering for Queue Triggers**

For queue-triggered functions, I configure the **poison queue** settings so failed messages don't get lost and can be reviewed later.

In summary, I combine proper exception handling with structured logging, clear response messages, and retry logic to make Azure Functions more reliable and easier to debug.

## 19. How to optimize function execution time?

To reduce how long a function takes to run (and to lower cost if using the Consumption Plan), I use these strategies:

1. **Avoid unnecessary code or logic**

I keep the function logic clean and short. If something is not needed during execution, I move it out or pre-process it elsewhere.

2. **Use async and non-blocking calls**

For external calls like HTTP requests or database access, I use async and await to avoid blocking the thread.

3. **Cache static data**

If I have configuration or data that doesn't change often, I cache it in memory during the function's lifetime

4. **Use connection pooling**

Instead of creating new database or HTTP connections every time, I reuse static clients like HttpClient or database connections. This reduces startup time.

5. **Minimize cold starts**

- I use the Premium Plan or Always On feature (in App Service Plan) to avoid cold starts.
- I reduce the number of dependencies to make cold starts faster.

6. **Parallelize independent operations**

If I can perform multiple tasks at once (like calling two APIs), I use Task.WhenAll to run them in parallel.

7. **Use batching when reading/writing data**

Instead of processing one record at a time, I read and write data in batches to reduce the number of external calls.

8. **Profile and monitor performance**

I use Application Insights to find which parts of the function are slow and optimize only those areas.

9. **Reduce payload size**

I keep the input and output payloads (like HTTP request bodies or queue messages) as small as possible for faster processing.

By combining these techniques, I make sure my function runs faster, handles more load, and costs less in serverless environments.

## 20. How to handle large message processing in Azure Functions?

When Azure Functions need to handle large messages, especially from sources like queues or HTTP requests, I follow these best practices to avoid failures and improve performance:

### 1. Avoid putting large content directly in the message

Instead of sending the full payload in a queue or service bus message, I upload the large content to a storage service (like Azure Blob Storage) and pass only the reference (like blob URL) in the message.

Example:

- Upload the large JSON or file to Blob Storage.
- Send a queue message with the blob URI and metadata.

### 2. Use Blob trigger for large files

If the data is already stored in a blob, I use a Blob Trigger Function. The function will be triggered as soon as the blob is uploaded. This avoids passing large content through queues or HTTP.

### 3. Stream the data instead of loading it all at once

When reading large content (from blob, HTTP, or DB), I stream it to avoid memory overflow:

```
import logging
```

```
import azure.functions as func
```

```
def main(blob: func.InputStream):
```

```
    try:
```

```
        for line in blob.read().decode('utf-8').splitlines():
```

```
            # Process each line
```

```
            logging.info(f"Line: {line}")
```

```
    except Exception as e:
```

```
        logging.error(f"Error while processing blob: {str(e)}")
```

### 4. Use durable functions for long-running or chunked processing

If the large data needs to be processed in steps or takes a long time, I use Durable Functions to split it into smaller activity functions. Each step can handle part of the data and resume if interrupted.

### 5. Set max message size if needed

I ensure the source system respects limits. For example:

- Azure Storage Queue: ~64 KB
- Azure Service Bus: 256 KB (standard), 1 MB (premium)

If messages exceed this, I offload data and pass references.

### 6. Increase timeout and memory limits (if needed)

For large processing, I make sure the function app has enough timeout and memory, especially if on the Premium Plan.

By offloading the data to storage, using references, and processing in chunks, I make sure the function can handle large messages without hitting platform limits.

## 21. How to optimize Event Hub processing for large-scale messages in Azure Functions?

When dealing with large-scale messages from Azure Event Hub in Azure Functions, I focus on both performance and reliability. Below are the best practices I follow using Python:

### 1. Use Event Hub Trigger with Batch Processing Enabled

I configure the Event Hub trigger to receive events in batches, which improves throughput by reducing the number of function invocations.

```
import logging
```

```
import azure.functions as func
```

```
def main(events: List[func.EventHubEvent]):
```

```
    for event in events:
```

```
        message = event.get_body().decode('utf-8')
```

```
        logging.info(f"Processing message: {message}")
```

```
    # Process each message
```

I make sure that in function.json, batch processing is enabled by setting cardinality to "many":

```
{
  "type": "eventHubTrigger",
  "name": "events",
  "direction": "in",
  "eventHubName": "myeventhub",
  "connection": "EventHubConnectionAppSetting",
  "cardinality": "many",
  "consumerGroup": "$Default"
}
```

### 2. Enable Checkpointing

Azure Functions with Event Hub binding automatically supports checkpointing. I ensure my function completes successfully to allow the system to checkpoint, avoiding message reprocessing.

### 3. Avoid Long Processing Time

I keep the processing of each message lightweight. If the logic is complex, I offload it to background systems like Azure Queue or Durable Functions.

#### 4. Handle Failures Gracefully

I wrap message processing inside try-except blocks and log errors instead of crashing the function.

```
def main(events: List[func.EventHubEvent]):  
    for event in events:  
        try:  
            message = event.get_body().decode('utf-8')  
            process_message(message)  
        except Exception as e:  
            logging.error(f"Failed to process message: {str(e)}")
```

#### 5. Scale Out with Dedicated Hosting Plan

For high-throughput scenarios, I prefer a Premium or Dedicated App Service Plan so that Azure Functions can scale out to handle more partitions and messages concurrently.

#### 6. Monitor and Tune Batch Size and Prefetch Count

I use host.json to tune batch size and prefetch count for better performance:

```
{  
  "version": "2.0",  
  "extensions": {  
    "eventHubs": {  
      "batchCheckpointFrequency": 1,  
      "eventProcessorOptions": {  
        "maxBatchSize": 100,  
        "prefetchCount": 300  
      }  
    }  
  }  
}
```

#### 7. Use Logging and Application Insights

I use logging and Application Insights to track message latency, errors, and throughput so I can keep optimizing based on real-time metrics.

By combining batch processing, efficient message handling, proper configuration, and good observability, I ensure my Azure Function can reliably process large-scale messages from Event Hub in production.



## 22. Write an Azure Function that processes messages from an Azure Queue Storage

### Function Code (Python):

```
import logging

import azure.functions as func

def main(msg: func.QueueMessage) -> None:

    try:

        # Read message from the queue

        message_body = msg.get_body().decode('utf-8')

        logging.info(f"Received queue message: {message_body}")

        # Sample processing logic

        processed_message = message_body.upper() # Example: convert to uppercase

        logging.info(f"Processed message: {processed_message}")

    except Exception as e:

        logging.error(f"Error while processing queue message: {str(e)}")
```

### function.json configuration:

```
{
  "scriptFile": "__init__.py",
  "bindings": [
    {
      "name": "msg",
      "type": "queueTrigger",
      "direction": "in",
      "queueName": "myqueue-items",
      "connection": "AzureWebJobsStorage"
    }
  ]
}
```

**How It Works:**

- The function is triggered whenever a new message is added to the Azure Queue named myqueue-items.
- It reads and decodes the message.
- It processes the message (in this case, converts it to uppercase).
- All important steps are logged using logging.info.

© Shubham Wadekar

## 23. Write a Durable Function that orchestrates two other functions and returns the combined result

### 1. Orchestrator Function (orchestrator\_function.py):

```
import azure.durable_functions as df
```

```
def orchestrator_function(context: df.DurableOrchestrationContext):
```

```
    result1 = yield context.call_activity('ActivityFunctionOne', "input for function one")
```

```
    result2 = yield context.call_activity('ActivityFunctionTwo', "input for function two")
```

```
    combined_result = {  
        "functionOneResult": result1,  
        "functionTwoResult": result2  
    }
```

```
    return combined_result
```

```
main = df.Orchestrator.create(orchestrator_function)
```

### 2. Activity Function One (ActivityFunctionOne/init.py):

```
import logging
```

```
def main(name: str) -> str:
```

```
    logging.info(f"ActivityFunctionOne received input: {name}")
```

```
    return f"Result from Function One using input: {name}"
```

### 3. Activity Function Two (ActivityFunctionTwo/init.py):

```
import logging
```

```
def main(name: str) -> str:
```

```
    logging.info(f"ActivityFunctionTwo received input: {name}")
```

```
    return f"Result from Function Two using input: {name}"
```

#### 4. HTTP Starter Function (HttpStart/init.py):

```
import logging

import azure.functions as func

import azure.durable_functions as df

async def main(req: func.HttpRequest, starter: str) -> func.HttpResponse:

    client = df.DurableOrchestrationClient(starter)

    instance_id = await client.start_new("orchestrator_function", None, None)

    logging.info(f"Started orchestration with ID = '{instance_id}'.")

    return client.create_check_status_response(req, instance_id)
```

#### What This Durable Function Does:

- Starts with an HTTP trigger.
- Calls two activity functions (FunctionOne and FunctionTwo) one after the other.
- Waits for both results.
- Combines and returns the results as a single JSON object.

## 24. Write a timer-triggered Azure Function that logs the current date and time to Azure Blob Storage

### 1. Function Code (Python – \_\_init\_\_.py):

```
import datetime

import logging

import azure.functions as func

def main(mytimer: func.TimerRequest, outputBlob: func.Out[str]) -> None:

    current_time = datetime.datetime.utcnow().isoformat()

    logging.info(f"Timer triggered at: {current_time}")

    # Write the timestamp to the blob

    blob_content = f"Function triggered at: {current_time}\n"

    outputBlob.set(blob_content)
```

### 2. function.json configuration:

```
{
  "scriptFile": "__init__.py",
  "bindings": [
    {
      "name": "mytimer",
      "type": "timerTrigger",
      "direction": "in",
      "schedule": "0 */5 * * * *" // every 5 minutes
    },
    {
      "name": "outputBlob",
      "type": "blob",
      "direction": "out",
      "path": "logs/timestamp.txt",
      "connection": "AzureWebJobsStorage"
    }
  ]
}
```

**What This Does:**

- Triggers every 5 minutes (you can change the schedule).
- Gets the current UTC time.
- Logs the time using logging.info.
- Writes the current time to a blob called timestamp.txt inside the logs container.

© Shubham Wadekar

## 25. Write an Azure Function that connects to an Azure Cosmos DB and retrieves a document by its ID

### 1. Function Code (\_\_init\_\_.py):

```
import logging

import azure.functions as func

import os

import json

from azure.cosmos import CosmosClient, exceptions

# Read environment variables

COSMOS_ENDPOINT = os.environ['COSMOS_DB_ENDPOINT']

COSMOS_KEY = os.environ['COSMOS_DB_KEY']

DATABASE_NAME = 'mydatabase'

CONTAINER_NAME = 'mycontainer'

def main(req: func.HttpRequest) -> func.HttpResponse:

    try:

        doc_id = req.params.get('id')

        partition_key = req.params.get('pk')

        if not doc_id or not partition_key:

            return func.HttpResponse("Missing 'id' or 'pk' in query string.", status_code=400)

        client = CosmosClient(COSMOS_ENDPOINT, COSMOS_KEY)

        database = client.get_database_client(DATABASE_NAME)

        container = database.get_container_client(CONTAINER_NAME)

        # Read the document by ID and partition key

        item = container.read_item(item=doc_id, partition_key=partition_key)

        return func.HttpResponse(

            json.dumps(item),

            status_code=200,

            mimetype="application/json"

        )
```

```
except exceptions.CosmosResourceNotFoundError:
```

```
    return func.HttpResponse("Document not found.", status_code=404)
```

```
except Exception as e:
```

```
    logging.error(f"Error: {str(e)}")
```

```
    return func.HttpResponse("Internal Server Error", status_code=500)
```

## 2. function.json Configuration:

```
{  
  "scriptFile": "__init__.py",  
  "bindings": [  
    {  
      "authLevel": "function",  
      "type": "httpTrigger",  
      "direction": "in",  
      "name": "req",  
      "methods": [ "get" ],  
      "route": "getdoc"  
    },  
    {  
      "type": "http",  
      "direction": "out",  
      "name": "$return"  
    }  
  ]  
}
```



### 3. Required Application Settings:

Set the following in local.settings.json or Azure App Settings:

```
{  
  "IsEncrypted": false,  
  "Values": {  
    "AzureWebJobsStorage": "<your-storage-connection-string>",  
    "FUNCTIONS_WORKER_RUNTIME": "python",  
    "COSMOS_DB_ENDPOINT": "https://<your-account>.documents.azure.com:443/",  
    "COSMOS_DB_KEY": "<your-cosmos-db-key>"  
  }  
}
```

#### How It Works:

- This is an HTTP-triggered function.
- It takes id and pk (partition key) from the query string.
- It uses the Cosmos SDK to connect, fetch the document, and return it as a JSON response.

**26. Write an Azure Function that uses an input binding to read data from an Azure Blob and an output binding to write to Azure Table Storage**

**1. Function Code (\_\_init\_\_.py):**

```
import logging

import azure.functions as func

import uuid

def main(inputBlob: func.InputStream, outputTable: func.Out[func.Table]):

    try:

        # Read data from blob

        blob_data = inputBlob.read().decode('utf-8')

        logging.info(f"Read from blob: {blob_data}")

        # Create a row for Table Storage

        row = {

            "PartitionKey": "logpartition",

            "RowKey": str(uuid.uuid4()),

            "BlobContent": blob_data

        }

        # Write to Azure Table

        outputTable.set(row)

        logging.info("Data written to Azure Table Storage.")

    except Exception as e:

        logging.error(f"Error: {str(e)}")
```

## 2. function.json configuration:

```
{
  "scriptFile": "__init__.py",
  "bindings": [
    {
      "name": "inputBlob",
      "type": "blob",
      "direction": "in",
      "path": "input-container/inputfile.txt",
      "connection": "AzureWebJobsStorage"
    },
    {
      "name": "outputTable",
      "type": "table",
      "direction": "out",
      "tableName": "BlobLogTable",
      "connection": "AzureWebJobsStorage"
    }
  ]
}
```

### How It Works:

- The function is triggered when needed (e.g., manual or timer).
- It reads the content of inputfile.txt from input-container.
- Then it writes the blob content into a row in BlobLogTable using a random RowKey.

### Required Setup:

- **Blob Storage** must contain the blob: input-container/inputfile.txt.
- **Azure Table Storage** must have a table named BlobLogTable.
- Both use the same connection: AzureWebJobsStorage (you can change it to another if needed).

**local.settings.json (for testing locally):**

```
{  
  "IsEncrypted": false,  
  "Values": {  
    "AzureWebJobsStorage": "<your-storage-connection-string>",  
    "FUNCTIONS_WORKER_RUNTIME": "python"  
  }  
}
```

© Shubham Wadekar