

AZURE COSMOS DB SCENARIO BASED Q&A

BY - SHUBHAM WADEKAR

Copyright © Shubham Wadekar. All rights reserved.

This material is for personal use only. No part may be copied, shared, resold, or published without written permission. Unauthorized distribution is strictly prohibited and may result in action.

For permissions or licensing, contact: [shubham.p.wadekar@gmail.com]

Disclaimer: This content is for educational purposes. Accuracy is attempted but not guaranteed. No job outcome is promised.

1. What are the typical use cases for Azure Cosmos DB?

Azure Cosmos DB is used in scenarios where we need high availability, low latency, and global scalability. Some common use cases are:

- IoT applications: For example, collecting telemetry data from thousands of devices in real time. Cosmos DB can handle this type of fast and massive data ingestion.
- Real-time recommendation engines: Like in e-commerce websites, where based on user behavior, we show recommendations instantly.
- Personalization systems: For apps like news feeds or video streaming platforms that deliver personalized content to users.
- Gaming leaderboards: Where we need very fast read and write operations and the data must always be available to players around the world.
- Fraud detection systems: Banks or fintech companies can use Cosmos DB to store transactions and detect patterns using fast queries.
- Retail and inventory management systems: Where product availability, pricing, and stock details need to be updated and accessed in real time.
- Content management systems: Apps that need to serve data to millions of users with low response times.

These use cases benefit from Cosmos DB's ability to automatically index data, support for multiple data models (document, graph, key-value), and global distribution.

2. Can you walk me through a sample application that uses Azure Cosmos DB?

Sure. Let's take an example of a food delivery application like Zomato or Swiggy.

In this app, we need to store information about users, restaurants, menus, orders, and delivery status. Here's how Cosmos DB fits in:

- We choose the SQL API of Cosmos DB because the data is in JSON format.
- User profiles are stored as JSON documents with details like name, phone, location, preferences.
- Restaurant data, menus, and item prices are also stored in JSON format.
- When a user places an order, an order document is created with fields like orderId, userId, restaurantId, itemList, totalPrice, and status.
- Real-time order tracking is done by continuously updating the delivery status in the order document.
- Since users and delivery agents can be from different cities, we replicate data in Cosmos DB across multiple regions. This way, data access is very fast regardless of the user's location.
- We use a logical partition key like userId or orderId to distribute the data for better performance and scale.

This application benefits from Cosmos DB because of fast response times, high availability, global replication, and automatic indexing without having to manage infrastructure.

3. How can you provision and scale throughput in Azure Cosmos DB?

In Cosmos DB, throughput is measured in Request Units per second, or RU/s. These RUs cover reads, writes, deletes, and queries. There are two main ways to provision throughput:

1. **Provisioned throughput:** We can assign RU/s to a container or to a database. For example, we can set 1000 RU/s for a container if we expect high read and write traffic. This ensures we always get the performance we need.
2. **Autoscale throughput:** Instead of setting a fixed RU/s, we set a maximum limit (like 4000 RU/s), and Cosmos DB automatically adjusts the RU/s based on traffic. This helps when the load changes a lot, like during daytime traffic spikes.

For scaling, Cosmos DB handles it automatically behind the scenes. If we partition our data properly using a good partition key (like `userId` or `orderId`), Cosmos DB distributes the data across physical partitions, and each partition gets its share of RU/s.

Also, we can scale across regions. By adding regions in the portal, Cosmos DB automatically replicates data and serves requests from the nearest location. This improves performance and ensures high availability even if one region goes down.

In short, we control throughput using RU/s, and scaling is simple as Cosmos DB takes care of the heavy lifting when we configure partition keys and regions correctly.

4. Why do I get throttled when I try to create many tables in the Table API at the same time?

Throttling in Cosmos DB happens when the number of requests exceeds the available throughput, which is measured in RU/s. When using the Table API, every operation, including creating tables, consumes RUs.

If we try to create many tables at once, it puts sudden high pressure on the system. Each `CreateTable` request takes a significant amount of RUs because it's a metadata operation that needs to be committed reliably.

Cosmos DB protects the performance and availability of the system by enforcing limits through throttling. So, if the total RU/s limit is reached due to multiple table creation requests happening at the same time, the service will start returning status code 429 (Request rate is large). This is a sign that we are exceeding the throughput we have provisioned.

To handle this, we can either:

- Create tables one at a time with small delays between requests.
- Increase the provisioned RU/s temporarily before the bulk table creation.
- Use retry logic with exponential backoff so that failed requests can be retried after a short wait.

This throttling is not a failure — it's a signal to slow down or increase capacity. Cosmos DB's SDKs also automatically retry on throttled requests if we use the default retry policy.

5. How does Azure Cosmos DB handle data indexing and querying?

Cosmos DB automatically indexes all the data we store, by default. We don't have to manually define or manage indexes like in traditional databases. This makes query performance very fast and predictable.

Here's how it works:

- Whenever we insert or update a document, Cosmos DB updates the index in the background.
- It supports both range indexes and spatial indexes, depending on the data types (for example, strings, numbers, dates, or geographic data).
- For document types (using SQL API), Cosmos DB uses a B-tree based index structure to allow fast reads on any property.
- Queries are done using SQL-like syntax, and because everything is indexed, the queries return results quickly without needing full scans.

We can also customize the indexing policy in some cases:

- If we don't want to index all fields (to save RUs on write-heavy workloads), we can exclude some paths.
- For write-heavy workloads with no need for queries, we can disable indexing altogether and enable it later when needed.

Cosmos DB provides an indexing metrics tab in the portal that shows how much RU is spent on queries, so we can tune the index policy if required.

In summary, Cosmos DB automatically manages indexes to give high performance during queries, but we can adjust the behavior using indexing policies if our use case demands optimization.

6. How does Azure Cosmos DB handle data consistency and availability?

Azure Cosmos DB gives us five different consistency levels to choose from, so we can balance between data freshness and availability based on our application's needs. These consistency levels are:

1. **Strong** – This guarantees the highest consistency. All reads will return the most recent write. But it might have higher latency and is not available across multiple regions.
2. **Bounded staleness** – This allows reads to be behind writes by a certain number of versions or time (like 5 seconds or 100 versions). It is useful when we want predictable lag.
3. **Session** – This is the default and most common level. It guarantees consistency within a user's session. The same user will always see their own writes, but other users may see slightly older data.
4. **Consistent prefix** – This ensures that data is read in the same order as it was written, but it can be stale.
5. **Eventual** – This gives the highest availability and lowest latency, but there's no guarantee about the order or freshness of the data.

Cosmos DB ensures high availability using automatic multi-region replication. If we enable multiple regions, the data is written to a primary region and replicated to other regions. If one region goes down, Cosmos DB automatically fails over to another region.

So in simple terms, Cosmos DB allows us to choose how fresh we want our data to be, and it automatically keeps our data available by storing copies in different locations.

7. How does Azure Cosmos DB handle data security and compliance?

Cosmos DB provides strong built-in security features to protect data both at rest and in transit.

1. **Encryption:** All data is encrypted at rest using Microsoft-managed keys by default. If needed, we can also use customer-managed keys from Azure Key Vault for more control. Data in transit is protected using TLS (Transport Layer Security), so it cannot be read during transmission.
2. **Access control:** We can control who can access the data using Azure Active Directory (AAD), Role-Based Access Control (RBAC), and resource tokens. Resource tokens are especially useful in client-side applications, where we can generate time-limited tokens to access specific containers or documents.
3. **Network security:** Cosmos DB supports private endpoints and IP firewall rules. This means we can restrict access only to specific networks or addresses, and block public internet access if required.
4. **Auditing and logging:** We can enable diagnostic logs and metrics through Azure Monitor. This helps track access patterns and detect unusual behavior.
5. **Compliance:** Cosmos DB is certified under many compliance standards like ISO, SOC, PCI DSS, HIPAA, and GDPR. This makes it suitable for use in regulated industries like healthcare, finance, and government.

So overall, Cosmos DB ensures that data is secure through encryption, identity control, network protection, and compliance with global standards, giving confidence for both small and enterprise-level applications.

8. How are data backups handled and managed by Azure Cosmos DB?

Azure Cosmos DB handles backups automatically. These backups are taken in the background and are fully managed by Microsoft. We don't have to manually schedule or trigger them.

There are two types of backup options in Cosmos DB:

1. **Periodic backup mode** (default for most accounts):

- Backups are taken automatically every four hours.
- Cosmos DB keeps the last two backups by default.
- These backups are stored in a separate storage system and are not visible to us directly.
- If we need to restore data, we have to create a support ticket, and Microsoft will restore the data into a new Cosmos DB account.

2. **Continuous backup mode** (optional and more flexible):

- This captures every change in the data and keeps the history for the last 30 days.
- We can restore data to any point in time within the last 30 days.
- Restores are self-service, meaning we can perform the restore ourselves using the Azure portal, CLI, or ARM templates.
- This is very helpful in situations like accidental deletes or logical errors.

When we restore from backup, the data is placed into a new Cosmos DB account, so that it doesn't overwrite existing data. From there, we can manually move or copy the restored data as needed.

In short, Cosmos DB handles all backups in the background, gives options for both periodic and continuous backups, and makes recovery easy through support or self-service tools depending on the mode we choose.

9. How can you achieve global distribution with Azure Cosmos DB?

Global distribution in Cosmos DB is very simple and powerful. It is built into the service, and we can enable it directly from the Azure portal or using code.

To achieve global distribution:

- We go to the Replicate data globally section in the Cosmos DB account in the Azure portal.
- From there, we select the regions where we want to replicate our data. For example, we can add West US, East Asia, or Europe.
- Cosmos DB automatically replicates our data across these regions using multi-master or single-master replication, depending on our configuration.

Benefits of this approach:

- **Low latency:** Cosmos DB serves data from the nearest region to the user, which improves performance.
- **High availability:** If one region goes down, the system automatically fails over to another region.
- **Disaster recovery:** Data is not lost because it exists in multiple locations.
- **Multi-region writes:** We can enable writes in multiple regions if needed. This is helpful in scenarios where users from different parts of the world are writing data at the same time.

In code, we can specify the preferred regions in the Cosmos DB SDK. It tries to connect to the nearest region first and uses others as backups.

So global distribution in Cosmos DB is just a few clicks away, and Cosmos DB takes care of the replication, consistency, and failover for us.

10. How to integrate Azure Cosmos DB with other Azure services?

Azure Cosmos DB can be easily integrated with many other Azure services, which makes it a great choice for building complete solutions. Here are some common integrations:

- **Azure Functions:** We can trigger serverless functions when documents are added or modified in Cosmos DB. This is useful for real-time processing or automation.
- **Azure Data Factory:** We can use Data Factory to move data in and out of Cosmos DB. This is often used for ETL processes or to connect with other databases and storage systems.
- **Azure Synapse Analytics:** We can connect Cosmos DB with Synapse Link to run big data analytics directly on operational data, without moving it or affecting performance.
- **Azure Logic Apps:** We can use Logic Apps to automate workflows involving Cosmos DB, like sending an email when new data is inserted or pushing data to another system.
- **Power BI:** We can visualize data from Cosmos DB in dashboards and reports. This is usually done using Azure Synapse Link or by exporting the data to a supported format.
- **Azure Event Hubs or Service Bus:** For event-driven architectures, we can send changes in Cosmos DB to Event Hubs or Service Bus using Change Feed, and process them in other services.

All of these integrations are supported natively or through connectors, making it easy to connect Cosmos DB with other parts of the Azure ecosystem.

So, Cosmos DB acts like a central data store, and we can plug it into almost any Azure service to build flexible, scalable, and automated solutions.

11. What is Databricks access in Cosmos DB?

Databricks access in Cosmos DB refers to the ability to connect Azure Databricks with Azure Cosmos DB so that we can read or write data between them. This is useful when we want to do big data analytics, data transformation, or machine learning on the data stored in Cosmos DB.

In a typical scenario, Cosmos DB is used to store JSON data through the SQL API. We can connect Azure Databricks using the Cosmos DB Spark Connector, which allows Spark in Databricks to interact with Cosmos DB containers.

Here's how it works:

- We configure the Cosmos DB endpoint, primary key, database name, and container name in the Spark session.
- Then, we can use Spark DataFrame APIs to read data from Cosmos DB or write data to it.
- We can also use Databricks notebooks to run transformation logic or analytics on the Cosmos DB data.

For example, we can train a machine learning model in Databricks using customer data stored in Cosmos DB. After training, we can write predictions back to Cosmos DB.

This integration is powerful because:

- Cosmos DB provides low-latency operational data storage.
- Databricks provides scalable compute for advanced analytics and AI.

So, Databricks access allows us to combine real-time data in Cosmos DB with large-scale analytics in Databricks for more intelligent applications.

12. How can you monitor and troubleshoot performance issues in Azure Cosmos DB?

To monitor and troubleshoot performance in Cosmos DB, we can use several built-in tools and metrics available in the Azure portal and through Azure Monitor.

Here's how we can approach it:

1. Metrics in the Azure Portal:

- Cosmos DB provides charts for key metrics like Request Units (RU) consumed, throttled requests (429 status), latency, availability, and storage usage.
- If we see throttling, it means we are exceeding the provisioned RU/s, and we might need to scale up or optimize our queries.

2. Diagnostic logs and insights:

- We can enable diagnostic logging and send logs to Log Analytics, Event Hubs, or Storage Account.
- With Log Analytics, we can run queries to find slow operations, failed requests, or RU bottlenecks.

3. Query metrics:

- When we run a query in the portal's Data Explorer, we can click on "Query Stats" to see how many RUs were consumed, how many documents were scanned, and the index usage.
- If RUs are too high, we might need to improve the indexing policy or rewrite the query.

4. Alerts:

- We can set up alerts using Azure Monitor to notify us when metrics cross certain thresholds, like high RU consumption, high latency, or throttling events.

5. Change Feed:

- We can monitor changes to data using Change Feed and use that for troubleshooting data flow in event-driven systems.

6. Best practices:

- Make sure to use a good partition key to evenly distribute the load.
- Use point reads or queries with filters on indexed properties.
- Avoid fetching large documents unnecessarily.

So, Cosmos DB gives us all the tools needed to monitor usage, detect issues, and optimize performance we just need to keep an eye on the metrics and follow best practices.

13. How can I get extra help with Azure Cosmos DB?

If I need extra help with Azure Cosmos DB, I usually follow a few reliable methods:

1. I start with the official Microsoft documentation. It has detailed guides, tutorials, and examples that are very easy to follow. This is helpful when I'm trying to understand a concept or how to use a feature.
2. I use Microsoft Learn, which provides free, hands-on training modules. It helps me practice using Cosmos DB with guided labs, especially when I'm trying something new like integrating with other services.
3. For quick questions or issues, I search or ask in community forums like Stack Overflow or Microsoft Q&A. Many times, someone has already faced the same issue and shared the solution.
4. If I face an error that I can't solve, and it's blocking my work, I use the Azure Portal to open a support ticket. If I have a paid support plan, Microsoft engineers help me directly. They respond quickly and help with deeper troubleshooting.
5. I also check GitHub repositories and sample code provided by Microsoft and the community. These help me understand how Cosmos DB is used in real-world applications.

So whenever I get stuck, I combine these resources to get answers or support, depending on the type of help I need.

14. How do I migrate an Azure Cosmos DB account to a different resource group or to a different subscription?

Moving a Cosmos DB account between resource groups or subscriptions is not a direct process. Microsoft does not allow the account itself to be moved. Instead, I need to follow a step-by-step approach to migrate the data:

1. I create a new Cosmos DB account in the target resource group or subscription where I want to move the data.
2. Then, I use one of the supported data migration tools, like Azure Data Factory, Cosmos DB Data Migration Tool, or custom scripts using Azure SDKs.
 - With Azure Data Factory, I create a pipeline that copies data from the source Cosmos DB account to the new account. I choose the source and sink datasets and configure the connection strings.
 - Alternatively, I use the Cosmos DB Data Migration Tool if I want a simple, UI-based option to export and import data.
3. If my application uses Change Feed, I make sure the application handles reprocessing after migration to avoid missing any changes.
4. After migration, I test the new setup thoroughly to ensure everything is working as expected. I compare document counts, validate indexes, and verify that performance is similar.
5. Once the new setup is verified, I update my application's connection string to point to the new Cosmos DB account.
6. Finally, I delete the old account only after I'm fully sure that everything has been successfully moved and tested.

So, even though I can't move the Cosmos DB account directly between subscriptions or resource groups, I can still migrate the data safely using tools and a structured approach.

15. How do I migrate an Azure Cosmos DB account to a different tenant?

Moving an Azure Cosmos DB account to a different tenant is not possible directly because each tenant is tied to a separate Azure Active Directory. Instead, I follow a migration process that involves creating a new account in the target tenant and moving the data manually.

Here's how I handle it:

1. First, I log in to the destination tenant using the correct Azure Active Directory and subscription.
2. I create a new Cosmos DB account in the destination tenant with the same configuration as the source (same API type, regions, consistency, and indexing policy).
3. To move the data, I use one of the following methods:
 - **Azure Data Factory:** This is my preferred option. I create a pipeline with the source Cosmos DB as the input and the target Cosmos DB as the output. This supports bulk copy and can scale well.
 - **Cosmos DB Data Migration Tool:** This lets me export data from the source and import it into the destination. It's good for one-time or small migrations.
 - **Custom code with Azure SDK:** I write a script in C# or Python to read data from the source Cosmos DB and write to the new one. This gives me full control but needs more work.
4. If the source account is actively being used, I make sure to handle changes that happen during migration. I use Change Feed to capture new inserts or updates and apply them to the new account after the bulk copy is done.
5. Once all data is moved and verified, I update my applications to use the new Cosmos DB account's connection string in the new tenant.
6. I monitor the new account for a few days to confirm it behaves as expected. After that, I decommission the old account from the source tenant.

So, although I cannot move the Cosmos DB account between tenants directly, I can migrate the data safely using these tools and steps.

16. What is the best way to bulk-insert documents into Cosmos DB?

When I need to insert a large number of documents into Cosmos DB quickly, I use bulk operations because they are optimized to save time and Request Units (RUs). The best way depends on the tool and API I'm using, but here are the general approaches I follow:

1. Azure Cosmos DB SDK with bulk support:

- I use the latest SDKs (like the .NET SDK v3 or Java SDK v4) that have built-in support for bulk operations.
- In my code, I enable bulk mode in the client options and then send multiple write operations in parallel using `Task.WhenAll()` in C#.
- Cosmos DB batches these operations internally and processes them more efficiently.

Example in C#:

```
CosmosClientOptions options = new CosmosClientOptions() { AllowBulkExecution = true };  
CosmosClient client = new CosmosClient(endpoint, key, options);  
Container container = client.GetContainer("myDatabase", "myContainer");
```

```
List<Task> tasks = new List<Task>();  
foreach (var doc in myDocuments)  
{  
    tasks.Add(container.CreateItemAsync(doc, new PartitionKey(doc.partitionKey)));  
}  
await Task.WhenAll(tasks);
```

2. Azure Data Factory:

- I use a copy activity to load data from sources like Azure Blob Storage, SQL Server, or another Cosmos DB account.
- I configure the sink settings to support parallel writes and batch size.
- This method is low-code and works well for loading millions of records.

3. Cosmos DB Data Migration Tool:

- I use this tool for one-time migrations. It supports CSV, JSON, and MongoDB as input formats and can insert data in bulk into Cosmos DB.

4. Best practices I follow:

- I choose a good partition key to distribute the writes evenly.
- I batch documents with similar partition keys to reduce RU usage.
- I monitor RU usage and throttling during bulk insert and apply retries with exponential backoff.

Using bulk mode with SDKs is the fastest and most flexible way when I'm writing custom code. For low-code solutions or scheduled loads, Azure Data Factory is the best option.

17. When I import data into Azure Table Storage, I never get a “quota full” notification. I receive this message when using the Table API. Is this restricting my application?

Yes, this message can restrict your application, and the reason has to do with how Azure Cosmos DB's Table API works differently compared to traditional Azure Table Storage.

In Azure Table Storage, quotas are soft, and the system rarely blocks writes unless the limits are extremely exceeded. But when I use the Table API in Azure Cosmos DB, the platform enforces stricter resource usage limits based on Request Units (RU/s) and storage capacity. This is because Cosmos DB is designed for performance and scalability, so it protects itself from overload by giving quota-related messages.

Here's what could be happening:

1. RU/s limit reached:

- If I'm doing a bulk insert or high-frequency writes, I might be exceeding the RU/s capacity that I've provisioned for my Cosmos DB account.
- When this happens, Cosmos DB returns a 429 error (Request rate is large) or sometimes messages that look like quota issues.
- This doesn't mean storage is full, but that the database is limiting the number of operations to protect performance.

2. Storage limits in Free Tier or low-capacity accounts:

- If I'm using the Free Tier or a small Cosmos DB setup, I may be hitting the total storage limit, which could trigger a quota-related message.
- Cosmos DB enforces actual storage limits based on the pricing tier.

3. Partition key issues:

- If many inserts are going to the same partition key, the partition might become a hotspot and hit throughput or size limits for that specific partition.

So yes, receiving a “quota full” or similar message while using Cosmos DB's Table API is a signal that the system is either out of throughput or approaching its physical limits. It's not just a warning it actively restricts your operations until usage falls back within the limits.

To solve this, I usually:

- Check the metrics in the Azure portal to confirm if it's RU/s or storage causing the problem.
- Scale up the provisioned RU/s or enable autoscale to let Cosmos DB handle the traffic.
- Review partitioning strategy to make sure writes are evenly distributed.
- Consider moving large-scale data import to off-peak hours.

Understanding these differences helps me design my application to avoid hitting these restrictions.

18. How do you choose the right API for your application in Cosmos DB?

When I choose the right API in Cosmos DB, I always start by looking at what my application needs and what kind of data model it already uses. Cosmos DB supports multiple APIs, and each one is designed to match a specific type of workload or existing system.

Here's how I decide:

1. SQL API:

- I choose this when I want to store data in JSON format and need rich query support with SQL-like syntax.
- It's best for new applications that are being built on Cosmos DB from scratch.
- It supports powerful indexing, filtering, and projection using familiar query language.

2. MongoDB API:

- I use this if my existing application already works with MongoDB and I want to switch to Cosmos DB without rewriting code.
- Cosmos DB handles MongoDB wire protocol, so I can use existing MongoDB drivers and tools.

3. Cassandra API:

- I select this when I have an application that uses Apache Cassandra.
- I can move my existing Cassandra app to Cosmos DB without major changes, while gaining benefits like global distribution and auto-scaling.

4. Gremlin API:

- I use this when I'm working with graph-based data, like relationships between people, social networks, or recommendation systems.
- Gremlin API supports the Gremlin query language for traversing graph structures.

5. Table API:

- I pick this if I already have an application that uses Azure Table Storage and I want to migrate to Cosmos DB for more scalability and global distribution.
- Cosmos DB's Table API supports the same schema as Azure Table Storage but with more powerful performance and throughput options.

6. ETCD API:

- This is used in cloud-native applications that need a distributed key-value store. It's useful for scenarios like Kubernetes configuration storage.

In summary, I choose the API based on:

- What my current application already uses (to reduce changes).
- The type of data I'm dealing with (document, key-value, columnar, graph).
- The query and performance needs of the application.

Once the API is selected, it cannot be changed for that account, so I make sure to evaluate this carefully at the start of the project.

19. How does Cosmos DB ensure data durability?

Cosmos DB ensures data durability by storing every piece of data across multiple replicas within the same region. It uses a technique called **quorum-based writes**, which means that a write is only considered successful when it is written to a majority of the replicas.

Here's how it works:

- Each region in Cosmos DB has four replicas of the data. When a write operation happens, the data is written to at least three out of those four replicas before returning success.
- This protects the data from hardware failures, power loss, or unexpected crashes because there are always extra copies available.
- All of these replicas are constantly synced, so even if one fails, Cosmos DB can continue to serve requests from the others.
- If I enable multi-region writes, the data is also replicated to other selected regions, making it durable across geographies.
- Behind the scenes, Cosmos DB uses write-ahead logs and journaling to make sure that even if a process fails during a write, the operation can be safely recovered.

So, the combination of multiple replicas, quorum-based writes, and cross-region replication ensures that once data is written to Cosmos DB, it is safe and won't be lost.

20. What is the Request Unit (RU) in Cosmos DB and how is it used?

Request Unit, or RU, is the performance currency in Cosmos DB. It is a way to measure how much processing power is required for an operation like reading, writing, or querying data.

Every operation in Cosmos DB — such as inserting a document, reading a record, or running a query — consumes a certain number of RUs. The amount depends on things like:

- Size of the document
- Type of operation (read, write, update, query)
- Complexity of the query
- Whether the fields being queried are indexed

Here's how I use it:

- When I create a Cosmos DB container, I can choose how many RUs per second I want to provision. This decides how much load the system can handle at a time.
- If I exceed the limit, Cosmos DB will throttle requests and return a 429 error, meaning I need to retry later or increase the RU/s.
- For unpredictable workloads, I use **autoscale mode**, which automatically adjusts the RUs based on the usage up to a maximum limit.

For example:

- A point read of a small document usually costs 1 RU.
- A write of a 1 KB document may cost around 5 RUs.
- A complex query filtering multiple fields might cost 100+ RUs depending on indexing and result size.

So, understanding and monitoring RUs helps me optimize performance and control costs in Cosmos DB.

21. How does indexing affect performance in Cosmos DB?

Indexing in Cosmos DB has a big impact on performance, especially for query and write operations.

By default, Cosmos DB automatically indexes all fields in all documents. This means:

- I can run queries on any field without creating a separate index.
- Queries are fast because Cosmos DB uses the index to locate data instead of scanning all documents.

However, indexing also affects write performance and RU consumption:

- When I insert or update a document, Cosmos DB must also update the index, which increases the RU cost.
- If I'm inserting a large number of documents, indexing every field can slow things down and cost more RUs.

That's why Cosmos DB allows custom indexing policies. I can:

- Exclude specific paths from being indexed if I know they won't be used in queries.
- Use indexing modes like lazy (for background indexing) or none (to skip indexing temporarily).
- Include only the fields I need for searching, to reduce write cost.

So, indexing improves query speed but increases the cost of writes. I always try to balance both by customizing the indexing policy based on how my application uses the data.

22. Discuss the ACID properties in the context of Cosmos DB transactions.

Cosmos DB supports ACID properties at the single-partition level, which means any transaction that happens within the same logical partition is guaranteed to be atomic, consistent, isolated, and durable.

Here's how each ACID property applies:

- **Atomicity:** If I perform multiple operations in a transaction (like insert and update), either all of them succeed or none of them are applied. For example, if I insert two documents in the same transaction and one fails, both are rolled back.
- **Consistency:** During the transaction, any read will return the most recent committed data. This ensures my data is always correct within that partition.
- **Isolation:** Cosmos DB uses snapshot isolation. Other operations cannot see changes made inside an ongoing transaction until it's completed. This prevents dirty reads and race conditions.
- **Durability:** Once the transaction is committed, the data is safely stored across multiple replicas. Even if there's a crash or failure, the data won't be lost.

Cosmos DB allows me to write stored procedures in JavaScript that run inside a single partition. These procedures can contain multiple operations and are fully ACID-compliant. However, cross-partition transactions are not natively ACID. To achieve that, I need to manage it in the application or use other patterns like eventual consistency or compensating transactions.

So, Cosmos DB provides full ACID transactions within one partition, which is good for most use cases when data is well-partitioned.

23. What are the best practices for choosing a partition key in a Cosmos DB container?

Choosing the right partition key is very important in Cosmos DB because it affects performance, scalability, and cost. I follow these best practices:

1. **Even data distribution:** I choose a key that spreads the data evenly across partitions. For example, if I'm storing orders, I use `userId` or `orderId` instead of `country` because `country` may lead to uneven distribution.
2. **High cardinality:** I use a field with many unique values to create more partitions and avoid hot spots. Fields like `customerId` or `sessionId` are better than something like `status`, which might only have a few values.
3. **Query efficiency:** I pick a key that matches my most common query filters. If I often query by `deviceId`, it makes sense to use that as the partition key.
4. **Write throughput:** To handle large-scale writes, the key should support spreading write load across partitions. For example, a social media app might use `userId` to prevent all writes going to one partition.
5. **Transactional boundaries:** Since transactions only work within a partition, I choose a key that groups related data together. For instance, all items in a shopping cart can have the same `cartId` as the partition key.
6. **Avoiding large partitions:** I make sure the data in each partition does not grow too big. Cosmos DB limits a logical partition to 20 GB of data.

If I'm unsure, I sometimes combine multiple fields to create a better key, like `userId#deviceId`.

So, the best partition key depends on how I store and access the data, and I always consider distribution, query patterns, and transaction requirements before choosing it.

24. What considerations should be taken into account when modeling data for a Cosmos DB instance?

When I model data for Cosmos DB, I don't follow the traditional relational database approach. Cosmos DB is a NoSQL database, so the design depends more on access patterns and performance needs. Here are the key things I consider:

1. **Denormalization:** Instead of spreading data across multiple tables like in SQL, I often combine related data into a single document. For example, an order document might include customer details, item list, and shipping info together. This reduces the number of joins and improves query performance.
2. **Access patterns first:** I design the structure based on how the application will read and write data. If I mostly read by userId, I make sure the user-related data is grouped and indexed accordingly.
3. **Partition key selection:** I carefully choose the partition key to ensure good distribution and support for queries and transactions, as explained earlier.
4. **Document size limits:** Each document can be up to 2 MB, so I avoid storing too much nested or repeated data. If the size gets too large, I split the data across documents and link them using a common key.
5. **Indexing strategy:** By default, Cosmos DB indexes all fields. If I don't need to query all fields, I customize the indexing policy to reduce RU usage and improve write performance.
6. **Change Feed needs:** If I plan to track changes in data (like real-time updates), I make sure my model supports it. This often involves including timestamps or status fields to monitor changes over time.
7. **Schema flexibility:** Cosmos DB is schema-less, but I still keep consistency in the document structure for better maintainability and easier querying.

In short, when I model data for Cosmos DB, I focus on what the app needs most — fast reads, efficient writes, and simple data access — and then shape the data to support that behavior, even if it means repeating or combining information.

25. How does partitioning impact the scalability and performance of a Cosmos DB application?

Partitioning is a key concept in Cosmos DB, and it directly impacts how well the application scales and performs.

Cosmos DB uses horizontal partitioning, which means it splits data into multiple physical partitions based on the partition key we define. This allows the system to distribute data and workloads evenly across different servers.

Here's how partitioning affects performance and scalability:

1. Scalability:

- More partitions allow Cosmos DB to handle more data and higher throughput.
- For example, if I provision 10,000 RU/s and my data is evenly partitioned, Cosmos DB spreads that RU budget across all the partitions.
- This means more users and larger workloads can be supported without performance issues.

2. Query performance:

- If the partition key is well chosen, queries can be directed to just one partition (this is called single-partition query), which is very fast and efficient.
- If the query needs to scan multiple partitions (cross-partition query), it takes more time and consumes more RUs.

3. Write throughput:

- Cosmos DB distributes write operations across partitions. If one partition receives most of the writes (called a hot partition), it becomes a bottleneck.
- But if writes are evenly distributed, the system can process them faster and more reliably.

In simple terms, partitioning lets Cosmos DB grow with the workload. A good partition key helps the app perform better, while a bad one can lead to throttling, slow queries, or even limits on storage per partition.

26. How might you handle hot partitions in Cosmos DB?

A **hot partition** happens when too many reads or writes go to the same partition, causing it to become overloaded. This reduces performance and can lead to request throttling or delays.

Here's how I handle hot partitions:

1. Choose a better partition key:

- I avoid low-cardinality keys like country or status, which don't spread data well.
- I use high-cardinality keys like userId, orderId, or sessionId so the load is balanced.

2. Add randomness to the key:

- If I can't change the key entirely, I sometimes add a suffix or prefix to the key to spread the load. For example, instead of userId, I use userId#1, userId#2, etc.

3. Use synthetic partition keys:

- I create a new key that combines two or more fields. For example, productId#regionId spreads access better than using just productId.

4. Limit request rate per user or client:

- I apply client-side throttling or use a message queue to spread writes over time instead of all at once.

5. Split large operations into smaller batches:

- Instead of writing or querying a lot of data in one go, I break the task into smaller parts that run on different partitions.

6. Monitor with Azure Metrics:

- I use the Cosmos DB metrics in the portal to track RU usage and identify which partition is hot. This helps me adjust my strategy quickly.

By fixing hot partitions, I make sure that no single partition is overloaded, which helps keep the application fast and scalable.

27. Describe some common anti-patterns in data modeling for Cosmos DB.

There are several mistakes I try to avoid when designing data models in Cosmos DB because they can affect performance, scalability, and cost. Here are the most common ones:

1. Using a bad partition key:

- Picking a low-cardinality or unevenly distributed key (like country or deviceType) can cause hot partitions and limit scalability.
- I always aim for a high-cardinality, evenly distributed key.

2. Trying to normalize everything like in relational databases:

- Cosmos DB works best when data is denormalized. If I split data into many small related documents, I end up making multiple reads or queries, which increases RU costs and complexity.

3. Overusing cross-partition queries:

- Designing queries that scan multiple partitions frequently is expensive and slow.
- I always try to keep most queries targeted to a single partition.

4. Ignoring document size limits:

- Cosmos DB has a 2 MB document size limit. If I keep adding nested arrays or embed too much data, the document becomes too large to store or read efficiently.

5. Indexing too much:

- By default, all fields are indexed. If I don't customize this, I may waste RUs on writes that update unnecessary indexes.
- I optimize indexing policies to include only the fields needed for queries.

6. Storing too many small writes in separate documents:

- If I'm storing logs or telemetry data as individual documents, it's better to group them (batch inserts) to reduce RU consumption.

7. Not handling schema changes properly:

- Cosmos DB is schema-less, but if I don't plan for future changes, I may end up with inconsistent document formats that make querying difficult.

To avoid these issues, I always design with access patterns in mind, test with real data, monitor usage, and update my data model when needed. This helps me keep the system efficient, fast, and cost-effective.

28. How do you perform SQL queries in Cosmos DB using the SQL API?

Cosmos DB supports SQL-like syntax through its SQL API, which lets me query JSON documents directly using familiar SELECT statements. This is the most commonly used API for Cosmos DB.

Here's how I perform SQL queries using the SQL API:

1. Using the Azure Portal:

- I go to my Cosmos DB account in the Azure portal.
- Under the "Data Explorer," I select the container I want to query.
- I click on "New SQL Query" and write a query like:

```
SELECT * FROM c WHERE c.city = "Mumbai"
```

- This will return all documents where the city field is Mumbai.

2. Using SDK in code:

- I use Cosmos DB SDKs (for .NET, Python, Node.js, etc.) to run SQL queries in my application.
- In .NET, for example:

```
var query = "SELECT * FROM c WHERE c.age > 30";  
var iterator = container.GetItemQueryIterator<Person>(query);
```

```
while (iterator.HasMoreResults)  
{  
    foreach (var item in await iterator.ReadNextAsync())  
    {  
        Console.WriteLine(item.Name);  
    }  
}
```

3. Query structure:

- SELECT * FROM c: retrieves all documents.
- c is just an alias for the document.
- I can use WHERE, ORDER BY, TOP, and JOIN for filtering and shaping results.

4. Parameterization:

- To avoid SQL injection and reuse queries, I use parameters:

```
var query = new QueryDefinition("SELECT * FROM c WHERE c.name = @name")  
    .WithParameter("@name", "Shubham");
```

5. RU impact:

- The more fields I query and the more data returned, the more RUs are consumed. Using indexed fields and limiting results helps improve performance.

So, with the SQL API, I can run powerful queries on JSON documents using a syntax that's easy to learn and similar to SQL Server or T-SQL, but adapted for JSON.

29. Can you use LINQ to query Cosmos DB? If so, explain how.

Yes, I can use LINQ (Language Integrated Query) to query Cosmos DB, especially when I use the .NET SDK. LINQ is useful because it allows me to write type-safe queries in C# without writing raw SQL strings.

Here's how I use LINQ with Cosmos DB:

1. Set up the Cosmos DB client and container:

```
CosmosClient client = new CosmosClient(endpoint, key);  
var container = client.GetContainer("myDatabase", "myContainer");
```

2. Use LINQ through the SDK's GetItemLinqQueryable method:

```
var query = container.GetItemLinqQueryable<Person>(allowSynchronousQueryExecution: true)  
    .Where(p => p.Age > 25 && p.City == "Delhi")  
    .ToList();
```

3. Things to know:

- LINQ gets translated into Cosmos DB SQL under the hood.
- Only supported LINQ operations will work. Some complex C# expressions may not be converted properly.
- Using LINQ makes the code easier to read and integrate with other .NET components.

4. Best practices:

- I always filter on the partition key when possible, so the query is scoped to a single partition.
- I avoid calling .ToList() directly on large result sets unless I want to load everything at once.

So yes, I use LINQ with Cosmos DB when working in .NET, and it helps me write cleaner and more maintainable code, especially in applications where logic and queries are tightly connected.

30. Explain how the Gremlin API can be used to query graph data in Cosmos DB.

Cosmos DB supports graph-based data using the **Gremlin API**. This is useful when I need to model and query relationships between entities, like in social networks, product recommendations, or organizational hierarchies.

Here's how I use the Gremlin API in Cosmos DB:

1. Modeling the data:

- In graph databases, data is stored as vertices (nodes) and edges (relationships).
- For example, in a social app, a person is a vertex, and the "follows" relationship is an edge.
- I store both vertices and edges in the same Cosmos DB container.

2. Creating a graph database:

- In the Azure portal, I create a Cosmos DB account and choose the Gremlin (graph) API.
- Then I create a graph (container) inside a database to hold the graph data.

3. Using Gremlin queries:

- Gremlin is a graph traversal language. I use it to navigate through connected nodes.
- Examples:

- Get all friends of a person:

```
g.V().has('name', 'Shubham').out('friend')
```

- Find people who follow someone who follows me:

```
g.V().has('name', 'Shubham').in('follows').in('follows')
```

- `g.V()` means start from all vertices.
- `out('friend')` follows outgoing edges labeled "friend".

4. Development tools:

- I use Gremlin SDKs for .NET, Java, or Node.js, or use tools like the Gremlin console to run queries.
- Cosmos DB also supports running Gremlin queries from Data Explorer in the portal.

5. Partitioning:

- Graph data must be partitioned properly using a good partition key (often id or some grouping key).
- Queries should be designed to stay within partitions when possible, to improve performance.

So, the Gremlin API in Cosmos DB allows me to work with connected data easily and run deep relationship queries using simple traversal steps, making it powerful for applications that depend on complex relationships.

31. Describe how Cosmos DB's MongoDB API differs from native MongoDB.

Cosmos DB provides a MongoDB-compatible API, which allows me to use MongoDB drivers and tools to interact with Cosmos DB as if it were a MongoDB server. However, there are some differences under the hood.

Here's what's similar:

- I can use MongoDB commands, collections, and query language like `find()`, `insertOne()`, `updateOne()`, and so on.
- I can connect using standard MongoDB drivers.
- Tools like MongoDB Compass or Robo 3T work with Cosmos DB.
- It supports BSON document format, JSON schema, and MongoDB wire protocol.

Now here's how Cosmos DB differs from native MongoDB:

1. Underlying engine:

- Cosmos DB is not built on MongoDB. It only mimics the API layer. Internally, it uses its own storage engine, indexing, and architecture.

2. Indexing:

- Cosmos DB automatically indexes all fields by default. In native MongoDB, I have to create indexes manually.
- Cosmos DB's indexing behavior may lead to higher write RU costs unless I customize it.

3. Scalability and global distribution:

- Cosmos DB has built-in support for global distribution and multi-region writes, which is not available by default in MongoDB.
- Cosmos DB also provides automatic failover and replication across regions.

4. Throughput model (RU/s):

- Cosmos DB charges and limits based on Request Units, not document size or IOPS like MongoDB.
- This means performance and cost tuning are different. I have to monitor RU consumption.

5. Feature support:

- Cosmos DB MongoDB API supports most but not all MongoDB features.
- Some advanced features like transactions across collections, aggregation pipeline stages, or certain operators may be limited or behave differently.

6. Data size and partitioning:

- Cosmos DB automatically manages partitions using a partition key. In MongoDB, I use sharding for similar scaling, but I have more control over it.
- Cosmos DB limits each logical partition to 20 GB.

In short, Cosmos DB's MongoDB API is great for migrating or extending MongoDB apps to Azure, but I always test my queries and features first, because the engine behind it is different and may not support 100% of MongoDB's native behavior.

32. Discuss the use of stored procedures, triggers, and UDFs (User-Defined Functions) in Cosmos DB.

In Cosmos DB, I can use stored procedures, triggers, and user-defined functions to write custom logic directly inside the database. These are written in JavaScript and are executed in the context of a single partition.

Here's how I use each one:

1. Stored procedures:

- A stored procedure is a group of operations bundled into one script.
- It runs as a single transaction within the same logical partition, so it follows ACID properties.
- I use stored procedures when I need to perform multiple operations together, like inserting several documents or updating related values in one atomic step.

Example:

```
function insertItems(items) {  
    var context = getContext();  
    var collection = context.getCollection();  
    var count = 0;  
  
    function insert(item) {  
        var isAccepted = collection.createDocument(collection.getSelfLink(), item, function(err) {  
            if (err) throw err;  
            count++;  
            if (count < items.length) {  
                insert(items[count]);  
            }  
        });  
  
        if (!isAccepted) throw new Error("Request not accepted");  
    }  
  
    insert(items[count]);  
}
```

2. Triggers:

- Triggers are small scripts that run automatically before or after a write operation.
- A pre-trigger runs before an insert or update, often used for validation or modifying the data.
- A post-trigger runs after the operation, useful for logging or cleanup.

For example, I use a pre-trigger to check if a field exists or meets a certain condition before inserting the document.

3. User-defined functions (UDFs):

- A UDF is a JavaScript function that I can call inside a SQL query.
- It runs on each document during the query and helps apply custom logic that cannot be written with standard SQL syntax.
- It doesn't modify data — it's only used for reading and transforming data.

Example:

```
function taxRate(amount) {  
  return amount * 0.18;  
}
```

Then I use it in a query:

```
SELECT c.id, udf.taxRate(c.totalAmount) AS tax FROM c
```

Important notes:

- All three (stored procedures, triggers, and UDFs) only work within a single partition.
- They consume Request Units, so I use them carefully to avoid high costs.
- Since they run inside Cosmos DB, they can reduce round-trips between the app and database.

I use these features when I need more control over logic inside the database, but I always design them to stay lightweight and focused to avoid performance issues.

33. What is the impact of cross-partition queries on the performance of Cosmos DB?

Cross-partition queries happen when a query cannot be satisfied by reading data from just one partition. Instead, Cosmos DB has to check multiple partitions to get the results.

Here's how this impacts performance:

1. More Request Units consumed:

- When a query scans multiple partitions, it uses more RUs compared to a single-partition query.
- This increases cost and may lead to throttling if RU limits are hit.

2. Slower query execution:

- Since Cosmos DB has to gather results from many partitions and combine them, the response time is usually longer.
- The more partitions involved, the higher the latency.

3. Limited query capabilities:

- Some queries, like ORDER BY, TOP, or JOIN, have limitations in cross-partition scenarios.
- They still work, but Cosmos DB might need to load more data or sort across partitions, which adds overhead.

4. More complex execution plans:

- The system has to coordinate the query across partitions, which adds background complexity.

To reduce the negative impact of cross-partition queries, I follow these best practices:

- I filter queries using the partition key when possible. This limits the query to a single partition.
- I use proper indexing to reduce the number of documents scanned.
- I use pagination (MaxItemCount) to retrieve data in smaller chunks.
- I monitor RU usage and adjust throughput if needed.

In summary, cross-partition queries are powerful but can be expensive and slower. So I always try to design my data model and queries to minimize them unless absolutely necessary.

34. How can you optimize query performance in a Cosmos DB database?

To optimize query performance in Cosmos DB, I focus on making my queries efficient and ensuring that the data model supports fast access. Here are the key steps I follow:

1. Use the partition key in queries:

- I always try to include the partition key in my queries. This ensures the query targets only one partition instead of scanning across many, which makes it faster and cheaper.

2. **Use filters early and avoid SELECT **:

- I avoid selecting all fields unless needed. I only return the required properties to reduce the amount of data read and the RU cost.
- I also use filters that match indexed fields early in the query.

3. Avoid cross-partition queries when possible:

- If I can rewrite the query to stay within one partition, it will perform much better.
- For example, filtering by userId or deviceId helps target the exact partition.

4. Optimize indexing:

- Cosmos DB automatically indexes all fields, but I customize the indexing policy to include only the fields I actually query.
- Removing unnecessary indexes reduces the size of the index and improves write performance.

5. Use pagination:

- Instead of fetching a large number of results at once, I use MaxItemCount with continuation tokens to page the results. This avoids large RU spikes and improves responsiveness.

6. Precompute and denormalize data:

- I store commonly needed values or summaries inside the documents so that queries can access them directly without extra joins or calculations.

7. Use efficient query patterns:

- I avoid expensive operations like CONTAINS or wildcard searches unless necessary.
- I use exact matches or range filters when possible, as they perform faster on indexed fields.

8. Monitor and test:

- I use the Query Metrics feature to see how much RU each query is using, and I refactor queries that are too costly.
- I test with real data and adjust queries to reduce RU consumption and latency.

By following these practices, I make sure that my Cosmos DB queries are fast, cost-effective, and scalable even with large datasets.

35. How can you manage and optimize throughput in Cosmos DB?

Managing and optimizing throughput in Cosmos DB is about balancing performance, scalability, and cost by correctly provisioning and using Request Units (RUs).

Here's how I manage it:

1. Choose the right throughput model:

- **Manual (provisioned) throughput:** I set a fixed RU/s value. Good when the workload is predictable.
- **Autoscale throughput:** Cosmos DB automatically adjusts RUs based on traffic (between 10% and 100% of the max limit). I use this when the traffic is unpredictable or spiky.

2. Distribute workload using partitioning:

- I choose a good partition key that spreads reads and writes evenly.
- This allows Cosmos DB to spread RU load across physical partitions and avoid hotspots.

3. Monitor and respond to throttling (429 errors):

- If I see frequent 429 errors, it means my operations are exceeding the allocated RU/s.
- I can handle this by retrying with exponential backoff or increasing the RU/s limit.

4. Optimize queries and data model:

- I reduce RU consumption by writing efficient queries and storing data in a way that avoids the need for costly joins or scans.

5. Use bulk operations where possible:

- Bulk inserts or updates are more efficient because Cosmos DB handles them with fewer network calls and internal optimizations.

6. Scale throughput at the right level:

- I can set throughput at the container level or the database level (shared across containers).
- For apps with multiple containers that don't all need high performance, I use shared throughput to save costs.

7. Use metrics and alerts:

- I monitor metrics like RU consumption, throttled requests, and average request latency.
- I set alerts when usage is high, so I can adjust throughput before it becomes a problem.

By carefully setting the right RU model, choosing a smart partition key, and keeping queries and writes efficient, I manage throughput in a way that keeps my Cosmos DB app fast and cost-effective.

36. Explain the use of the indexing policy in Cosmos DB.

The indexing policy in Cosmos DB controls how documents are indexed and which fields are included in the index. Indexing allows me to query data quickly, but indexing everything by default can increase storage size and RU consumption, especially during writes.

Here's how I use indexing policy:

1. Default behavior:

- Cosmos DB automatically indexes all fields in all documents. This means I don't have to manually define indexes, and I can query any field immediately.
- This is great for flexibility but may lead to higher RU usage on write operations.

2. Custom indexing policy:

- I can customize the indexing policy to improve performance and reduce costs.
- I do this by:
 - Including only specific fields that are commonly queried.
 - Excluding fields that are never queried.
 - Controlling the precision and types of index (range, spatial, or composite).

3. Indexing modes:

- **Consistent:** Indexes are updated immediately with every write. This is the default and ensures all queries see up-to-date indexes.
- **Lazy:** Indexes are updated in the background, which reduces write cost but may delay query accuracy.
- **None:** Disables indexing entirely for a container. I use this for write-heavy containers where I don't need queries.

4. Composite indexes:

- These are useful when I need to sort or filter on multiple fields together, especially for ORDER BY queries with more than one field.
- Without composite indexes, such queries are not efficient.

5. Index paths:

- I can define **included paths** and **excluded paths** to precisely control which properties are indexed.
- For example:

```
{
  "indexingMode": "consistent",
  "includedPaths": [
    { "path": "/name/?" },
    { "path": "/createdDate/?" }
  ],
}
```

```
"excludedPaths": [  
  { "path": "/*" }  
]  
}
```

This indexes only the name and createdDate fields, while skipping everything else.

So, indexing policy is important because it lets me control the trade-off between query performance and write efficiency. By customizing it to match my real queries, I reduce RU usage and make the database more optimized.

© Shubham Wadekar

37. Discuss the relevance of time-to-live (TTL) settings in Cosmos DB.

Time-to-live (TTL) in Cosmos DB is a setting that allows me to automatically delete documents after a certain period of time. It's very useful when I want to keep my data store clean and avoid storing outdated or unnecessary information.

Here's how TTL works and why it's helpful:

1. Enabling TTL:

- I enable TTL at the container level, and optionally at the document level.
- Once enabled, Cosmos DB checks each document's age and deletes it after the TTL period expires.

2. How TTL is set:

- I can specify the TTL value in seconds at the container level.
- Documents inherit this setting unless they override it with their own TTL.
- Example: If I set TTL to 86400 seconds (24 hours), each document will be removed automatically 24 hours after it is created or last modified.

3. Use cases:

- Temporary data like sessions, cache, telemetry, logs, or notifications.
- Automatically cleaning up expired items without writing any background job or cleanup code.

4. Benefits:

- Keeps the container size under control.
- Reduces storage costs.
- Improves performance by removing stale data that would otherwise be scanned or indexed.

5. Important notes:

- TTL deletions are handled by the system in the background and are not immediate.
- TTL operations consume a small number of RUs, but it's usually very efficient.
- I can disable TTL at any time if I want to stop auto-deletion.

In short, TTL is a simple but powerful way to manage data lifecycle in Cosmos DB. It helps me automate cleanup and ensure the system only keeps what's relevant, without manual work.

38. What are some common performance bottlenecks in Cosmos DB and how can you address them?

There are several common performance bottlenecks in Cosmos DB, and I always try to design my solution in a way that avoids them. Here's what I've seen and how I deal with each:

1. **Poor partition key choice**

If the partition key does not distribute data evenly, one partition may become too busy, which slows things down. This is called a hot partition. To fix this, I choose a partition key with high cardinality and even access patterns, such as `userId` or `deviceId`.

2. **Cross-partition queries**

When queries don't include the partition key, Cosmos DB has to search across many partitions, which increases RU consumption and slows down performance. I include the partition key in my queries wherever possible to keep them targeted and efficient.

3. **Large document sizes**

Cosmos DB has a 2 MB limit per document. If a document becomes too large due to nested arrays or unnecessary fields, it impacts read and write performance. I avoid storing too much data in one document and use reference documents when needed.

4. **Over-indexing**

By default, all fields are indexed. If I don't customize the indexing policy, every write operation takes more RUs because all fields are updated in the index. I reduce write cost by only indexing the fields I need for queries.

5. **Unoptimized queries**

Queries that use operations like `contains` or `sort` on unindexed fields can be slow and consume more RUs. I optimize my queries to use indexed fields and avoid scanning large datasets unnecessarily.

6. **Not handling throttling (429 errors)**

If the request rate exceeds the available RUs, the database will throttle requests. I handle this by adding retry logic with exponential backoff in my application and monitoring RU usage to adjust throughput if needed.

7. **Frequent updates on the same document**

Cosmos DB stores documents as whole units. Frequent updates, especially on large documents, can be costly. I reduce the number of updates or split large documents into smaller ones if they change frequently.

To avoid these bottlenecks, I always monitor performance metrics, use the Cosmos DB query metrics tool, and test with realistic data loads before going live.

39. Discuss best practices for using the change feed to maintain materialized views in Cosmos DB.

The change feed in Cosmos DB lets me listen to changes in a container in real time. This is very useful when I want to maintain a materialized view, which is like a pre-computed summary or report based on data from the main container.

Here are the best practices I follow:

- 1. Use the same partition key**

When maintaining a materialized view, I try to use the same partition key in both the source and the view container. This helps me group related changes and process them efficiently.

- 2. Process changes in order**

Cosmos DB guarantees the order of changes within a partition. I use this to safely process changes one by one and ensure my view is always consistent.

- 3. Use Azure Functions or Azure Data Factory**

I often use Azure Functions with a Cosmos DB trigger to automatically react to new changes. This is a serverless and cost-effective way to keep materialized views updated.

- 4. Filter unnecessary data**

I only process the changes that affect my materialized view. If a document field that doesn't matter to the view changes, I skip it to save processing time and cost.

- 5. Keep processing idempotent**

I make sure my logic can safely handle retries. If a change is processed more than once (due to retries), the result should still be correct. This avoids duplicates or errors.

- 6. Handle deletions separately**

Change feed only shows inserts and updates, not deletions. If I need to track deletes, I either use a soft delete flag in the document or implement another mechanism to track what was removed.

- 7. Scale out using lease containers**

When working with high volumes, I use the change feed processor library. It uses a lease container to manage distributed workers and balances the load across multiple instances automatically.

- 8. Monitor and log failures**

I add logging and monitoring so that if something goes wrong while processing a change, I can trace it and reprocess it if needed.

Using the change feed with these best practices helps me build reliable, real-time data pipelines and keep my summary views accurate without constantly re-querying the full dataset.

40. How does eventual consistency in Cosmos DB work and when would you use it?

Eventual consistency means that when data is written to Cosmos DB, it may not be immediately available in all regions or replicas. Over time, the data becomes consistent across all locations, but there may be a short delay. This model gives the best performance and the lowest latency, especially for global applications.

Here's how it works in simple terms:

- When I write data to Cosmos DB using eventual consistency, the write is accepted and saved right away in the primary region.
- However, if I immediately try to read that same data from another region or replica, I might get an older version or no result yet.
- Eventually, the system makes sure that all regions have the latest version of the data, but it doesn't guarantee how fast that happens.

This model is useful in situations where:

- Reading slightly old data is acceptable.
- The application is read-heavy and performance matters more than immediate accuracy.
- There are multiple regions and I want to reduce the cost and latency of cross-region replication.

For example, I would use eventual consistency in:

- Product catalog pages where I just need to show a list of items, and it's fine if a recently added product shows up after a few seconds.
- News feeds, logs, or social media timelines where freshness is important but not critical.
- Analytics dashboards that can tolerate minor delays in updates.

So, I use eventual consistency when I want the highest availability and lowest latency and when strong accuracy is not required for every read.

41. What consistency level would you choose for a shopping cart application in Cosmos DB and why?

For a shopping cart application, I would choose **session consistency**.

Here's why:

- In a shopping cart, the user expects to see the most recent changes they made — like when they add or remove an item.
- However, it's not necessary for all users or all regions to see those changes instantly.
- Session consistency ensures that within a user session, all reads reflect the writes that user has made. This gives a good balance between performance and correctness.

For example:

- If I add a product to my cart and then immediately view the cart, I will see that product.
- But other users won't be affected by this data, so global strict consistency is not needed.

Session consistency is also more efficient than strong or bounded staleness, which can add latency or RU cost. So, for shopping cart scenarios where I care about user-specific consistency but not global consistency, session consistency is the best fit.

42. How does Cosmos DB handle global distribution?

Cosmos DB is designed to be globally distributed, which means I can replicate my data to multiple Azure regions all over the world. This helps me build applications that are fast and available to users regardless of their location.

Here's how Cosmos DB handles global distribution:

1. Multiple regions

I can add one or more Azure regions to my Cosmos DB account with just a few clicks. Once added, Cosmos DB automatically replicates data to those regions in the background.

2. Read and write access

I can configure any of the regions to be readable. If I enable multi-region writes, I can also allow writing from multiple regions at the same time.

3. Automatic failover

Cosmos DB provides automatic failover. If the primary region goes down, another region can take over without data loss, based on the chosen consistency level.

4. Low-latency access

When users connect to the closest region, they get faster responses. This improves user experience, especially for global applications.

5. Region-specific endpoints

Cosmos DB SDKs automatically route requests to the nearest available region using the account's configuration. This helps reduce network latency.

6. Data replication

Cosmos DB uses asynchronous replication across regions. The actual speed of replication depends on the chosen consistency level — stronger consistency means slower replication, while weaker consistency is faster.

By distributing data globally, Cosmos DB helps me build applications that are resilient, highly available, and perform well for users anywhere in the world.

43. Explain multi-region writes and how they affect consistency and performance.

Multi-region writes allow me to write data to any of the configured Cosmos DB regions, not just the primary one. This is helpful when I have users spread across different parts of the world and want them to experience low-latency write operations.

Here's how multi-region writes work and how they impact the system:

1. Performance benefits

- With multi-region writes, users can send data to the nearest region, which reduces write latency.
- This improves the experience for applications that handle user input or transactions from different locations.

2. High availability

- If one write region goes down, others can still accept writes. This ensures there's no downtime.
- Cosmos DB also automatically resolves conflicts if the same item is updated in multiple regions.

3. Conflict resolution

- When the same item is updated in more than one region at the same time, Cosmos DB uses a conflict resolution policy.
- By default, it uses the last write wins strategy, based on a timestamp. But I can also define a custom function to resolve conflicts.

4. Consistency trade-offs

- Enabling multi-region writes limits me to using consistency levels below strong. Strong consistency is only available in single-region write mode.
- So, if I need strict ordering and global consistency, I should not enable multi-region writes. But for most real-time applications, session or bounded staleness is usually good enough.

5. Cost consideration

- Enabling multi-region writes might slightly increase the cost due to more replication and potential conflict handling, but it's often worth it for the performance and availability benefits.

So, I use multi-region writes when I want to offer fast, reliable write performance across the globe and I'm okay with relaxing the consistency level a little. It's very helpful for modern apps like chat platforms, IoT systems, or global e-commerce platforms.

44. What are the benefits of geo-replication in Cosmos DB?

Geo-replication in Cosmos DB allows me to replicate data across multiple Azure regions. This brings several important benefits, especially for applications that need high availability, global reach, and fast performance.

Here are the key benefits:

1. **High availability**

Geo-replication ensures that even if one Azure region goes down, my data is still available in other regions. Cosmos DB automatically fails over to another region so that the application continues to work without downtime.

2. **Low latency for global users**

By placing the data closer to users in different parts of the world, reads and writes happen faster. For example, if my users are in India and the US, I can replicate data to both regions to reduce latency for both.

3. **Disaster recovery**

In case of natural disasters, outages, or major failures in one region, I don't lose data. Since the data is also stored in other regions, recovery is much faster and safer.

4. **Scalability across regions**

I can scale my app globally by handling traffic in multiple regions. This is useful for apps that grow quickly or operate in many countries.

5. **Automatic replication**

Cosmos DB handles all the replication automatically. I don't have to write any code or manage the sync. It keeps all regions up to date based on the chosen consistency level.

6. **Support for multi-region writes**

Geo-replication allows me to enable multi-region writes, which means users can write data to the nearest region. This makes write operations faster and increases fault tolerance.

7. **Compliance support**

Some countries have data residency requirements. By choosing specific regions for replication, I can meet those legal or compliance needs.

So, geo-replication gives me peace of mind, better performance, and more control over where and how data is stored. It's a major advantage of using Cosmos DB for building modern cloud applications.

45. How do you configure failover priorities for regions in Cosmos DB?

Configuring failover priorities in Cosmos DB is how I define the order in which other regions should take over if the main region fails. This helps me control how my application behaves during an outage.

Here's how I do it:

1. **Access the Azure portal**

I go to my Cosmos DB account in the Azure portal.

2. **Go to the "Replicate data globally" section**

This is where I can see all the regions where my Cosmos DB data is replicated.

3. **Set failover priorities**

Next to each region, I'll see a number. This number shows the priority order. Lower numbers have higher priority. I drag and drop the regions to change the order.

Priority 0 is the primary region (first to handle requests).

Priority 1 is the first backup region (used if primary fails), and so on.

4. **Enable automatic failover**

I turn on the automatic failover option so that Cosmos DB will switch to the next region in the list if the current one goes down.

5. **Manual failover (optional)**

If needed, I can also trigger a manual failover from the portal. This is useful for testing or for planned maintenance.

6. **Application behavior**

Cosmos DB SDKs automatically detect failover and re-route traffic to the next region, so my application continues to work without making any changes.

In summary, configuring failover priorities lets me decide which region should take over next during an outage. I always make sure to set this based on where most of my users are, so they get the best performance and availability even during unexpected failures.

46. How does Cosmos DB integrate with Azure Active Directory?

Cosmos DB integrates with Azure Active Directory (AAD) to provide secure and centralized identity and access management. This means I can control who can access Cosmos DB resources using the same identities and permissions defined in AAD, rather than managing separate keys or credentials.

Here's how the integration works:

1. **Role-based access control (RBAC)**

Cosmos DB supports Azure RBAC, so I can assign roles like Reader, Contributor, or Cosmos DB-specific roles to users, groups, or service principals in AAD. These roles define what actions they can perform on Cosmos DB resources, such as reading data, writing data, or managing the database.

2. **Authentication with Azure AD tokens**

Instead of using primary keys or connection strings, applications can authenticate to Cosmos DB using Azure AD tokens obtained through OAuth2. This is more secure and manageable, especially in enterprise environments.

3. **Granular access management**

Using Azure AD, I can control access at different levels — subscription, resource group, or individual Cosmos DB accounts and containers — depending on my organization's security needs.

4. **Integration with Managed Identities**

If my application is running in Azure (like in Azure Functions or Azure VMs), I can use Managed Identities to automatically get tokens and authenticate to Cosmos DB without storing credentials anywhere.

5. **Audit and compliance**

Because all access is tied to Azure AD identities, I get detailed logging and auditing capabilities. This helps meet security and compliance requirements.

In summary, integrating Cosmos DB with Azure Active Directory lets me manage access securely and conveniently, reduce risks from key leaks, and apply enterprise-grade identity controls.

47. How does Cosmos DB handle data at rest and in transit encryption?

Cosmos DB ensures data is secure both when it is stored (at rest) and when it is moving between clients and servers (in transit) by using strong encryption methods.

1. Encryption at rest

- All data stored in Cosmos DB is automatically encrypted using AES-256 encryption, which is a very strong and widely accepted standard.
- This encryption protects data on disks, backups, and replicas.
- The encryption keys are managed by Microsoft by default, but I can also use Customer-Managed Keys (CMK) stored in Azure Key Vault for more control over the encryption keys.

2. Encryption in transit

- When data moves between my application and Cosmos DB, it uses TLS (Transport Layer Security) encryption.
- This protects the data from being intercepted or tampered with during communication over the internet.
- All Cosmos DB endpoints enforce TLS by default.

3. Additional security

- Cosmos DB also supports IP firewalls and virtual network service endpoints, which restrict access to authorized networks, further protecting data while in transit.

Overall, Cosmos DB uses industry-standard encryption both at rest and in transit to ensure that data remains confidential and secure throughout its lifecycle.

48. Explain how you can secure access to Cosmos DB data using role-based access control (RBAC).

I secure access to Cosmos DB data using role-based access control (RBAC) by leveraging Azure Active Directory (AAD) identities and assigning specific roles that define what actions users or applications can perform.

Here's how I do it:

- 1. Use Azure AD identities**

Instead of relying on shared keys or connection strings, I configure Cosmos DB to accept authentication tokens issued by Azure AD. This lets me manage access using users, groups, or service principals in Azure AD.

- 2. Assign roles**

Azure provides built-in roles for Cosmos DB such as Cosmos DB Account Reader, Cosmos DB Contributor, and data-plane roles like Cosmos DB Built-in Data Reader or Data Contributor. I assign these roles to the right identities based on the least privilege principle — giving only the permissions needed.

- 3. Granular access control**

Roles can be assigned at different scopes — the entire subscription, a resource group, or a specific Cosmos DB account or container — allowing precise control over who can read, write, or manage resources.

- 4. Application integration**

My applications authenticate with Azure AD to obtain tokens, and these tokens carry the assigned roles. When the app calls Cosmos DB APIs, Cosmos DB verifies the token and authorizes actions accordingly.

- 5. Audit and monitor**

Since RBAC uses Azure AD, all access attempts and role assignments can be logged and monitored, helping me track who accessed what and detect any unauthorized attempts.

By using RBAC, I improve security by eliminating shared keys, simplifying access management, and ensuring only authorized identities can perform specific operations on Cosmos DB.

49. How does Cosmos DB integrate with other Azure services such as Azure Functions?

Cosmos DB integrates seamlessly with many Azure services, including Azure Functions, to build scalable, serverless, and event-driven applications.

Here's how the integration with Azure Functions works:

1. Cosmos DB trigger

Azure Functions can be set up with a Cosmos DB trigger, which listens to the change feed — a continuous stream of inserts and updates happening in a Cosmos DB container.

- Whenever new data is added or modified, the function is automatically triggered.
- This allows real-time processing of changes without polling or complex infrastructure.

2. Binding support

Azure Functions support input and output bindings for Cosmos DB:

- Input bindings let functions read data from Cosmos DB easily, by providing parameters or queries.
- Output bindings allow functions to write or update documents in Cosmos DB simply by returning data from the function.

3. Serverless architecture

Using Cosmos DB with Azure Functions enables event-driven architectures where the function runs only when needed, scaling automatically, and paying only for execution time.

4. Use cases

- Real-time data processing like updating materialized views or sending notifications.
- Data validation or enrichment as data arrives.
- ETL pipelines where Azure Functions transform data before storing or forwarding it.

5. Security and identity

Functions can authenticate to Cosmos DB using managed identities, so I don't have to manage connection strings or keys manually.

In summary, the integration between Cosmos DB and Azure Functions makes it easy to build reactive, scalable, and cost-effective solutions that respond immediately to data changes in the database.

50. What options are available to scale out a Cosmos DB collection?

To scale out a Cosmos DB collection, I primarily rely on partitioning and throughput management. Here are the main options:

1. Use logical partitioning

- Cosmos DB uses a partition key to divide the data into logical partitions.
- When I choose a good partition key with many unique values, Cosmos DB automatically distributes the data and requests across multiple physical partitions.
- This allows the collection to grow in size and handle more requests by adding more partitions behind the scenes.

2. Provision or autoscale throughput

- I can provision a higher throughput (measured in Request Units per second or RU/s) for the collection to handle more operations per second.
- Alternatively, I use autoscale throughput, which automatically scales the RU/s based on traffic, making it easier to handle variable workloads.

3. Add regions for global distribution

- By replicating the collection to multiple regions, I can scale read and write operations geographically.
- With multi-region writes enabled, I can scale write throughput globally.

4. Use bulk operations

- To handle large data loads efficiently, I use bulk insert or update APIs, which process many documents in fewer calls.

5. Optimize data model and indexing

- A well-designed data model and indexing strategy reduces RU consumption, allowing the same throughput to support more operations.

In short, scaling out a Cosmos DB collection is mainly about choosing the right partition key, adjusting throughput, and distributing data and traffic efficiently across partitions and regions.

51. Discuss how Cosmos DB fits into a microservices architecture.

Cosmos DB is well suited for microservices architectures because of its flexibility, scalability, and low-latency performance. Here's how it fits in:

1. Decentralized data management

- Each microservice can have its own Cosmos DB container or database, enabling data isolation and independence.
- This reduces coupling between services and makes it easier to develop, deploy, and scale them separately.

2. Flexible data models

- Cosmos DB supports multiple data models (document, key-value, graph, column-family), allowing each microservice to choose the best model for its needs.
- This flexibility supports diverse microservices with different data requirements.

3. Scalability

- Cosmos DB automatically scales throughput and storage based on the workload, matching the dynamic nature of microservices.
- Services that require more resources can scale independently.

4. Global distribution and low latency

- Cosmos DB's global replication allows microservices deployed worldwide to access data with low latency.
- It supports multi-region writes, helping microservices stay consistent in distributed systems.

5. Event-driven architecture support

- Cosmos DB's change feed enables event-driven microservices, where one service reacts to data changes made by another.
- This helps build reactive and loosely coupled microservices.

6. Strong integration with Azure ecosystem

- Cosmos DB works well with Azure Functions, Logic Apps, and Event Grid, which are often used in microservices for orchestration and event processing.

7. Security and compliance

- With Azure AD integration and RBAC, microservices can have secure and controlled access to data.

Overall, Cosmos DB supports microservices by providing a scalable, flexible, and globally distributed data platform that helps services remain independent while enabling seamless communication and data sharing.

52. Can you integrate Cosmos DB with Azure Event Hubs or Azure Service Bus? If yes, how?

Yes, Cosmos DB can be integrated with both Azure Event Hubs and Azure Service Bus, though the integration patterns differ depending on the use case.

1. Integrating Cosmos DB with Azure Event Hubs

- Cosmos DB has a feature called the change feed, which captures inserts and updates in the database as an ordered, reliable log of changes.
- I can build a process (using Azure Functions, Azure Stream Analytics, or a custom application) that listens to the Cosmos DB change feed and pushes those changes as events into Azure Event Hubs.
- This allows me to stream database changes to downstream systems, analytics platforms, or event-driven architectures.

2. Integrating Cosmos DB with Azure Service Bus

- Azure Service Bus is typically used for messaging between applications or services.
- I can use Azure Functions or other middleware to listen to Cosmos DB's change feed and send messages to Service Bus queues or topics when specific data changes occur.
- Conversely, I can have applications that receive messages from Service Bus and write data into Cosmos DB.

3. How to implement the integration

- Use Azure Functions with Cosmos DB trigger to react to data changes.
- Inside the function code, send events or messages to Event Hubs or Service Bus.
- This creates a near real-time pipeline from Cosmos DB to messaging or event systems.

4. Benefits

- This integration supports event-driven architectures, real-time processing, and decoupled communication between services.

In summary, Cosmos DB integrates well with Azure Event Hubs and Service Bus by using the change feed and Azure Functions to enable real-time data streaming and messaging.

53. How can Cosmos DB be used for real-time analytics?

Cosmos DB supports real-time analytics by providing fast, scalable data ingestion and immediate access to fresh data, combined with features that enable streaming and integration with analytics tools.

Here's how I use Cosmos DB for real-time analytics:

1. Ingest data rapidly

- Cosmos DB's low-latency writes and ability to scale throughput means I can capture large volumes of data from devices, applications, or users in near real-time.

2. Use the change feed

- The change feed lets me capture all inserts and updates as a continuous stream.
- I use this feed to trigger downstream analytics processes without polling or batch jobs.

3. Integrate with Azure Stream Analytics or Azure Functions

- I connect the change feed to Azure Stream Analytics jobs that perform real-time aggregation, filtering, or anomaly detection.
- Alternatively, Azure Functions can process the change feed to update dashboards, trigger alerts, or push data into other analytical stores.

4. Store pre-aggregated or materialized views

- I create materialized views or summary tables within Cosmos DB or another store to support fast queries and dashboards.

5. Use Cosmos DB's SQL API for real-time querying

- Cosmos DB's rich query capabilities allow me to run low-latency queries over fresh data directly, supporting live dashboards and reports.

6. Global distribution

- If I have users worldwide, Cosmos DB replicates data globally so analytics can be run closer to the data source or users for better performance.

7. Integration with Power BI and other BI tools

- Cosmos DB can be connected to Power BI through connectors or Azure Synapse Link for near real-time visualization.

In essence, Cosmos DB acts as both the operational database and the streaming data source for analytics, enabling real-time insights by combining fast data capture, event-driven processing, and immediate query capability.

54. Explain the integration options between Cosmos DB and Azure Stream Analytics.

Cosmos DB and Azure Stream Analytics work very well together when I need to process and analyze real-time streaming data. Azure Stream Analytics can both read data from and write data to Cosmos DB, depending on what I want to achieve.

Here are the main integration options:

1. Cosmos DB as a sink (output)

- This is the most common scenario. I use Azure Stream Analytics to process data coming from sources like Azure Event Hubs or IoT Hub, and then write the results directly into a Cosmos DB container.
- I define Cosmos DB as an output in the Stream Analytics job and provide the database name, container name, and authentication credentials.
- This is useful for writing processed, filtered, or aggregated data into Cosmos DB for dashboards, applications, or further queries.

2. Cosmos DB as a source (input)

- While Azure Stream Analytics does not support Cosmos DB as a direct input source, I can work around this by using the Cosmos DB change feed.
- I set up an Azure Function or Azure Data Factory to read the change feed from Cosmos DB and push those changes to an Event Hub or Blob Storage.
- Then I configure Stream Analytics to read from that source and process it in real-time.

3. Combining with other services

- I often connect Stream Analytics jobs with Power BI for visualization, and Cosmos DB acts as the storage layer that keeps the latest analytics results or historical data.
- Stream Analytics can also write to multiple outputs, so I can store some processed results in Cosmos DB and send alerts or events elsewhere.

In summary, Azure Stream Analytics can write real-time processed data into Cosmos DB directly and can indirectly consume data from Cosmos DB using the change feed, enabling powerful end-to-end streaming analytics pipelines.

55. Discuss patterns for processing and analyzing streaming data in Cosmos DB.

There are several patterns I follow when I need to process and analyze streaming data using Cosmos DB. These help me build efficient, scalable, and real-time data solutions.

1. Ingestion with Cosmos DB and change feed

- I use Cosmos DB as a fast, scalable ingestion point where devices or apps push streaming data like telemetry, logs, or transactions.
- I then use the change feed to detect these changes in real-time and trigger downstream processes.

2. Lambda-like pattern

- I separate the real-time layer (using Cosmos DB + Stream Analytics or Functions) from the batch or long-term analytics layer.
- Real-time updates go into Cosmos DB for quick actions or dashboards, while batched historical data is moved to a data warehouse like Azure Synapse for deep analysis.

3. Materialized views

- I use the change feed with Azure Functions to compute rolling summaries (like counts, averages, latest values) and write those into a separate Cosmos DB container.
- This allows my application to quickly query pre-processed results without running expensive queries over raw data.

4. Alerting and anomaly detection

- Stream Analytics jobs can process real-time data and detect patterns or thresholds (like temperature > 100°C).
- When such conditions are met, I write the alert to Cosmos DB and also notify users or services.

5. Enrich and transform

- While writing to Cosmos DB from Stream Analytics, I enrich the data with additional fields (timestamps, device info) and apply filters to store only relevant records.

6. Archival and lifecycle management

- I use Time-to-Live (TTL) in Cosmos DB to automatically remove old data after a certain period.
- This keeps the database lean and focused on recent, relevant streaming data.

7. Visualization

- I connect Power BI or custom dashboards to Cosmos DB to show live data updates, like real-time customer orders, IoT telemetry, or user activity.

These patterns allow me to build solutions that are not only real-time but also scalable and manageable, giving both fast response and deep insight using Cosmos DB as a core part of the streaming data architecture.

56. What SDKs are available for Cosmos DB and how do they cater to different languages?

Cosmos DB provides SDKs in many popular programming languages so that developers can easily interact with it using their preferred tools. These SDKs are designed to handle everything from creating and querying documents to managing containers, handling throughput, and working with features like change feed and partitioning.

Here are the main SDKs available and how they cater to different languages:

1. .NET SDK

- This is one of the most feature-rich SDKs and is often the first to get new capabilities.
- Supports LINQ queries, async operations, change feed processing, and fine-grained control over indexing and consistency.
- Ideal for C# and .NET-based applications, including Azure Functions and ASP.NET apps.

2. Java SDK

- Offers full support for Cosmos DB APIs, including partitioning, indexing, and connection management.
- Good for enterprise-grade Java applications and Spring Boot services.

3. Node.js SDK

- Lightweight and well-suited for JavaScript and TypeScript apps, including those built with Express.js or running in serverless environments.
- Handles document operations, queries, and session consistency easily.

4. Python SDK

- Popular in data science and machine learning projects.
- Simple interface to insert, read, update, and query documents.
- Often used in scripts and notebooks or in Flask and Django applications.

5. Go SDK

- Good for building fast, concurrent cloud-native services.
- Includes core features like item operations, partition key handling, and consistency settings.

6. Other SDKs

- Cosmos DB also has SDKs and support for languages like PHP, Ruby, and C++, though these are usually community-driven or have limited official features.

7. REST API and ODBC

- For platforms or languages without direct SDKs, I can use the Cosmos DB REST API or ODBC driver to interact with the database.

Each SDK supports native idioms of the language (like promises in JavaScript, await/async in .NET, or context management in Go) to make development easier and more consistent with the application style.

57. Explain how optimistic concurrency control works in Cosmos DB.

Cosmos DB uses optimistic concurrency control to make sure that two users or processes don't accidentally overwrite each other's changes. It does this using something called ETag, which stands for entity tag.

Here's how it works in simple steps:

1. ETag is like a version number

Every time a document is updated in Cosmos DB, the system generates a new ETag for that document.

Think of it like a hidden timestamp or version ID that changes with each update.

2. When I read a document

Along with the data, I also get the current ETag.

If I plan to update the document, I can use this ETag to make sure no one else has changed it before my update is applied.

3. When I write or update

I send the ETag back with the update request using an access condition (usually If-Match).

Cosmos DB checks whether the ETag matches the current version in the database.

4. If the ETag matches

That means the document hasn't changed since I read it, so the update goes through.

5. If the ETag doesn't match

It means someone else has updated the document after I read it.

In that case, Cosmos DB returns a precondition failed response (HTTP 412), and my update is rejected.

This approach avoids locking the data and keeps the system fast and scalable, but still prevents data from being accidentally overwritten.

Optimistic concurrency control is very useful in distributed systems where many users or services may be trying to update the same data at the same time. It helps ensure data accuracy without slowing things down.

58. What tools does Cosmos DB provide for local development and testing?

Cosmos DB provides a few helpful tools that allow me to develop and test applications locally without needing access to the live Azure environment. This makes it easier to build and debug applications faster and at no cost during development.

Here are the main tools:

1. Cosmos DB Emulator

- The emulator is a free tool I can install on my Windows machine to run a local instance of Cosmos DB.
- It supports most features of the SQL API, including partitioning, consistency levels, indexing, and query execution.
- It uses a local endpoint like `https://localhost:8081/`, and I use a fixed key provided by the emulator for access.
- It's great for unit tests, debugging queries, or developing apps offline.

2. Azure Storage Emulator and Azurite (for Table API)

- If I use the Table API, I can test using Azurite or the older Storage Emulator to simulate table storage locally.

3. SDK support for local connections

- All Cosmos DB SDKs support connecting to the emulator by pointing to its local endpoint and using its master key.
- This makes it easy to run the same code against both local and cloud environments just by changing the connection string.

4. Cosmos DB Data Explorer in Emulator

- The emulator includes a web-based Data Explorer UI that works just like the portal version.
- I can create databases, containers, run queries, and inspect data directly in the browser.

5. Docker version of the emulator

- If I use Linux or want container-based testing, Microsoft also provides a preview version of the Cosmos DB emulator as a Docker container.
- It allows me to include Cosmos DB in automated test environments using Docker Compose.

Using these tools, I can build and test my Cosmos DB-based applications locally with almost the same experience and features as the cloud version.

59. Discuss best practices for managing the Cosmos DB lifecycle within a CI/CD pipeline.

When managing Cosmos DB in a CI/CD pipeline, I follow best practices that ensure my deployments are consistent, safe, and automated. These practices help me handle everything from creating the database to updating containers and seeding test data.

Here are the key best practices:

1. Use Infrastructure as Code (IaC)

- I use tools like ARM templates, Bicep, or Terraform to define Cosmos DB resources such as accounts, databases, and containers.
- This ensures the infrastructure is version-controlled, repeatable, and consistent across environments.

2. Separate environments

- I maintain separate Cosmos DB instances for development, testing, staging, and production.
- Each environment is managed through the pipeline with separate configurations and permissions.

3. Parameterization and secrets

- I use pipeline variables and secret stores like Azure Key Vault to manage settings such as account names, keys, and connection strings securely.

4. Automated testing with Emulator

- For unit and integration testing, I use the Cosmos DB Emulator in the pipeline, either directly on a Windows agent or through a Docker container.
- This avoids using the live database and keeps test runs fast and isolated.

5. Apply schema and indexing changes as code

- Cosmos DB is schema-less, but I still document and version my expected document structures.
- I include index policy updates and TTL settings in my IaC files to track changes over time.

6. Data seeding and migration scripts

- I include scripts in the pipeline to insert reference or test data after deployment.
- For updates or data migrations, I use versioned scripts run through automation tools or Functions.

7. Monitoring and alerts setup

- I configure alerts for RU usage, latency, throttling, and errors as part of the CI/CD setup, so each environment is monitored correctly after deployment.

8. Avoid hard deletion in CI/CD

- In production, I avoid deleting Cosmos DB accounts or containers as part of automation to prevent accidental data loss.

By following these practices, I make sure that Cosmos DB resources are managed in a reliable, secure, and automated way that fits into the overall DevOps lifecycle.

60. Outline the backup and restore options available in Cosmos DB.

Cosmos DB provides automated, continuous backup and restore capabilities to protect data and help recover from accidental deletion or corruption. These options are mostly managed by Azure and require minimal setup.

Here's how the backup and restore options work:

1. Automatic backups (continuous backup)

- Cosmos DB automatically performs continuous backups of data without needing any manual configuration.
- Backups are stored securely and managed by Microsoft.
- With continuous backup enabled, I can restore data to any point in time within the last 30 days.

2. Point-in-time restore (PITR)

- If data is accidentally deleted or corrupted, I can perform a point-in-time restore.
- This allows me to restore the state of a database or container as it was at a specific moment.
- The restored data goes into a new Cosmos DB account, so I can review it or migrate it safely.
- Only Cosmos DB accounts that have continuous backup mode enabled can use this feature.

3. Backup configuration

- When creating a Cosmos DB account, I choose between periodic backup (older option, with snapshots every 4 hours) and continuous backup (the newer and more flexible option).
- Continuous backup is the recommended approach for most applications.

4. No manual snapshots needed

- Cosmos DB handles all backups automatically in the background. I don't need to manage any backup files or schedule.

5. Restore process

- To restore, I use Azure Portal, Azure CLI, or ARM templates.
- I choose the restore point, and Azure will create a new Cosmos DB account with the restored data.

6. Limitations

- Backups are region-specific. If the region is unavailable, the restore might be delayed.
- Backup retention is limited to 30 days in continuous mode.

In short, Cosmos DB provides a reliable, hands-off backup system with the ability to recover data exactly as it was at any moment in the past 30 days using point-in-time restore.

61. What monitoring tools and metrics are important to keep an eye on for Cosmos DB performance?

Monitoring Cosmos DB is important to ensure that my application runs smoothly and efficiently. Azure provides built-in tools and key performance metrics to help me understand how Cosmos DB is behaving and identify any issues.

Here are the main tools and important metrics:

1. Azure Monitor and Insights

- Azure Monitor provides charts, alerts, and logs for Cosmos DB.
- Cosmos DB Insights is a dashboard inside the Azure portal that shows performance, usage, throttling, and storage details.

2. Key metrics to monitor

Request Units (RU/s) usage

- Shows how much of the provisioned throughput is being used.
- High or near-maximum usage may mean I need to scale up or optimize queries.

Throttled requests (429 errors)

- Indicates requests are being rate-limited due to RU limits.
- Frequent throttling affects performance and user experience.

Latency (server-side)

- Measures how long Cosmos DB takes to process read and write requests.
- Helps identify delays in data access.

Availability

- Shows if the service is up and responding.
- Downtime or degraded performance is visible here.

Total Requests

- Gives an idea of how busy the system is.
- Helps in capacity planning and scaling decisions.

Data storage

- Tracks the size of data stored in each container.
- Useful for cost management and partition planning.

3. Diagnostic Logs

- Logs provide detailed information about operations, errors, and query performance.
- These logs can be sent to Log Analytics or a storage account for deeper analysis.

4. Alerts

- I set up alerts based on thresholds (e.g., if RU usage exceeds 80% or if throttling occurs).
- Helps in catching performance problems early before they impact users.

5. Query metrics

- I can use SDKs or the Data Explorer to check individual query performance.
- Shows RU consumption and duration of specific queries.

In summary, by using Azure Monitor and keeping a close watch on RU usage, throttling, latency, and query performance, I can keep Cosmos DB running smoothly and adjust resources or queries before any issues affect the application.

© Shubham Wadekar

62. Explain the process of disaster recovery in the context of Cosmos DB.

Cosmos DB is built with high availability and disaster recovery in mind. The platform is globally distributed by design, and it includes several built-in features that automatically handle disaster recovery with minimal manual intervention.

Here's how the disaster recovery process works:

1. Automatic multi-region replication

I can configure my Cosmos DB account to replicate data across multiple Azure regions.

This ensures that if one region becomes unavailable due to a disaster (like an outage or natural event), the other region(s) can take over.

2. Automatic failover

When I enable automatic failover, Cosmos DB will automatically route traffic to the next region in my defined priority list if the primary region goes down.

This ensures my application continues to work without any code changes.

3. Manual failover (if needed)

If I don't enable automatic failover, I can manually trigger a failover through the Azure portal, CLI, or ARM templates.

This gives me more control over when and how failover happens.

4. Multi-region writes

If I enable multi-region writes, all regions can accept both read and write operations.

This provides better resilience because even if one region is down, users in other regions can still write data.

5. Data consistency options

Cosmos DB uses a distributed consistency model, and I can choose the right level (eventual, session, strong, etc.) to balance performance and availability across regions.

During a disaster, Cosmos DB ensures that the chosen consistency level is maintained while switching regions.

6. Point-in-time restore for data loss recovery

If data is corrupted or accidentally deleted (not a region-level failure), I can use point-in-time restore (with continuous backup enabled) to recover data from any point within the last 30 days.

7. Monitoring and alerting

I use Azure Monitor to track region health, replication lag, and latency.

I set up alerts to notify me of region outages, degraded performance, or failover events.

In short, Cosmos DB handles most disaster recovery scenarios automatically through global distribution, regional failover, and continuous backup. As a developer or admin, I just need to configure replication and failover settings properly to be ready for such events.

63. Can you monitor the throughput usage and performance metrics of your Cosmos DB account? If so, how?

Yes, Cosmos DB provides several ways to monitor throughput usage and performance. These tools help me track how well the database is performing and detect any issues like throttling or high latency.

Here's how I monitor performance:

1. Azure Monitor

Azure Monitor automatically collects metrics from my Cosmos DB account.

I can view charts and reports on metrics like RU/s usage, throttled requests, latency, availability, and storage usage.

2. Cosmos DB Insights

This is a pre-built dashboard in the Azure portal under the "Insights" section of Cosmos DB.

It gives me a visual overview of important performance indicators like:

- Total requests
- RU consumption per container
- Number of throttled requests
- Read and write latency
- Indexing and storage usage

3. Metrics explorer

I can use Metrics Explorer in the portal to create custom charts.

For example, I can create a graph showing RU usage over time to help understand peak load periods.

4. Diagnostic logs

I can enable diagnostic settings to send logs and metrics to Log Analytics, a storage account, or Event Hubs.

These logs provide deeper visibility into operations and queries.

5. Query metrics (via SDKs or Data Explorer)

When I run queries, I can check their RU cost and execution time either in the Data Explorer or through the SDKs.

This helps me identify which queries are expensive and need optimization.

6. Alerts

I set up alerts for specific thresholds, such as:

- RU consumption > 80%
- Throttled requests > 0
- Latency > 500 ms

These alerts notify me by email or through other channels when something needs attention.

By regularly monitoring these metrics, I can scale throughput appropriately, optimize queries, and ensure the Cosmos DB environment is healthy and performing well.

© Shubham Wadekar